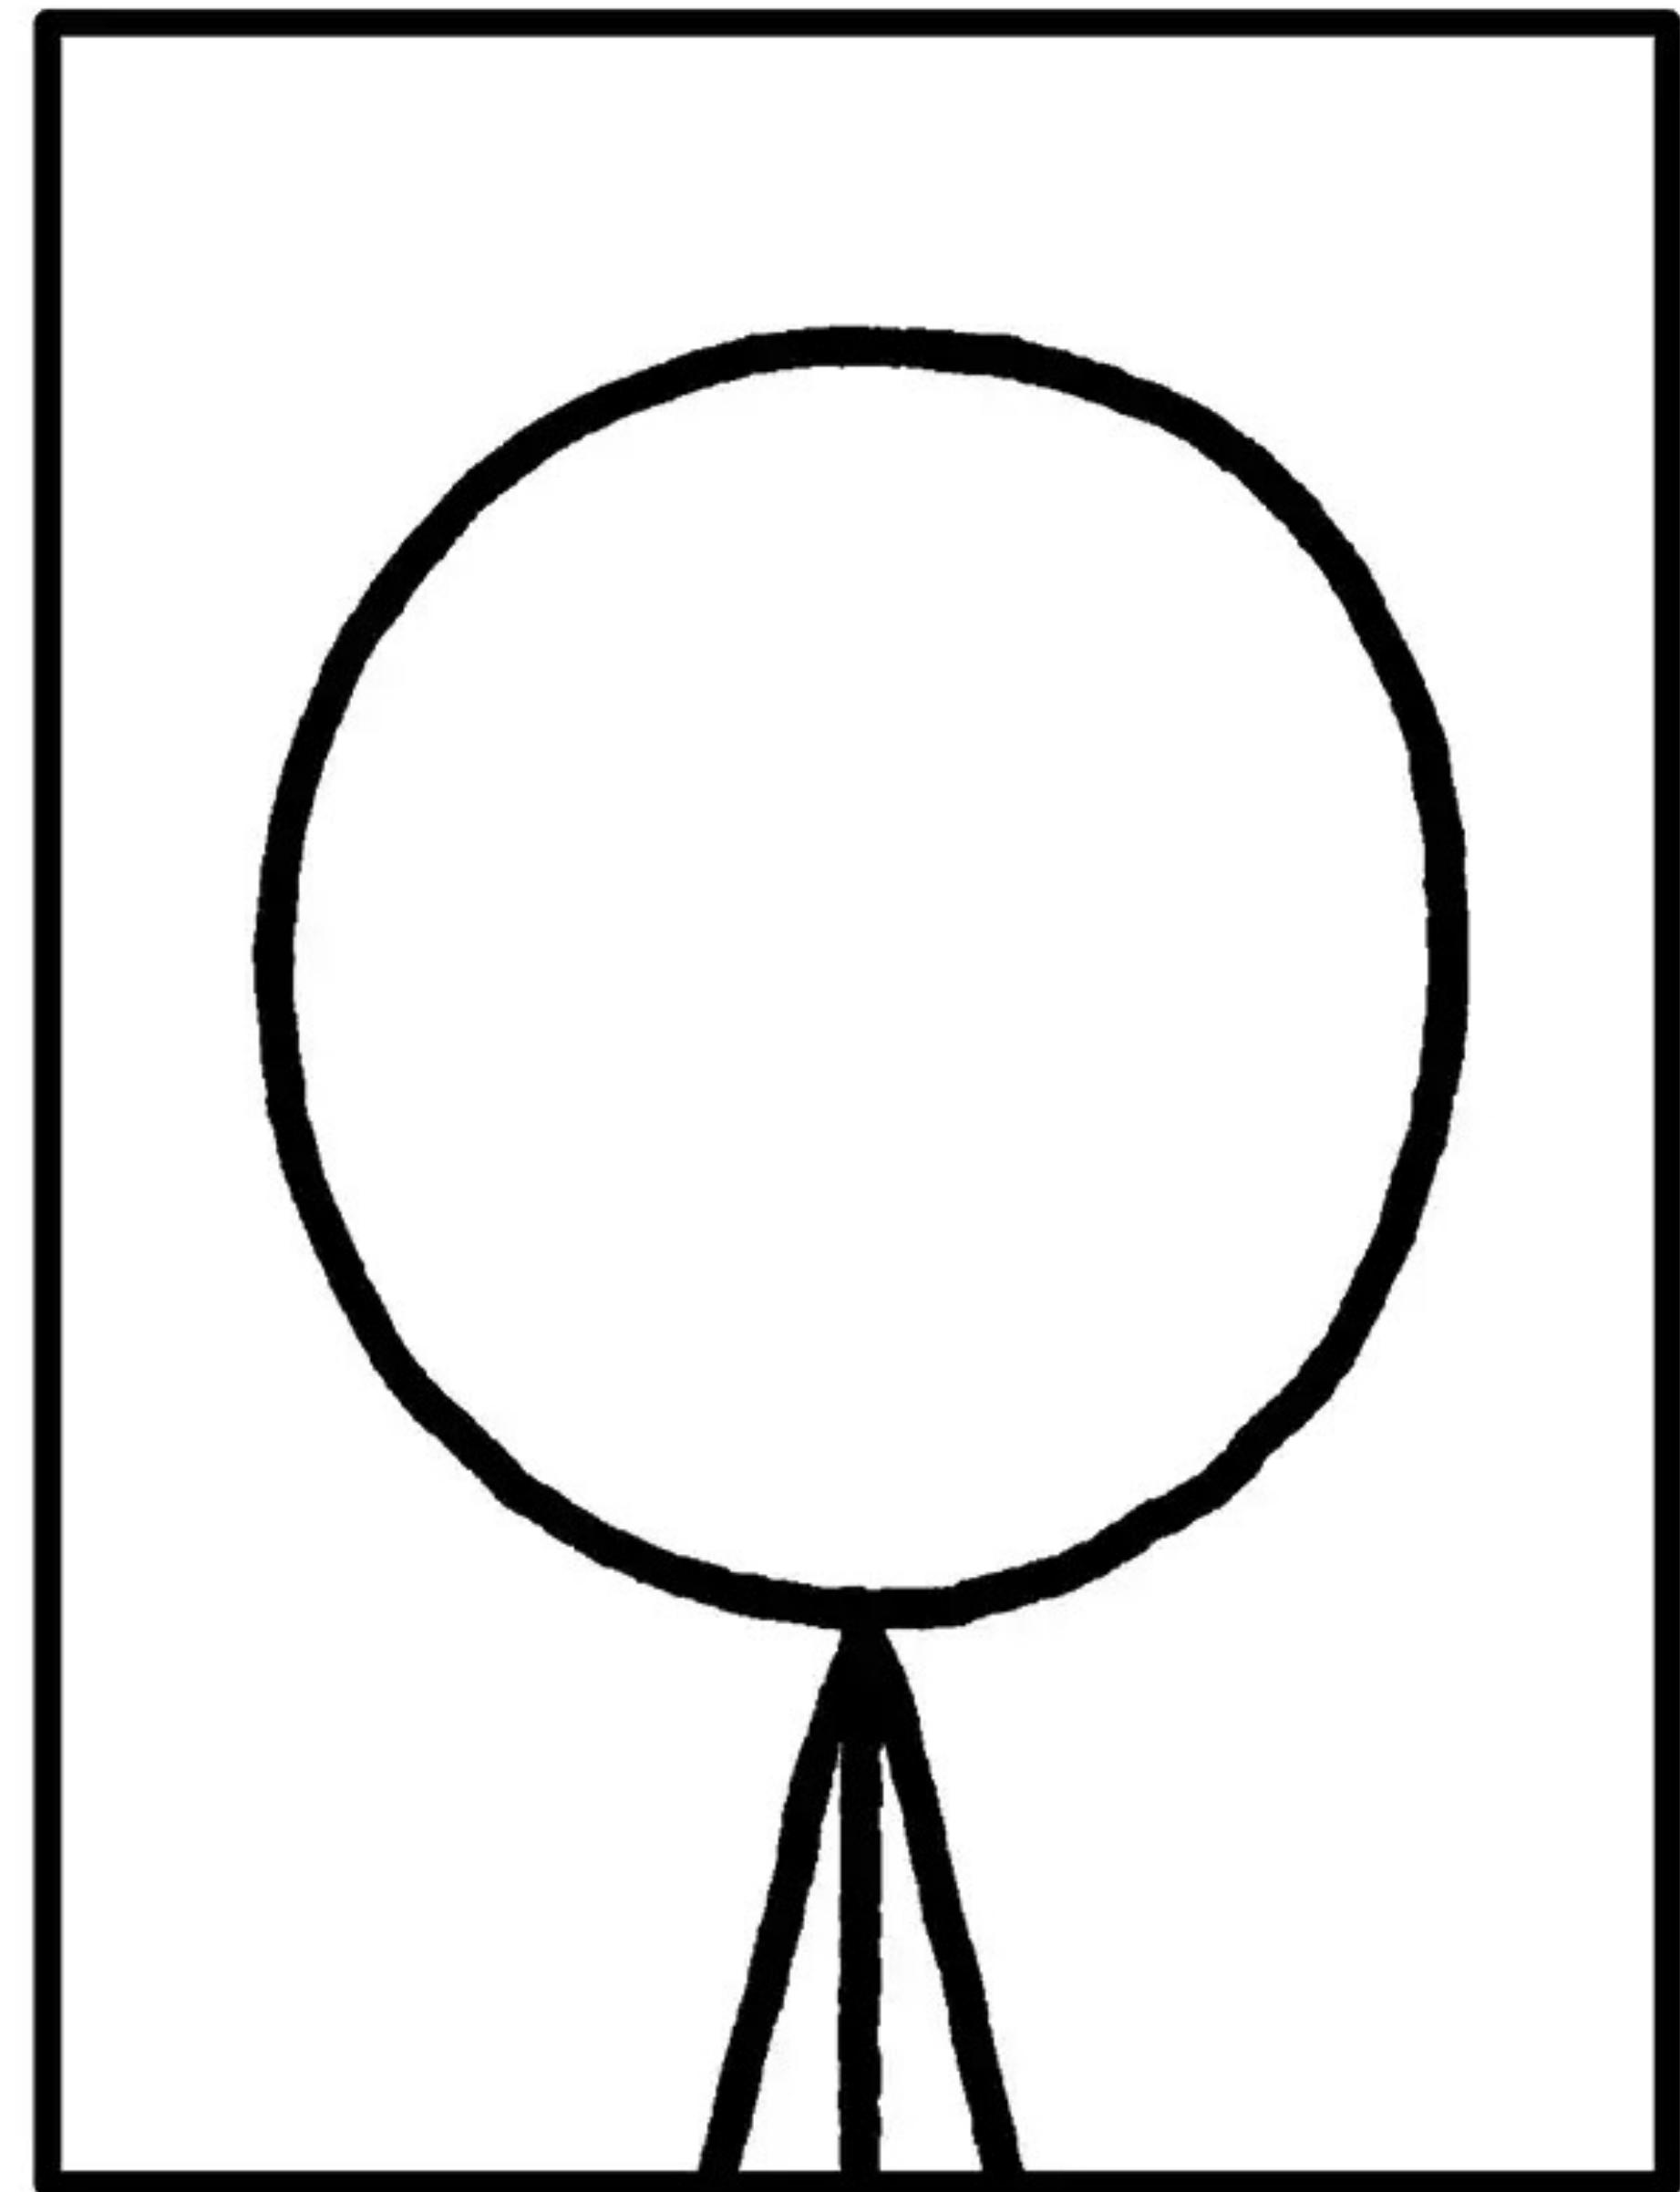


Bayesian regression modeling: Theory & practice

Part 1: Bayesian basics & simple linear regression

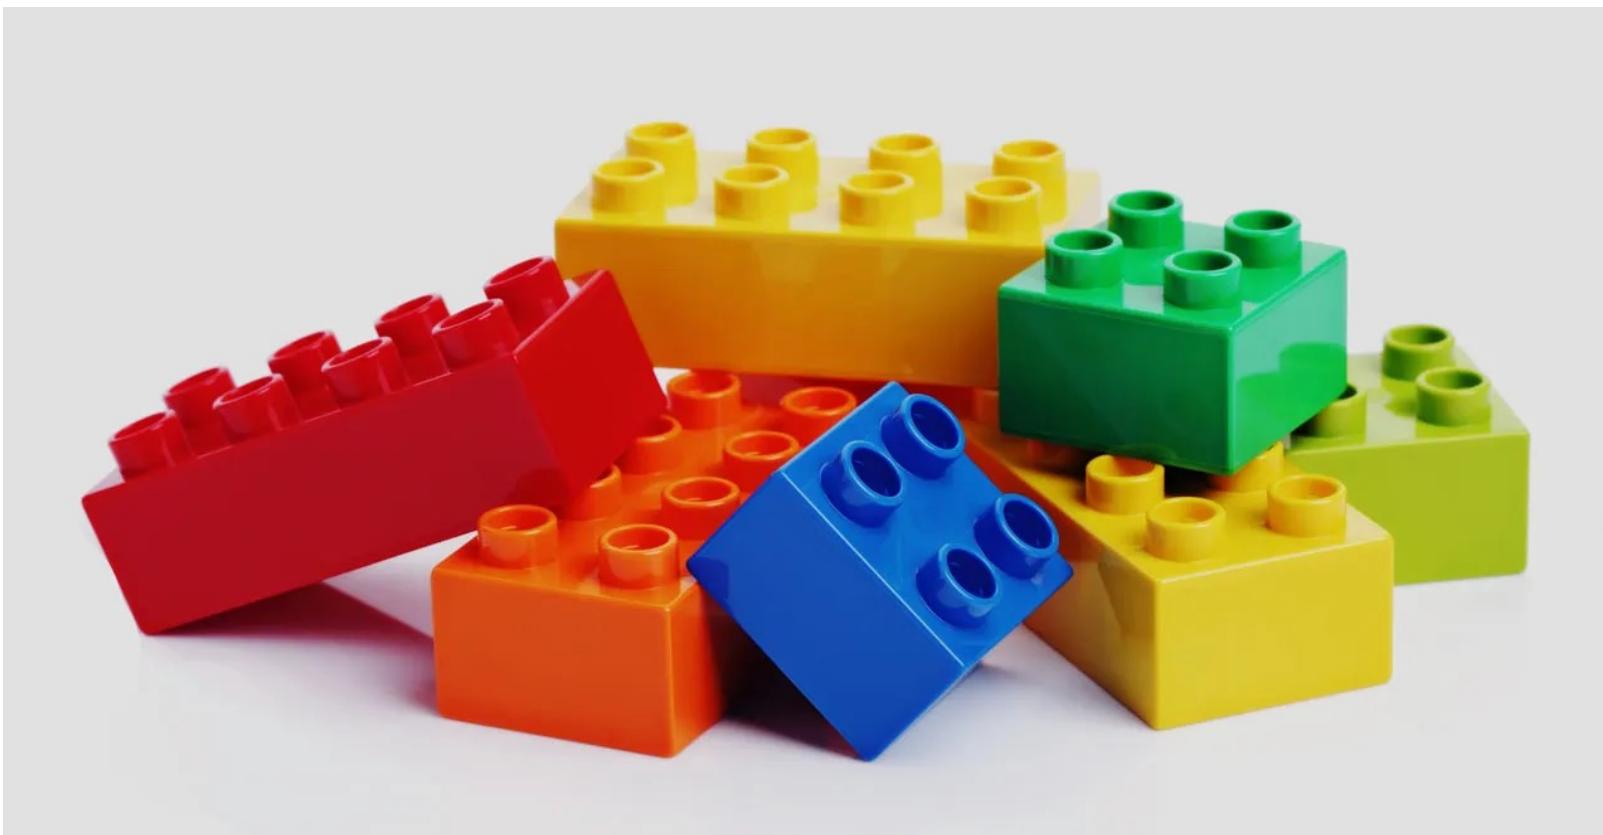
Michael Franke



Hi! It's me.

Pros of BDA

- ▶ well-founded & totally general
- ▶ easily extensible / customizable
- ▶ more informative / insightful
- ▶ stimulates view: “models as tools”



Cons of BDA

- ▶ not yet fully digested by community
- ▶ possibly computationally complex
- ▶ less ready-made, more hands-on
- ▶ requires thinking (wait, that's a pro!)
 - last two points less valid than 10 years ago

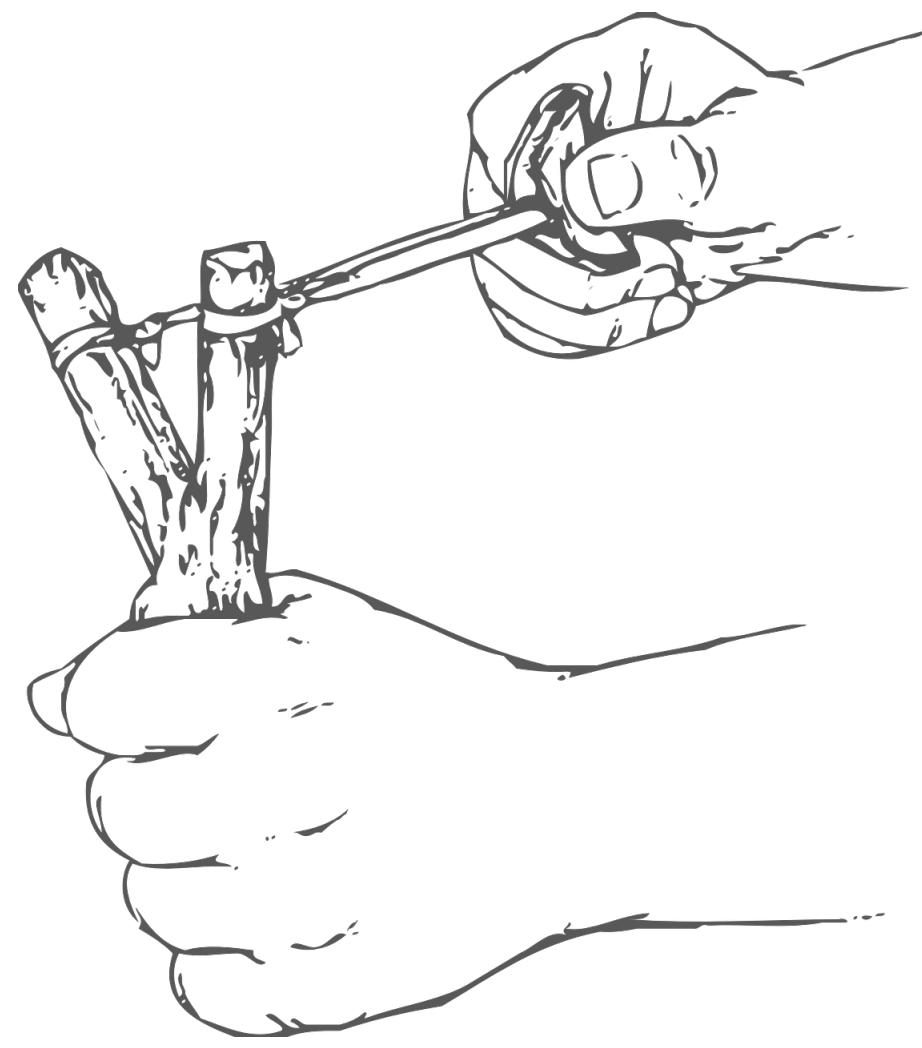




this course

Main learning goals

1. become comfortable in applying multi-level GLMS
 - a. determine the appropriate (kind of) model for a given problem
 - b. implement, run and interpret the Bayesian model
 - c. draw conclusions regarding evidence for/against research questions
2. become oriented in the landscape “beyond standard LMs”
 - a. know enough to determine what is relevant for you
 - b. know enough to quickly find and cover what you need
3. understand key concepts of intermediate BDA
 - a. multi-level modeling
 - b. non-standard link functions, non-linear models ...
 - c. model comparison w/ Bayes factors and cross-validation
 - d. Bayesian computation (MCMC)
 - e. causal inference





no coins, please!

How this works

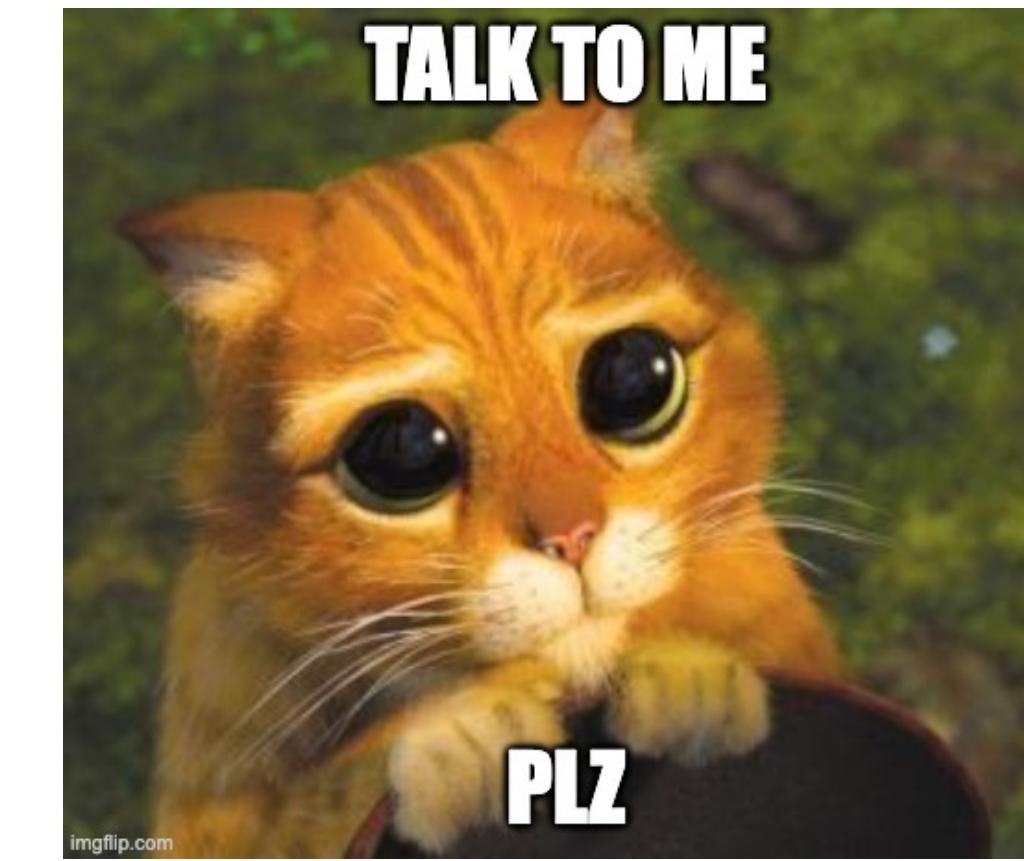
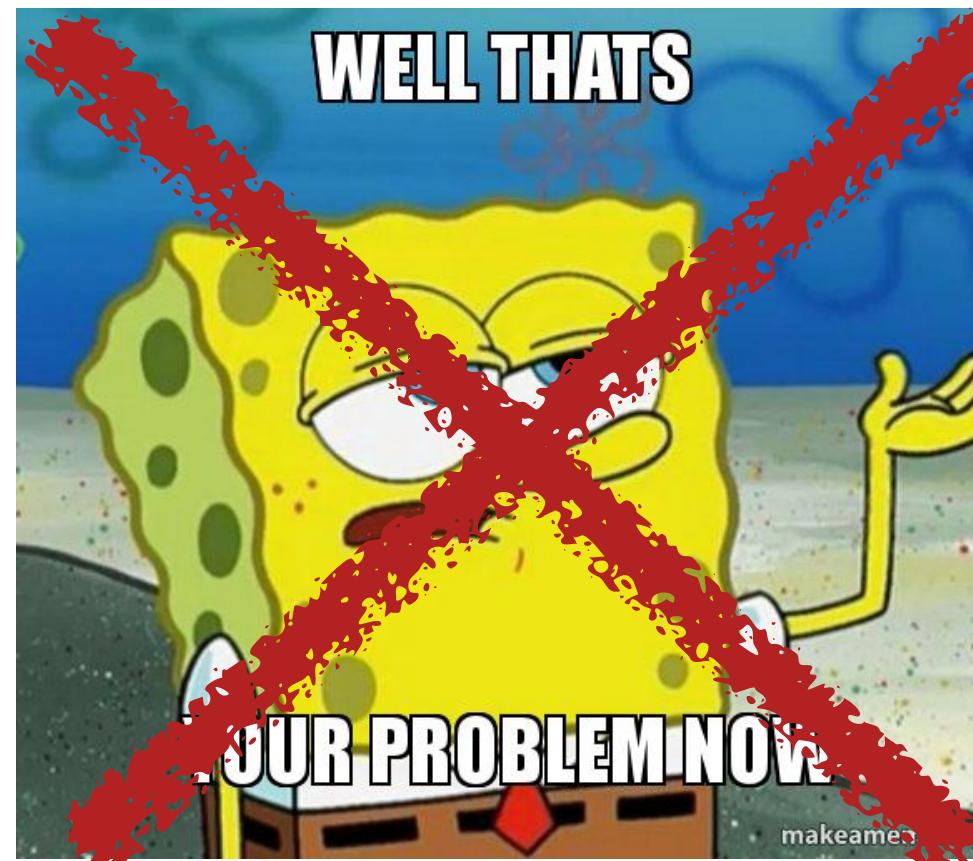


- ▶ class from 2:00 – 6:00 pm
 - w/ some shorter and one longer break
 - some live demos
- ▶ web-book w/ tutorials and exercises
 - live version: [here](#)
 - GitHub code: [here](#)
- ▶ discussion of exercises
 - on Slack (or other means)
 - first part of class
- ▶ contribute (if you want)
 - e.g., send pull-requests to populate the cheat sheet

Here is the a plan

Day 1	Day 2	Day 3	Day 4	Day 5
course structure BDA basics	priors predictions	GLMs & beyond	model comparison	non-linear models
Bayesian regression categorical predictors	multi-level models	Bayesian computation	causal inference	🛠

exercises exercises exercises exercises exercises exercises exercises



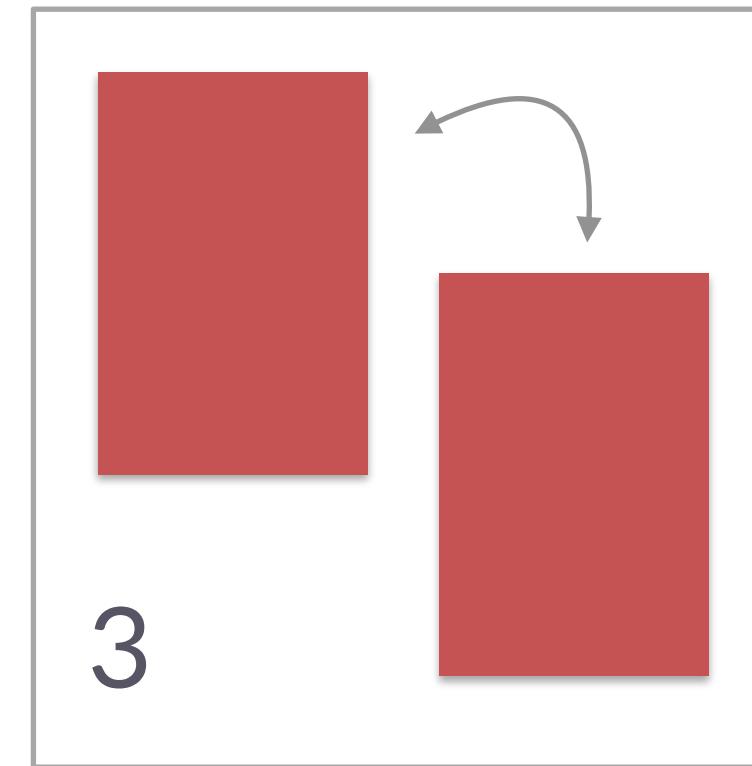
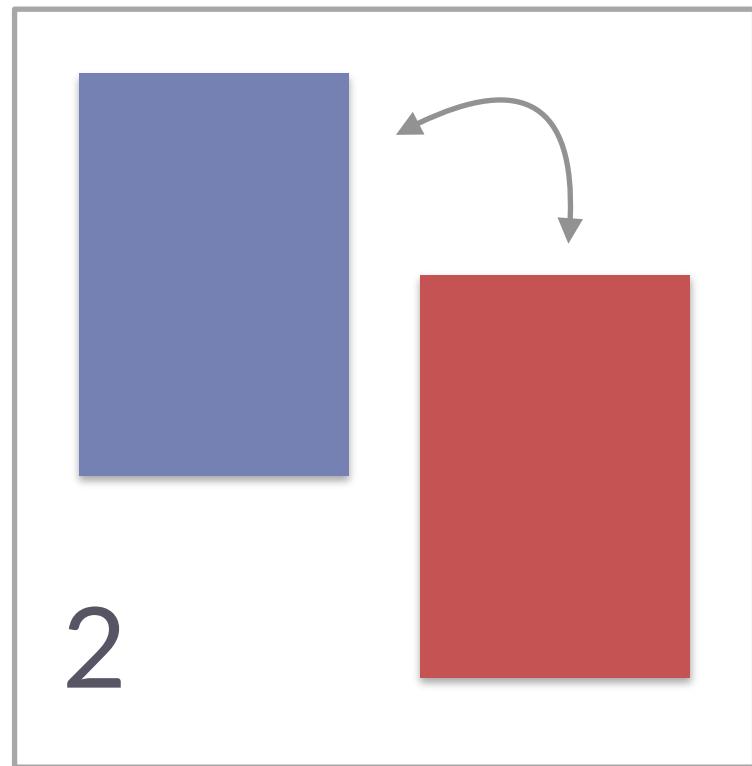
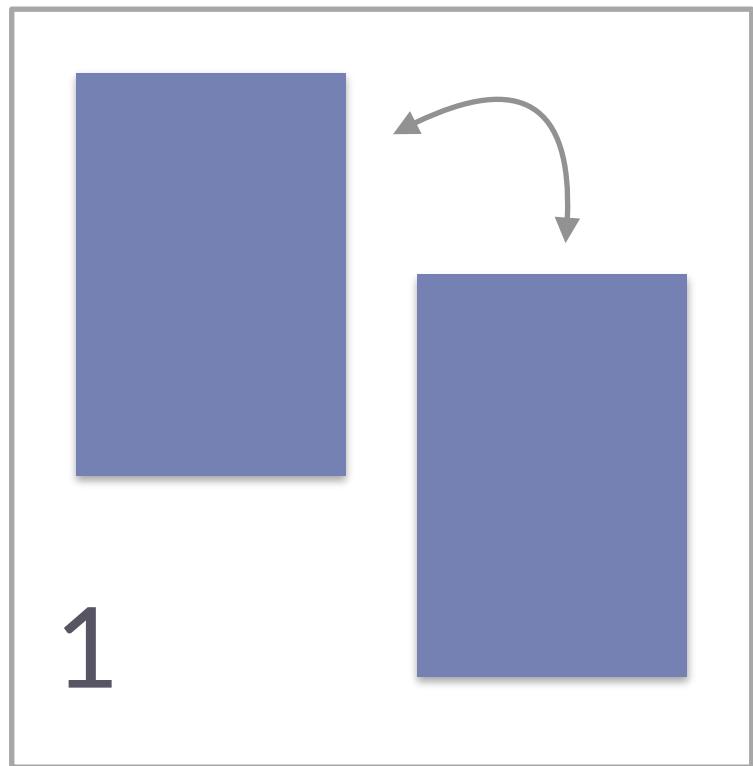


Bayesian modeling

Three-card problem

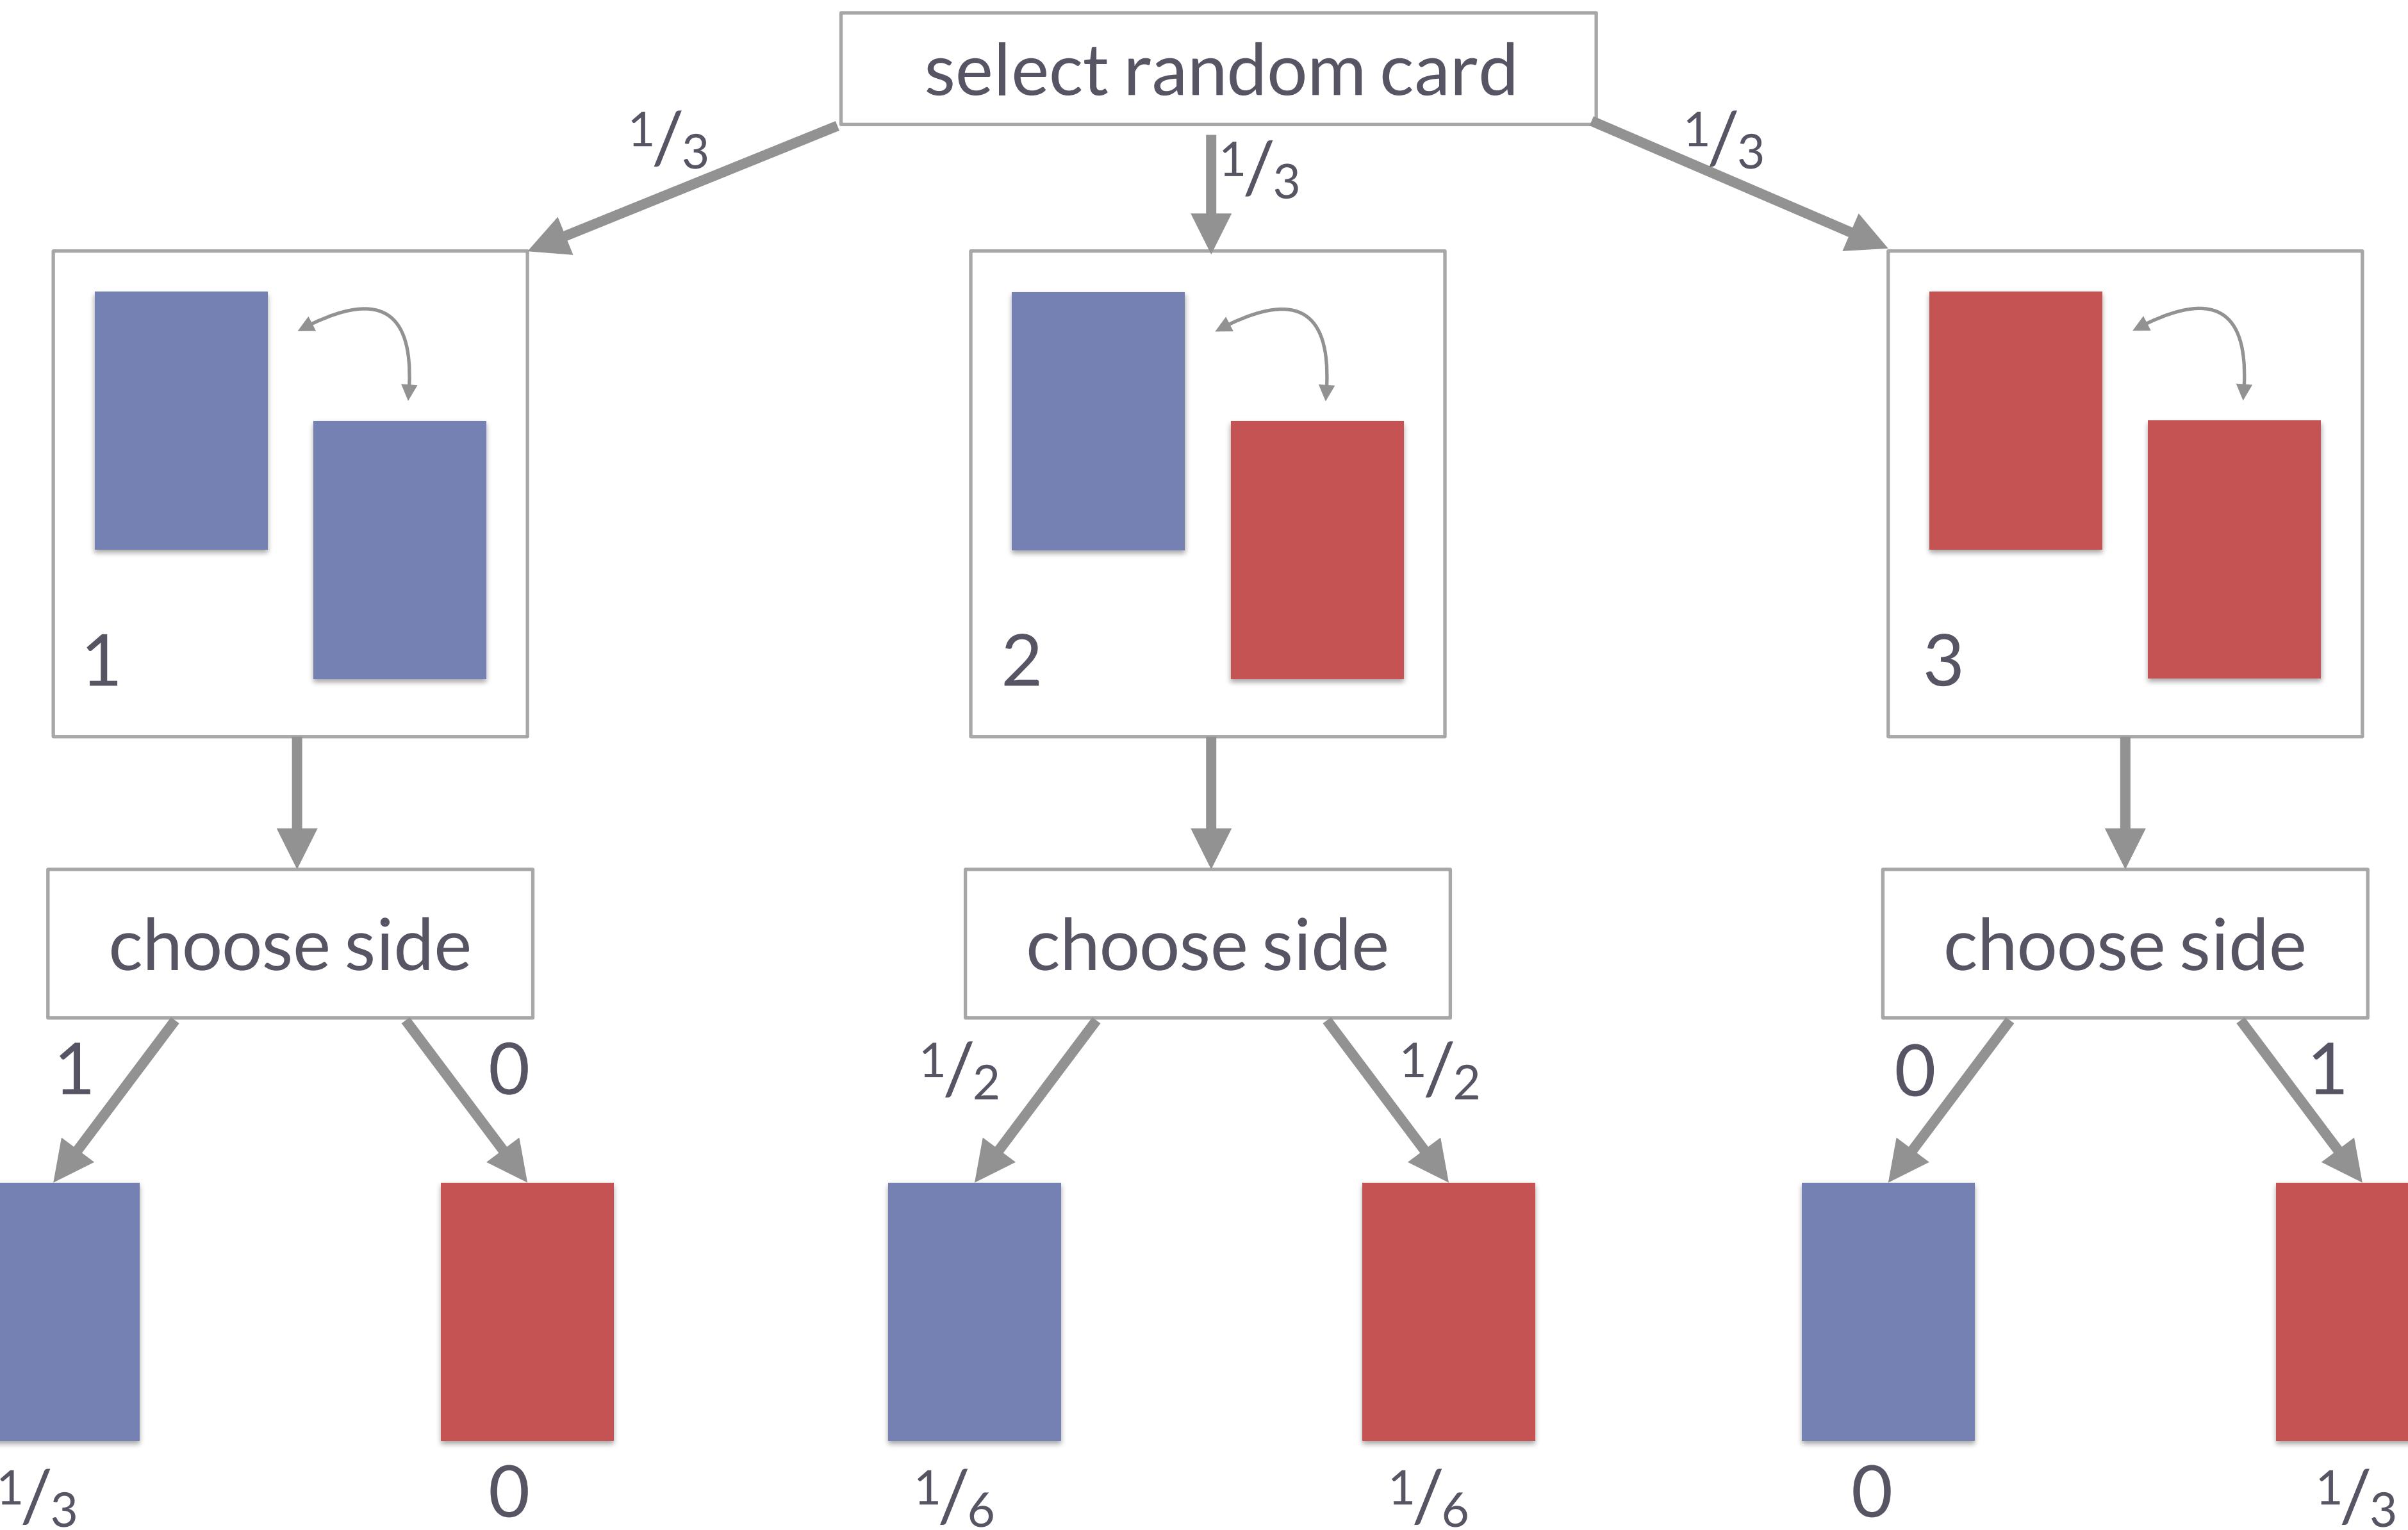
problem statement

- ▶ Sample a card (uniformly at random).
- ▶ Choose a side of that card to reveal (uniformly at random).
- ▶ What's the probability that the side you do not see is **BLUE**, given that the side you see is **BLUE**?



Three-card problem

data-generating process



Conditional probability and Bayes rule

for the three-card problem

- conditional probability

$$P(A | B) = \frac{P(A \cap B)}{P(B)}$$

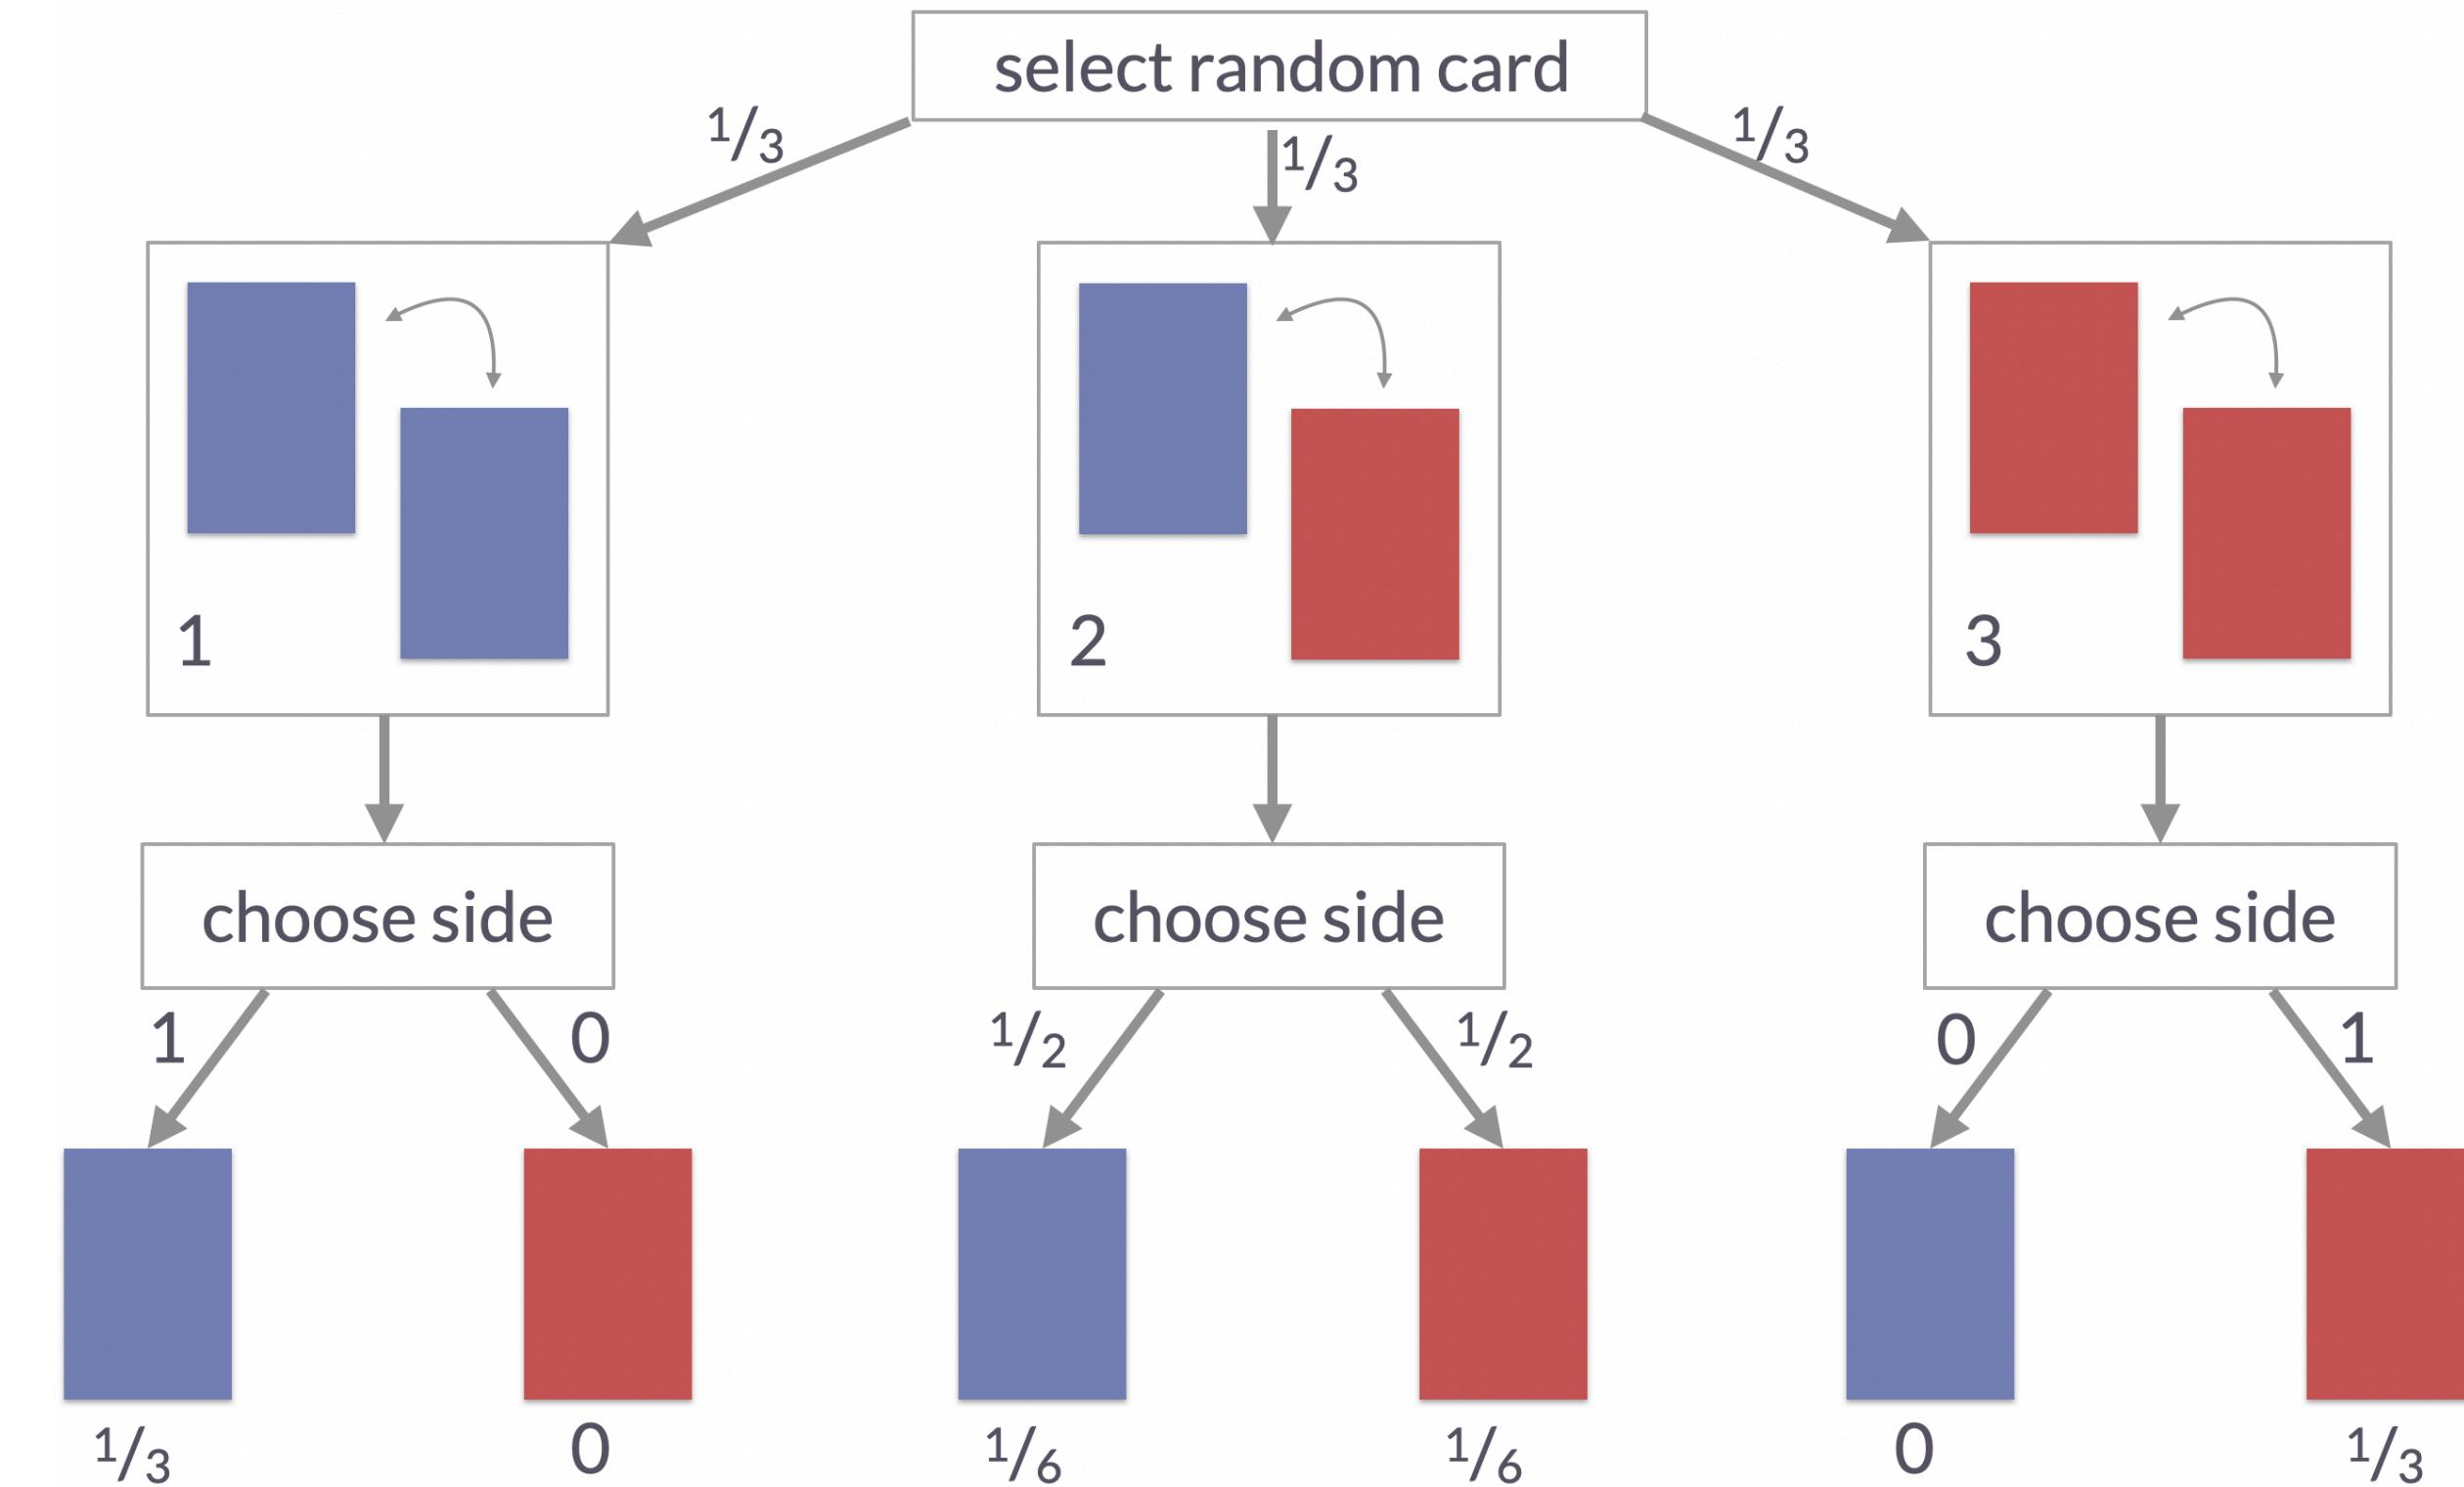
- Bayes rule

$$P(A | B) = \frac{P(B | A)P(A)}{P(B)}$$

- Applied to three-card problem:

$$P(\text{card 1} | \text{blue}) = \frac{P(\text{blue} | \text{card 1}) P(\text{card 1})}{P(\text{blue})}$$

$$= \frac{1 \times \frac{1}{3}}{\frac{1}{2}} = \frac{2}{3}$$



“reasoning from observed effect to latent cause via a model of the data-generating process”

Bayesian data analysis

in a nutshell

- ▶ BDA is about what we *should* believe given:
 - some observable data, and
 - our model of how this data was generated
(a.k.a. **the data-generating process**)

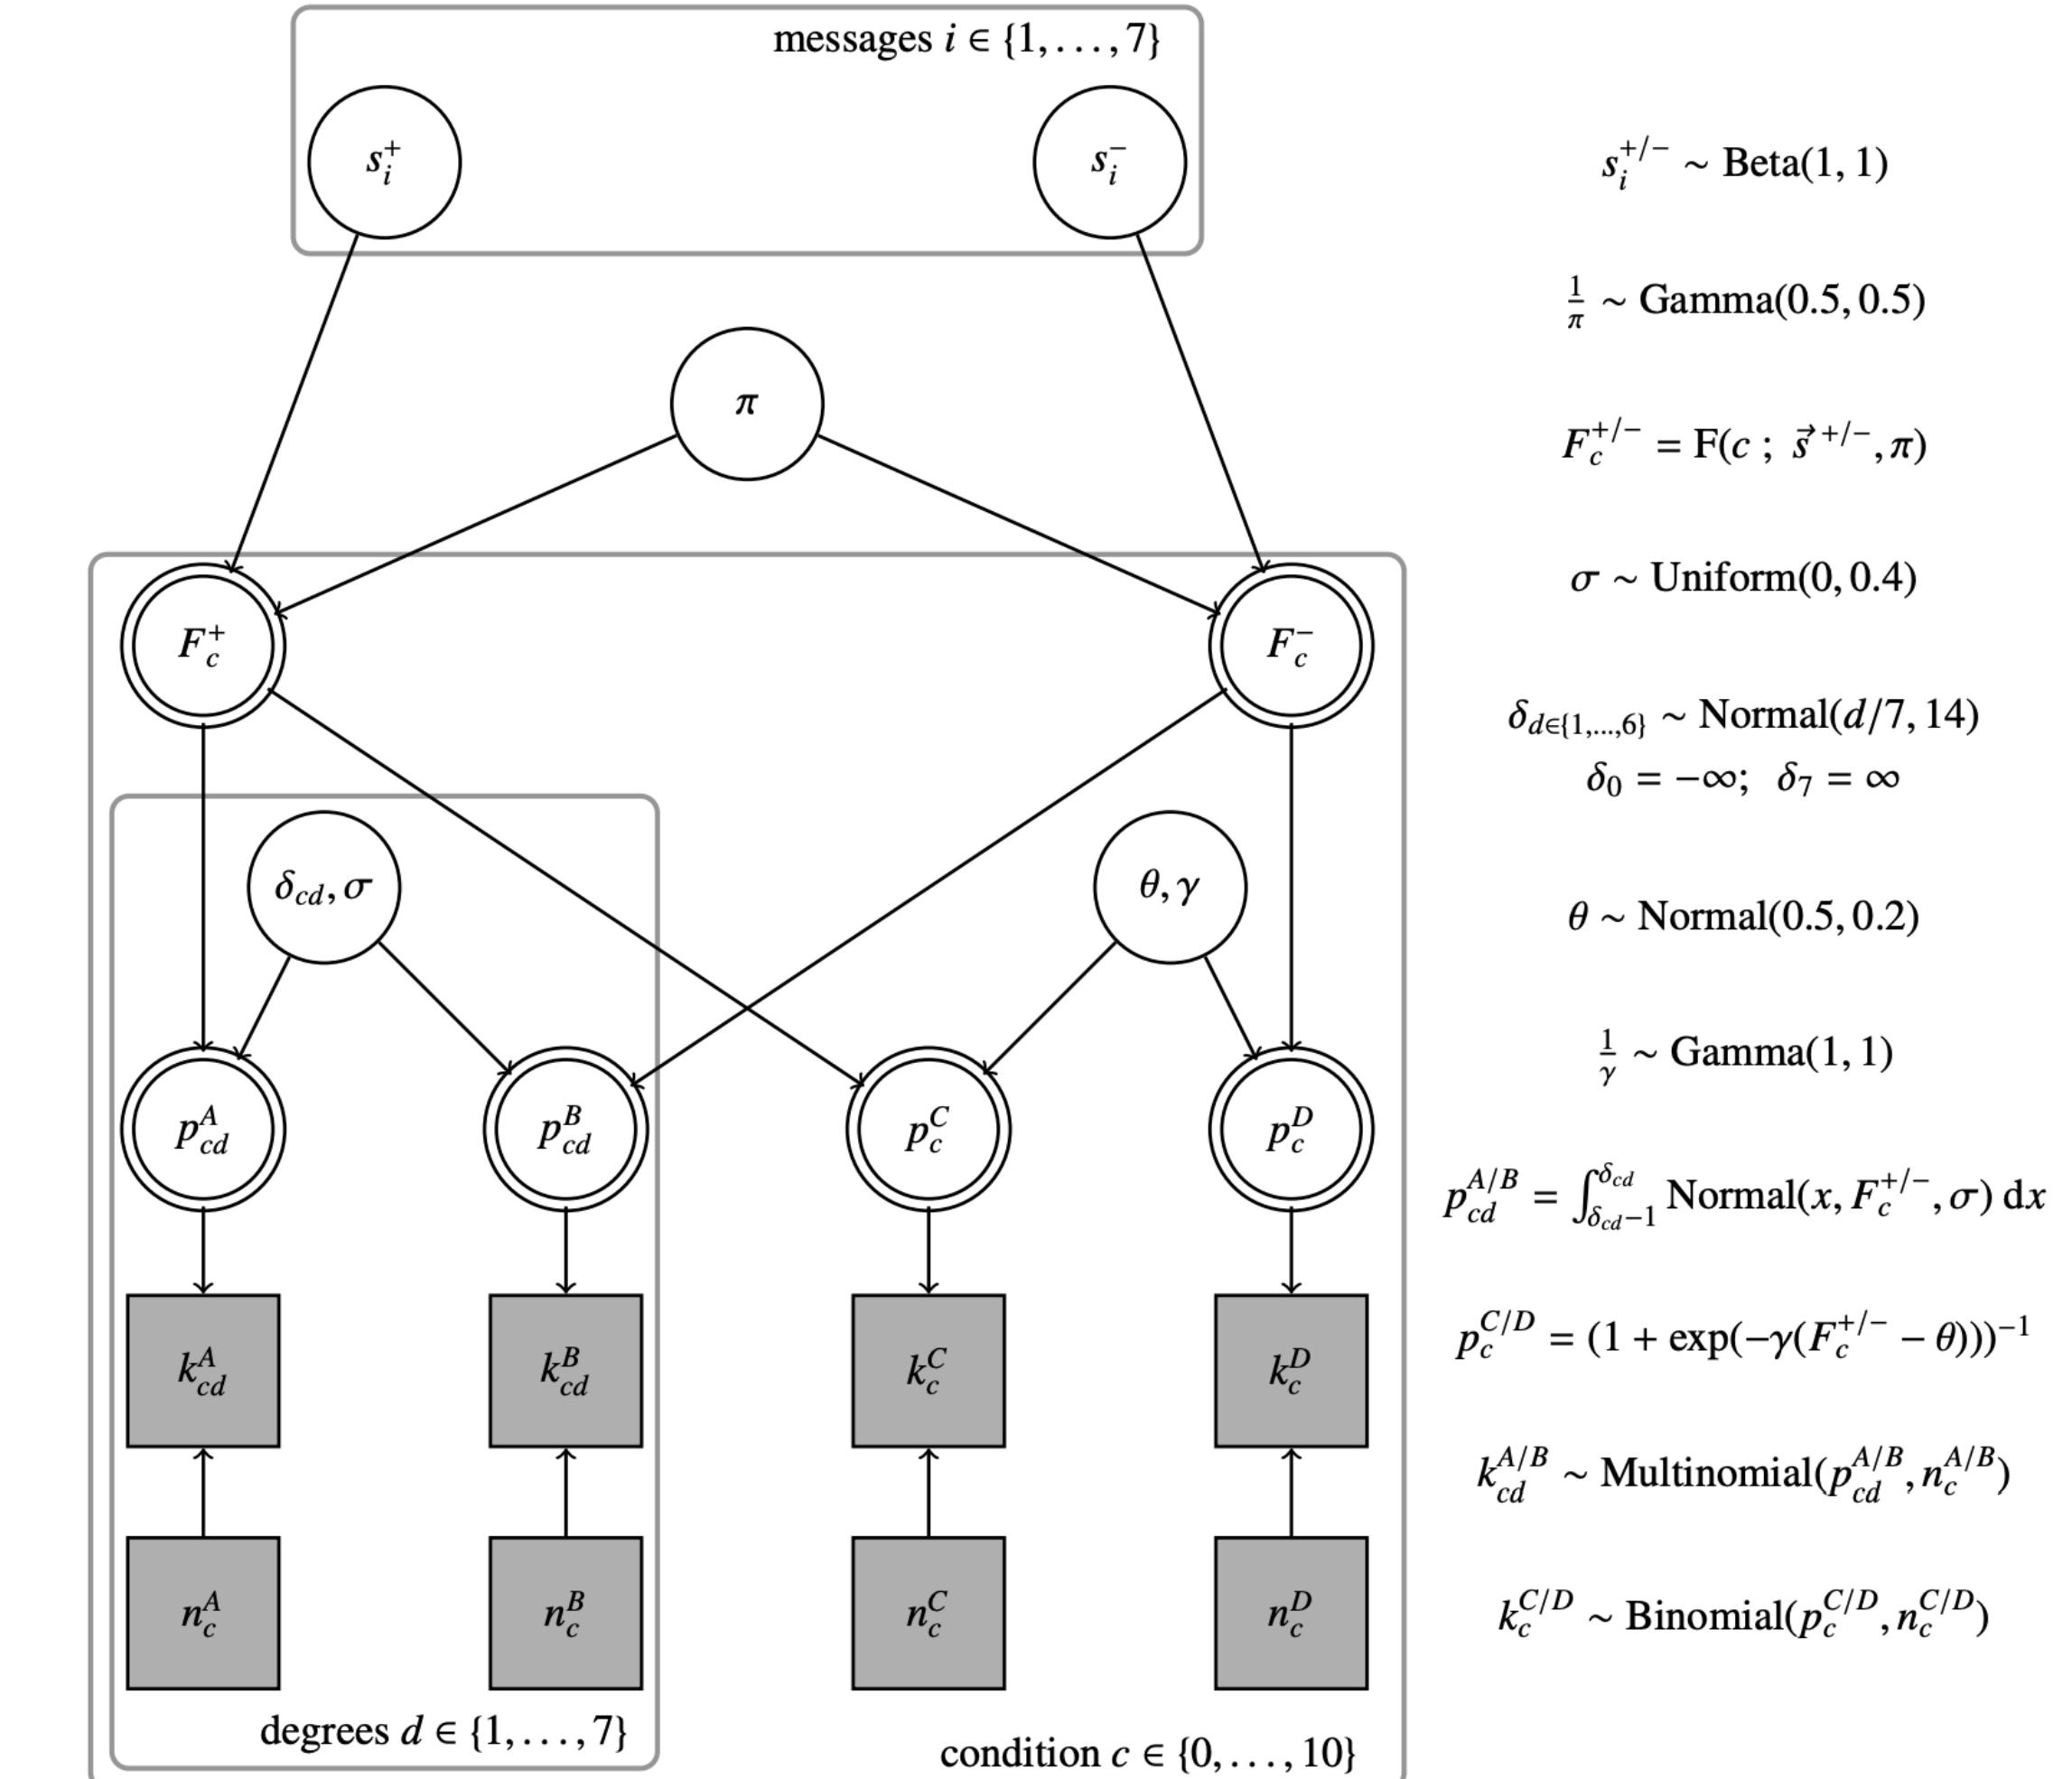
▶ our best friend will be **Bayes rule**

- e.g., for **parameter inference**:

$$P(\theta | D) \propto \underbrace{P(\theta)}_{\text{posterior}} \times \underbrace{P(D | \theta)}_{\text{prior likelihood}}$$

- or, for **model comparison**:

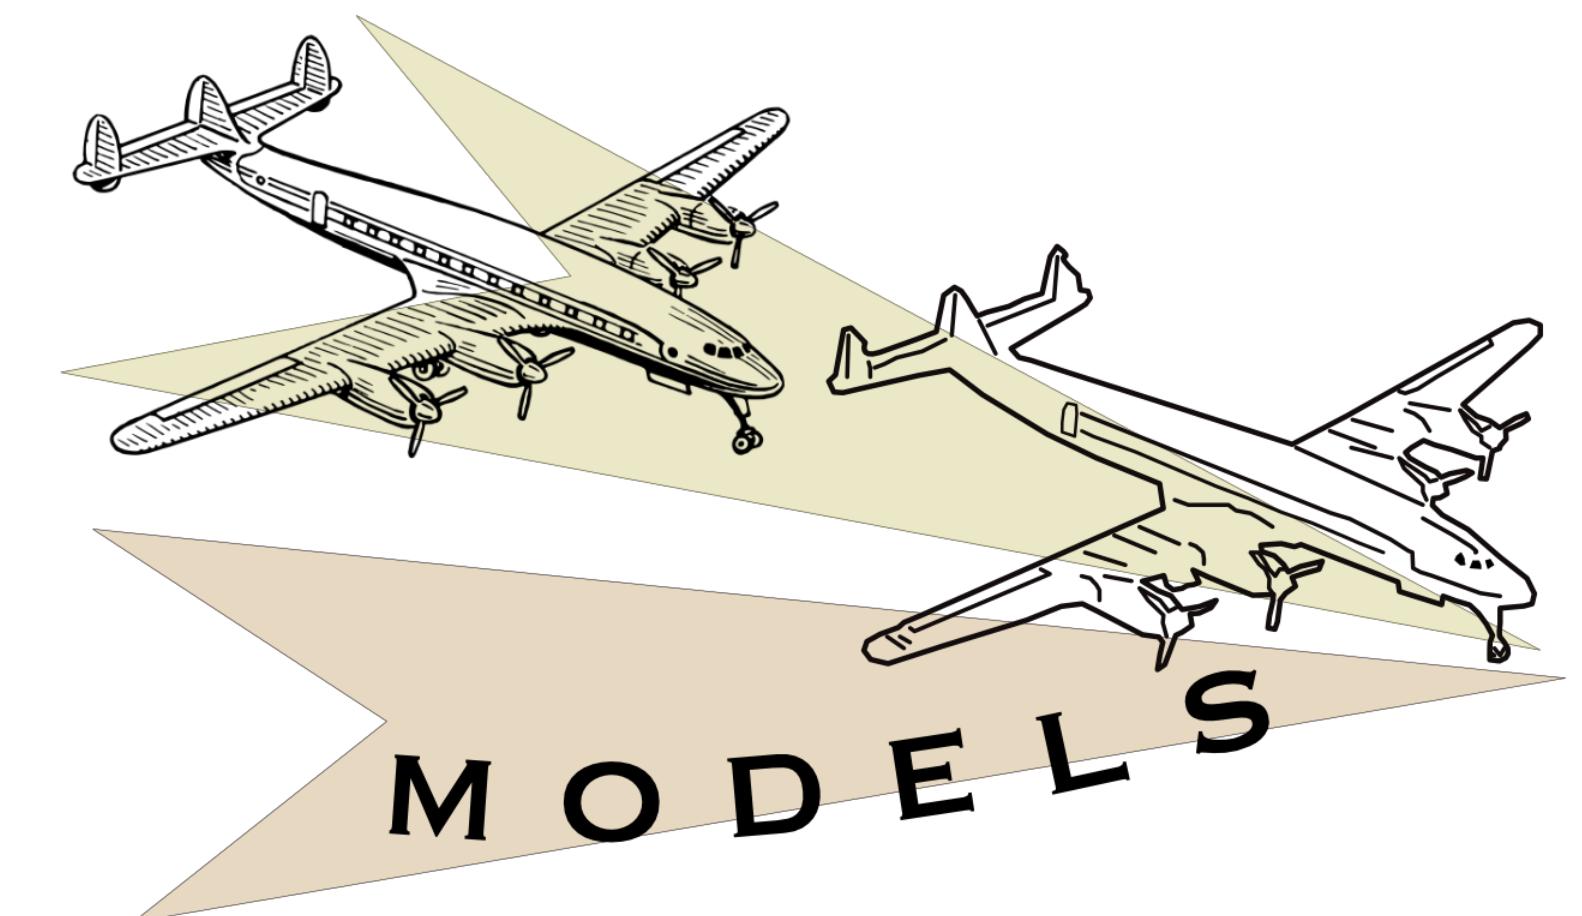
$$\frac{P(M_1 | D)}{P(M_2 | D)} = \underbrace{\frac{P(D | M_1)}{P(D | M_2)}}_{\text{Bayes factor}} \times \underbrace{\frac{P(M_1)}{P(M_2)}}_{\text{prior odds}}$$

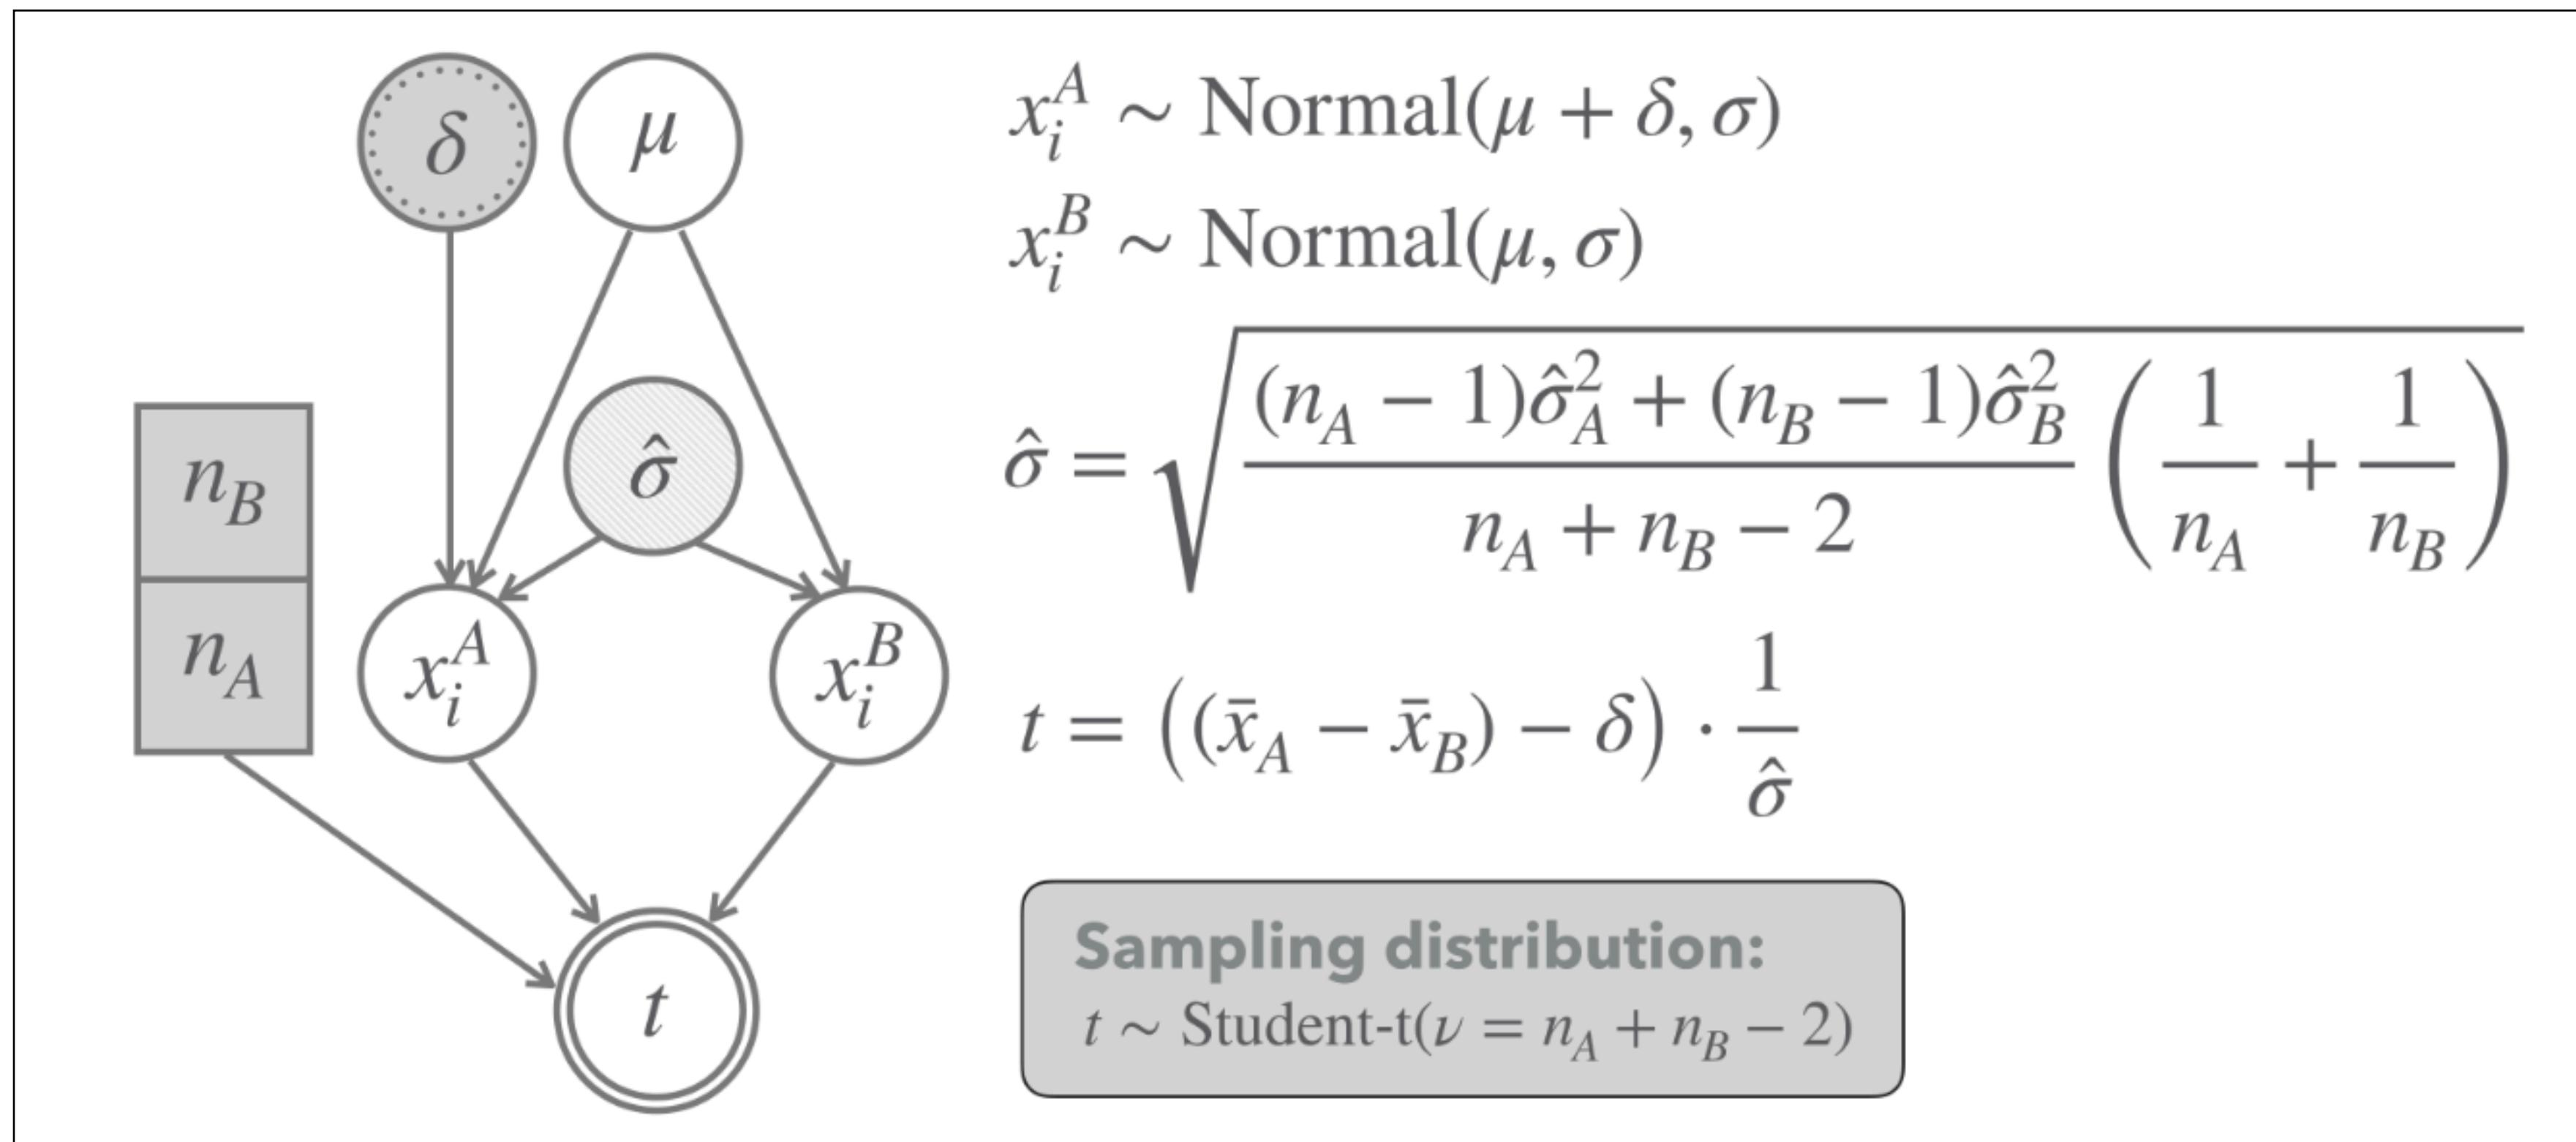
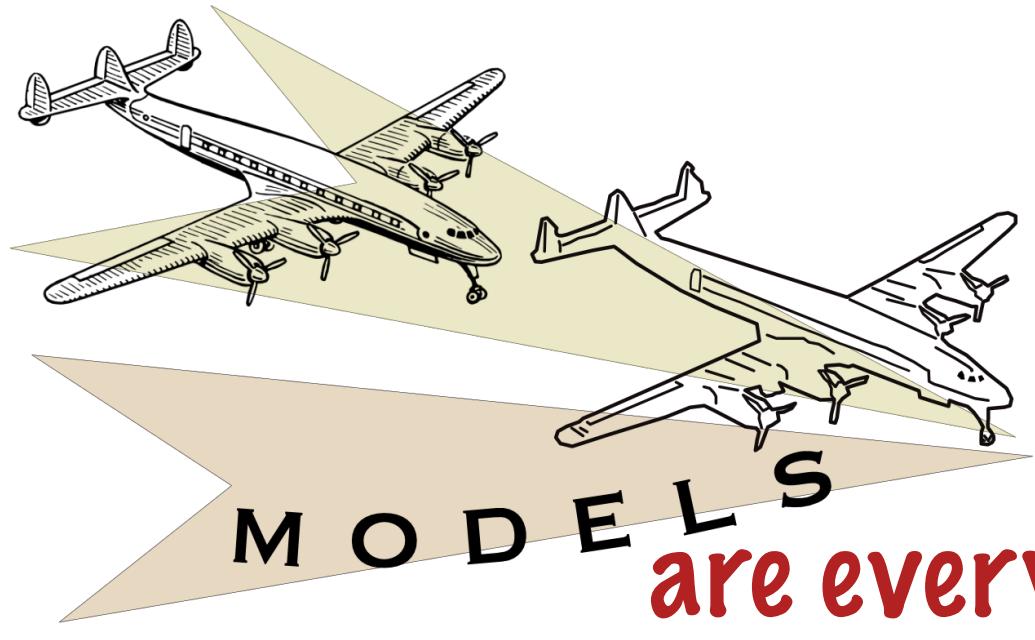


Statistical models

likelihoods from a data-generating process

- ▶ A **statistical model** is a condensed formal representation, following common conventional practices of formalization, of the assumptions we make about what the data is and how it might have been generated by some (usually: stochastic) process.
- ▶ “All models are wrong, but some are useful.” (Box 1979)
- ▶ a **Bayesian statistical model** of stochastic process generating data D consists of:
 - a vector of **parameters** θ
 - a **likelihood function**: $P(D | \theta)$
 - a **prior** distribution: $P(\theta)$
- ▶ among other things, we can use a model for **inference**:
 - posterior distribution: $P(\theta | D) \propto P(D | \theta) P(\theta)$





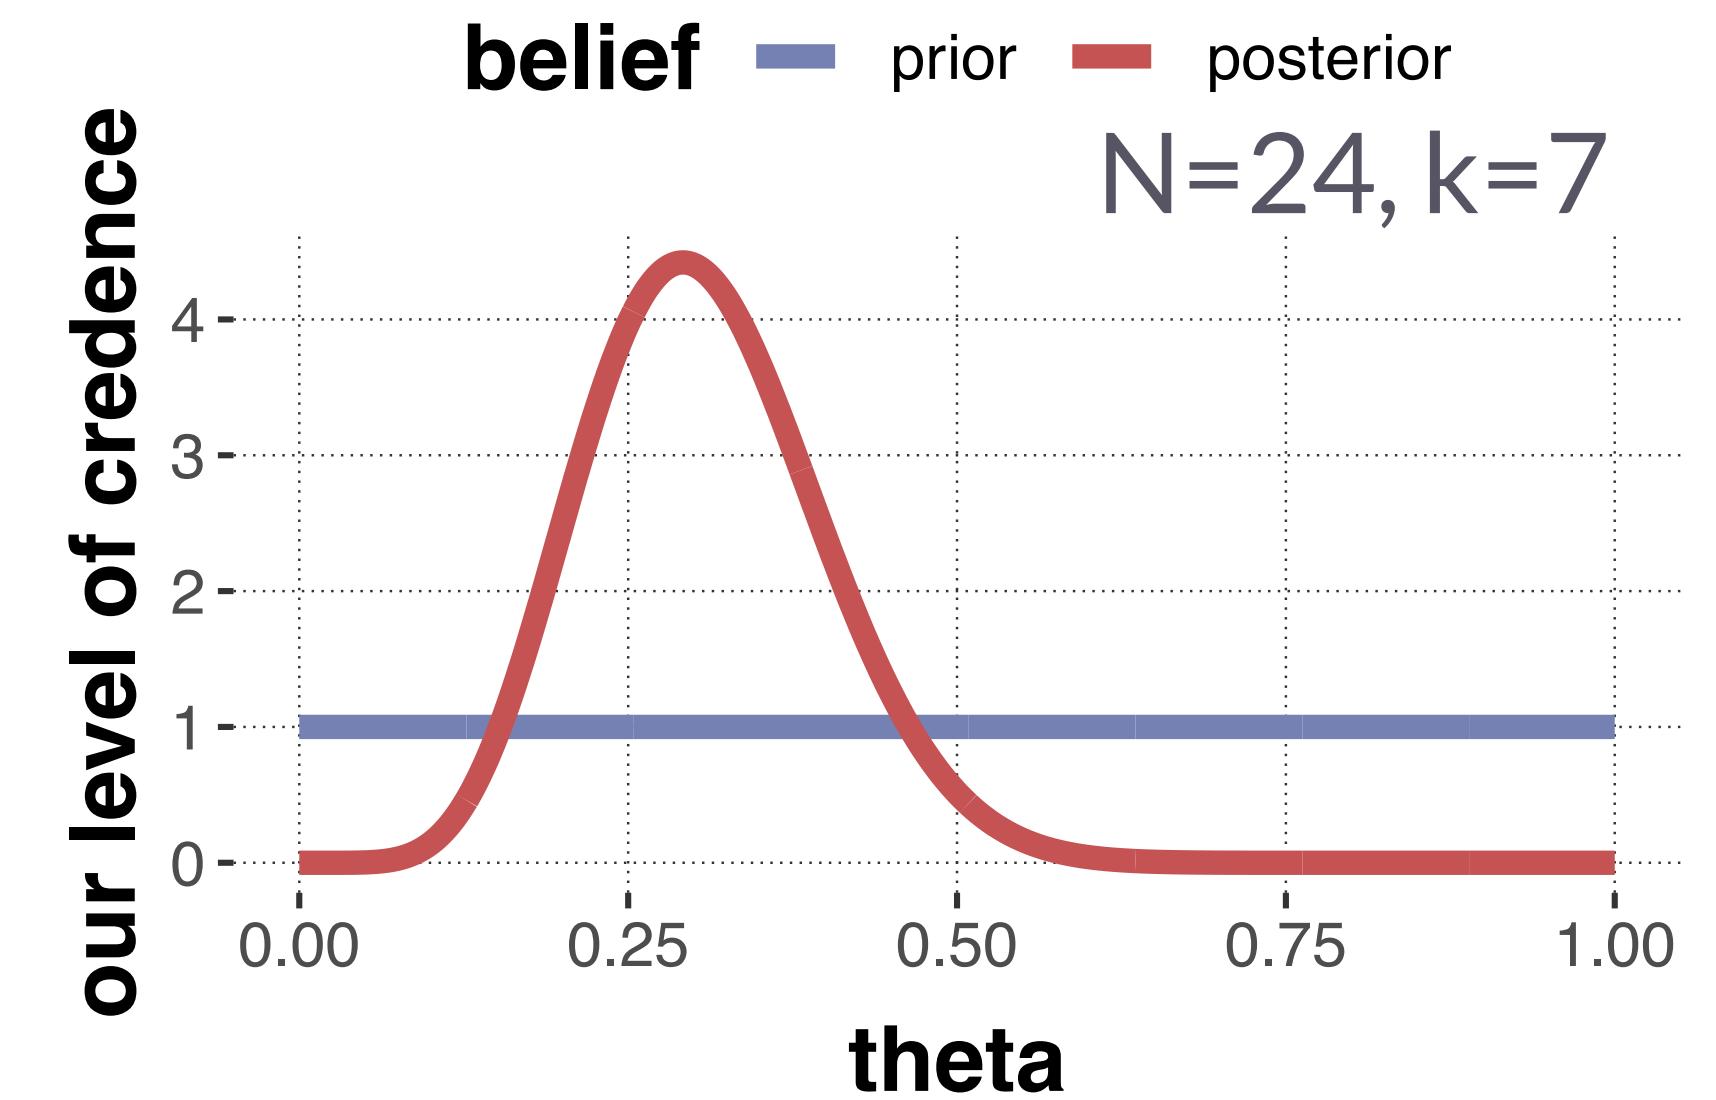
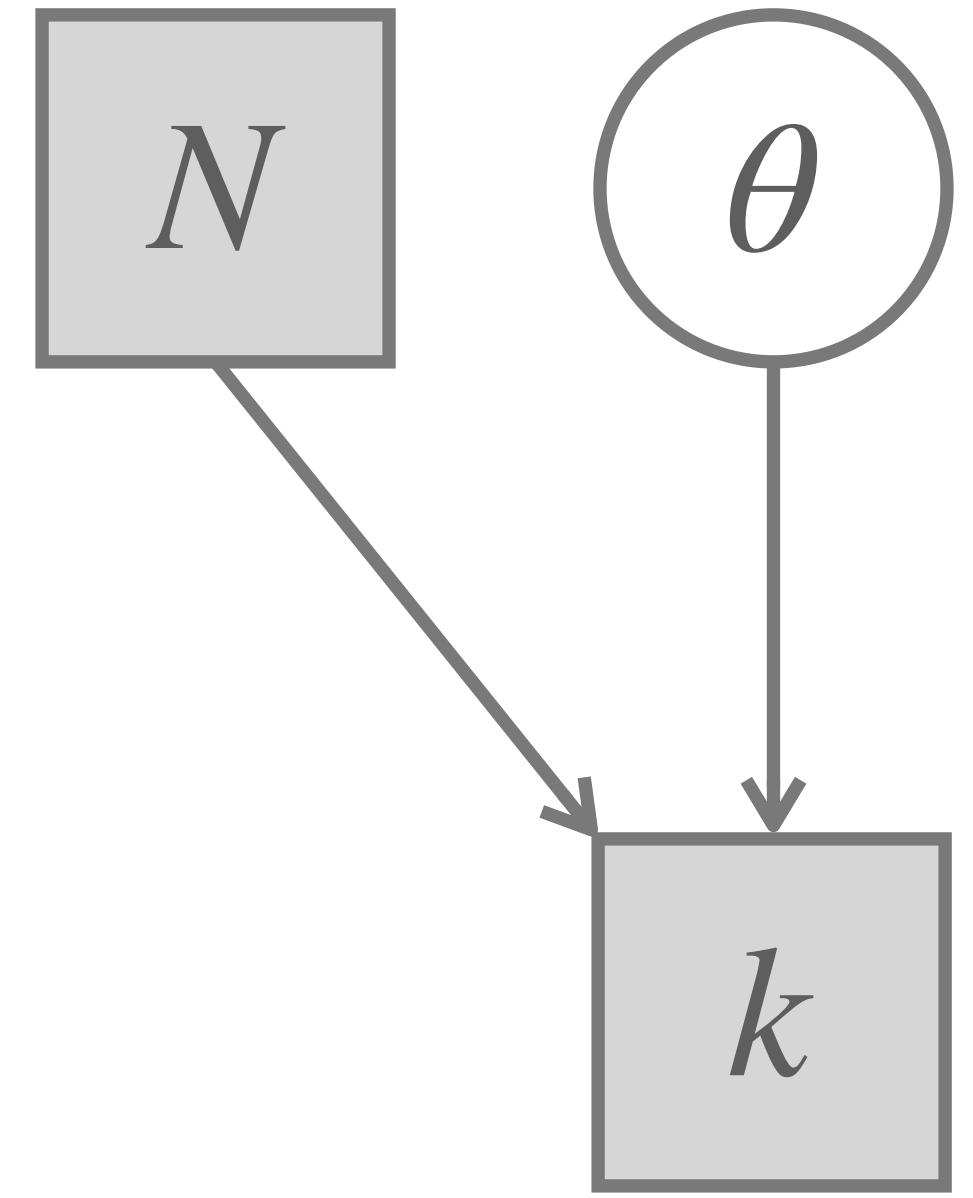
model of the data-generating process buried inside a two-sample t-test

read more [here](#)

Binomial model

the 'coin-flip' model

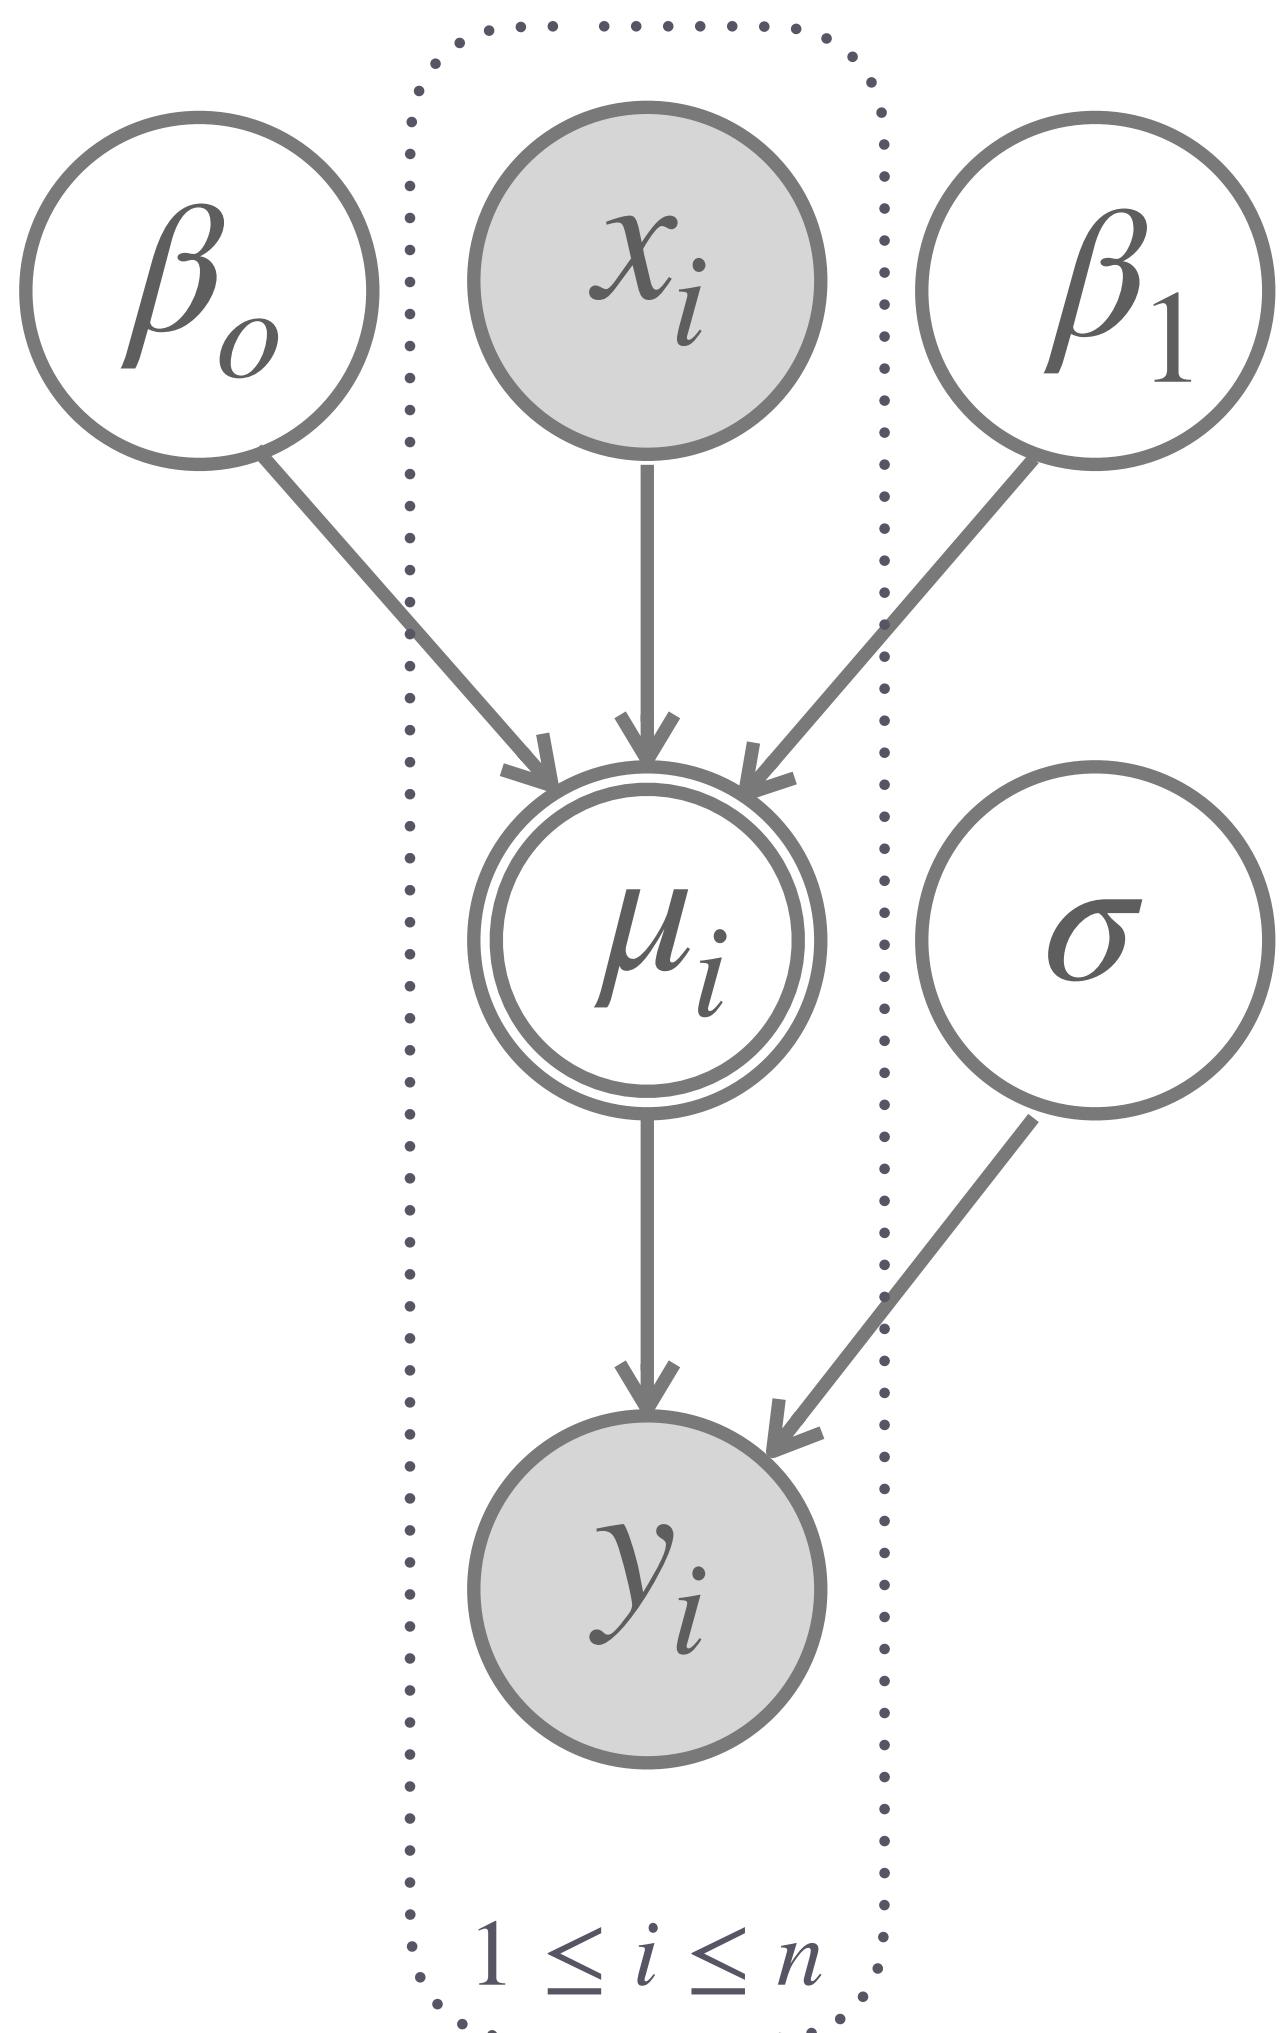
- ▶ data: pair of numbers $D = \{k, N\}$
 - N is the number of tosses
 - k is the number of heads (successes)
- ▶ variable:
 - θ is the number of heads (successes)
- ▶ uninformed prior:
$$\theta \sim \text{Beta}(1,1)$$
- ▶ likelihood function:
$$k \sim \text{Binomial}(\theta, N)$$
- ▶ conventions for model graphs:
 - circles / squares: continuous / discrete variables
 - white / gray nodes: latent / observed variables



Simple linear regression model

for a single predictor variable

- ▶ data: n pairs of numbers $D = \{\langle x_1, y_1 \rangle, \dots \langle x_n, y_n \rangle\}$
 - x_i is the i -th observation of the **independent / predictor variable**
 - y_i is the i -th observation of the **dependent / response variable**
- ▶ parameters:
 - β_0 is the **intercept** parameter
 - β_1 is the **slope** parameter
 - σ is the standard deviation of a normal distribution
- ▶ derived variable: [shown in node w/ double lines]
 - μ_i is the linear predictor for observation i
- ▶ priors (uninformed):
$$\beta_0, \beta_1 \sim \text{Uniform}(-\infty, \infty) \quad \log(\sigma^2) \sim \text{Uniform}(-\infty, \infty)$$
- ▶ likelihood:
$$y_i \sim \text{Normal}(\mu_i, \sigma) \quad \mu_i = \beta_0 + x_1 \cdot \beta_1$$





Simple linear regression

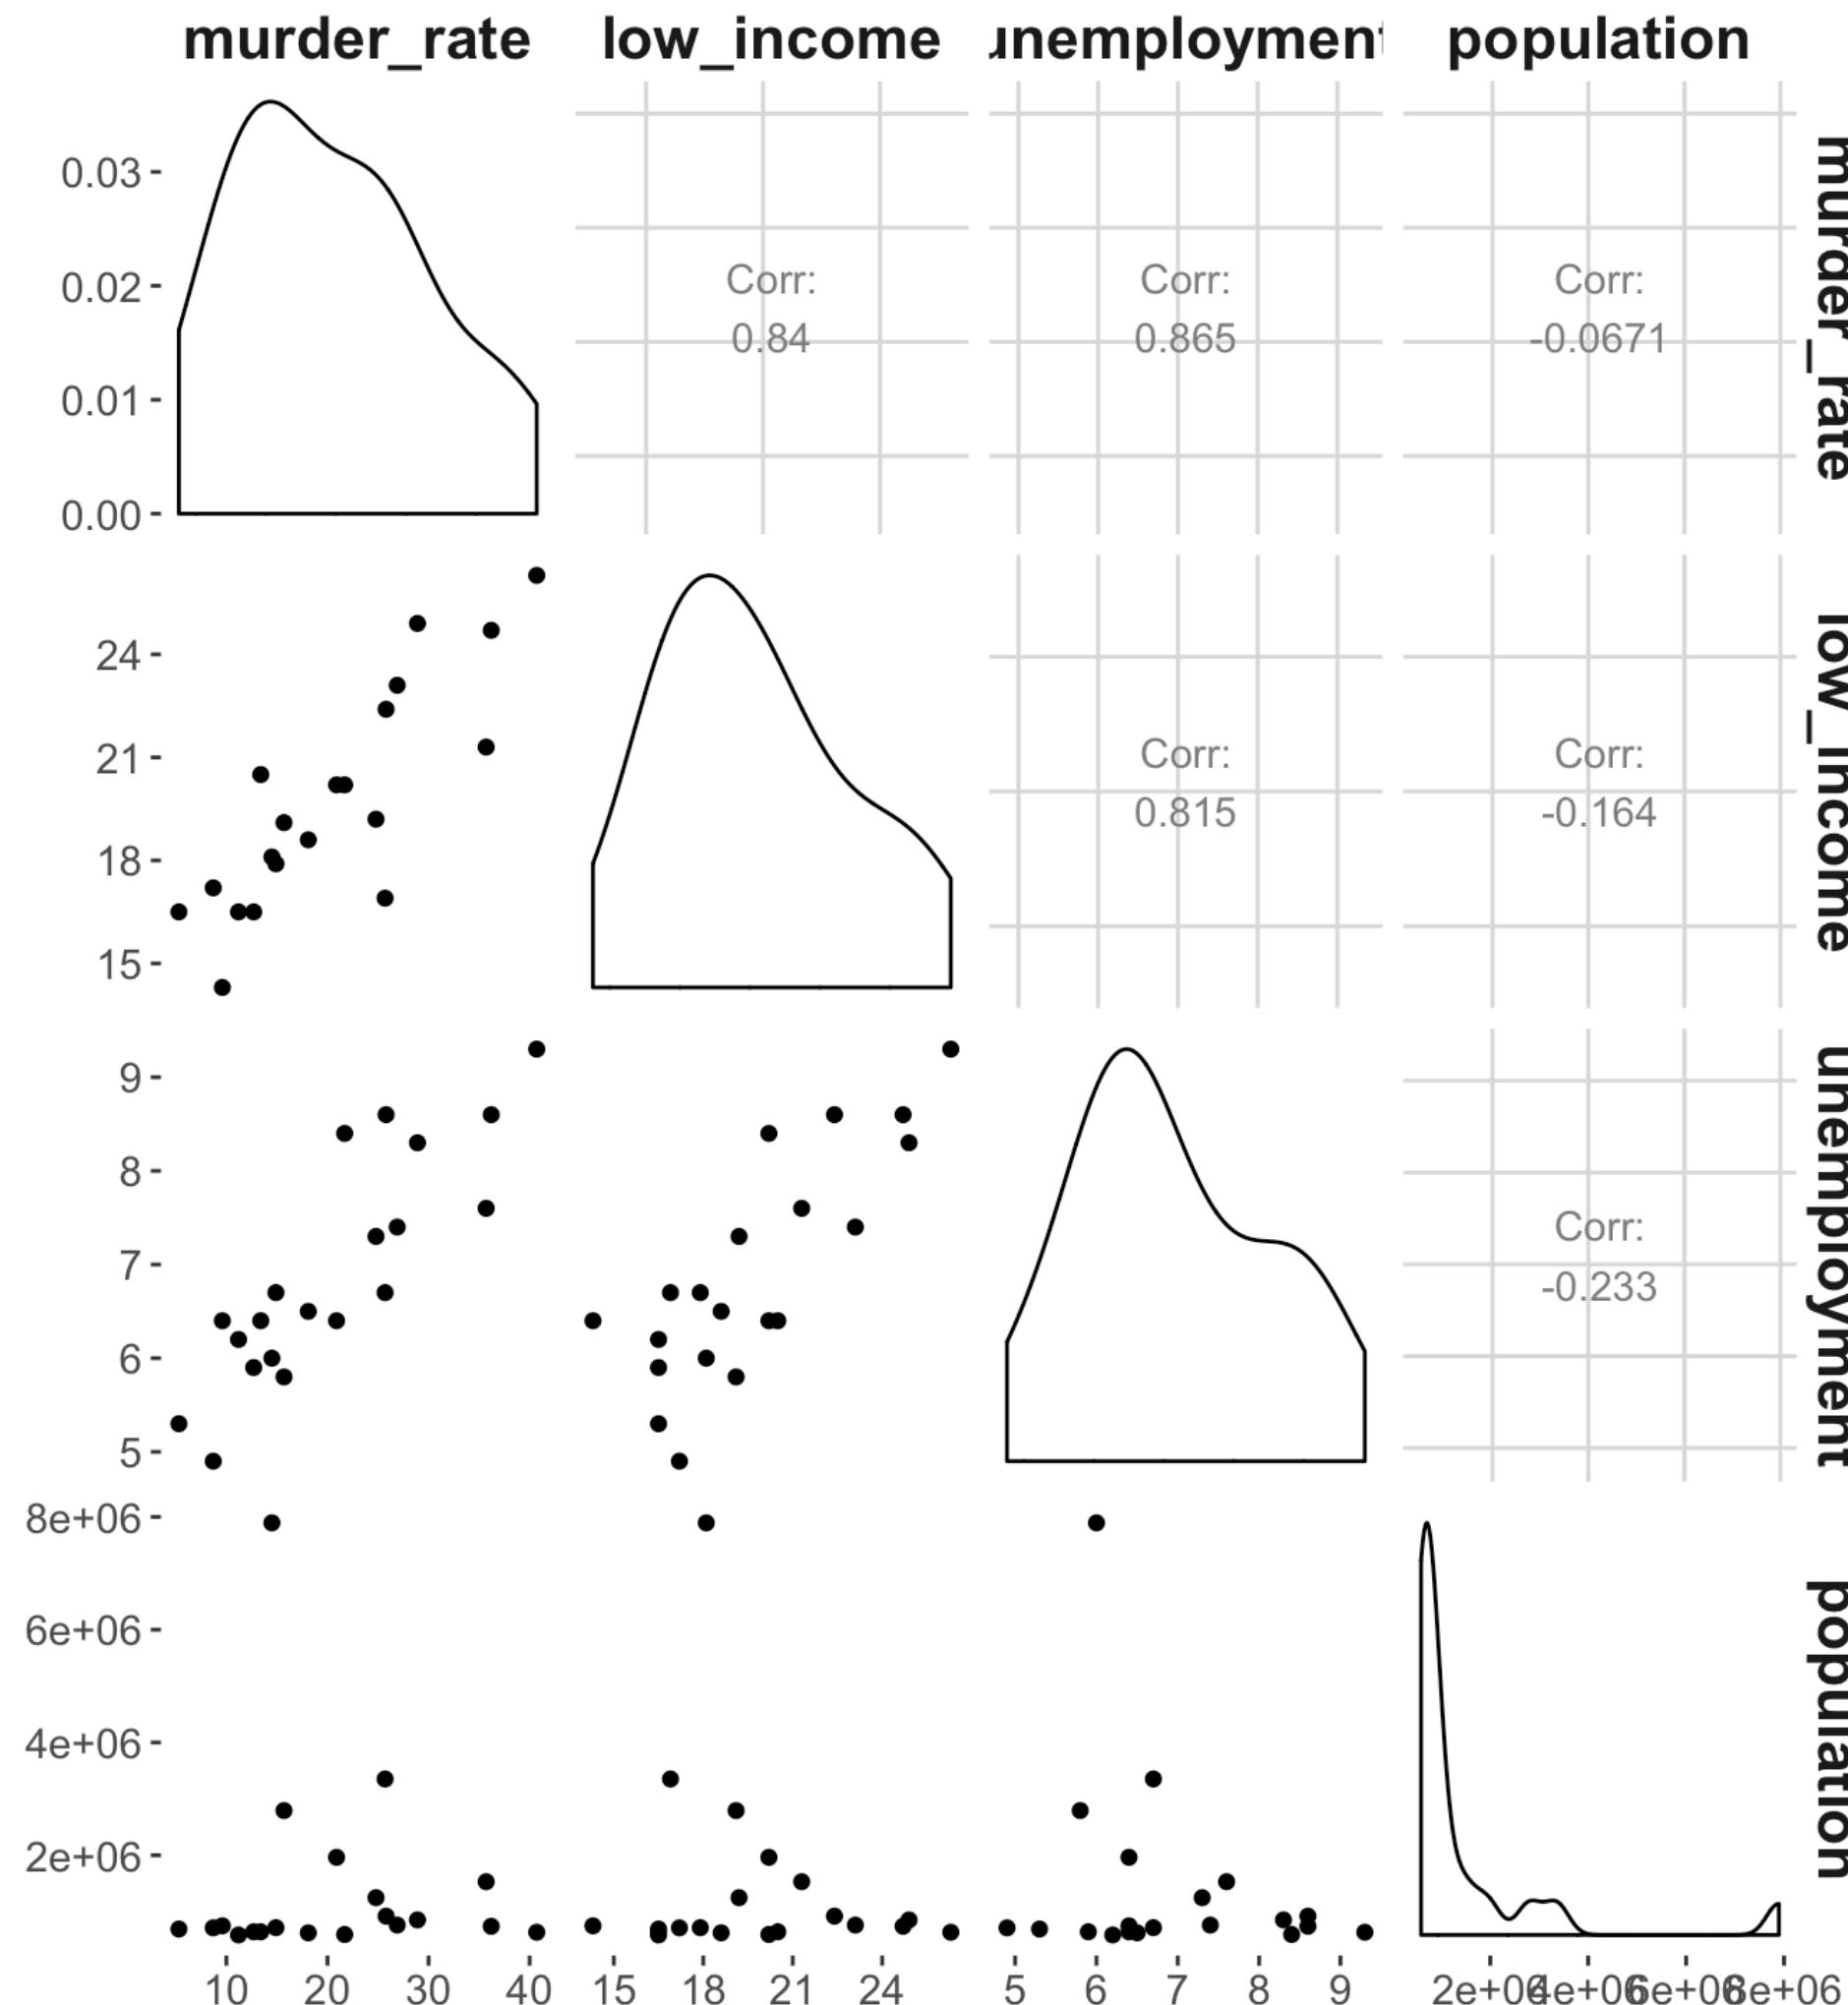
likelihood & Bayesian posterior

Murder data

annual murder rate, average income, unemployment rates and population

```
## # A tibble: 20 x 4
##   murder_rate low_income unemployment population
##       <dbl>      <dbl>        <dbl>     <dbl>
## 1       11.2      16.5        6.2    587000
## 2       13.4      20.5        6.4    643000
## 3       40.7      26.3        9.3   635000
## 4        5.3      16.5        5.3    692000
## 5       24.8      19.2        7.3  1248000
## 6       12.7      16.5        5.9    643000
## 7       20.9      20.2        6.4  1964000
## 8       35.7      21.3        7.6  1531000
## 9        8.7      17.2        4.9    713000
## 10      9.6      14.3        6.4    749000
## 11      14.5      18.1        6     7895000
## 12      26.9      23.1        7.4    762000
## 13      15.7      19.1        5.8   2793000
## 14      36.2      24.7        8.6    741000
## 15      18.1      18.6        6.5   625000
## 16      28.9      24.9        8.3   854000
## 17      14.9      17.9        6.7   716000
## 18      25.8      22.4        8.6   921000
## 19      21.7      20.2        8.4   595000
## 20      25.7      16.9        6.7  3353000
```

Murder rate data



annual murders per
million inhabitants

percentage inhabitants
with low income

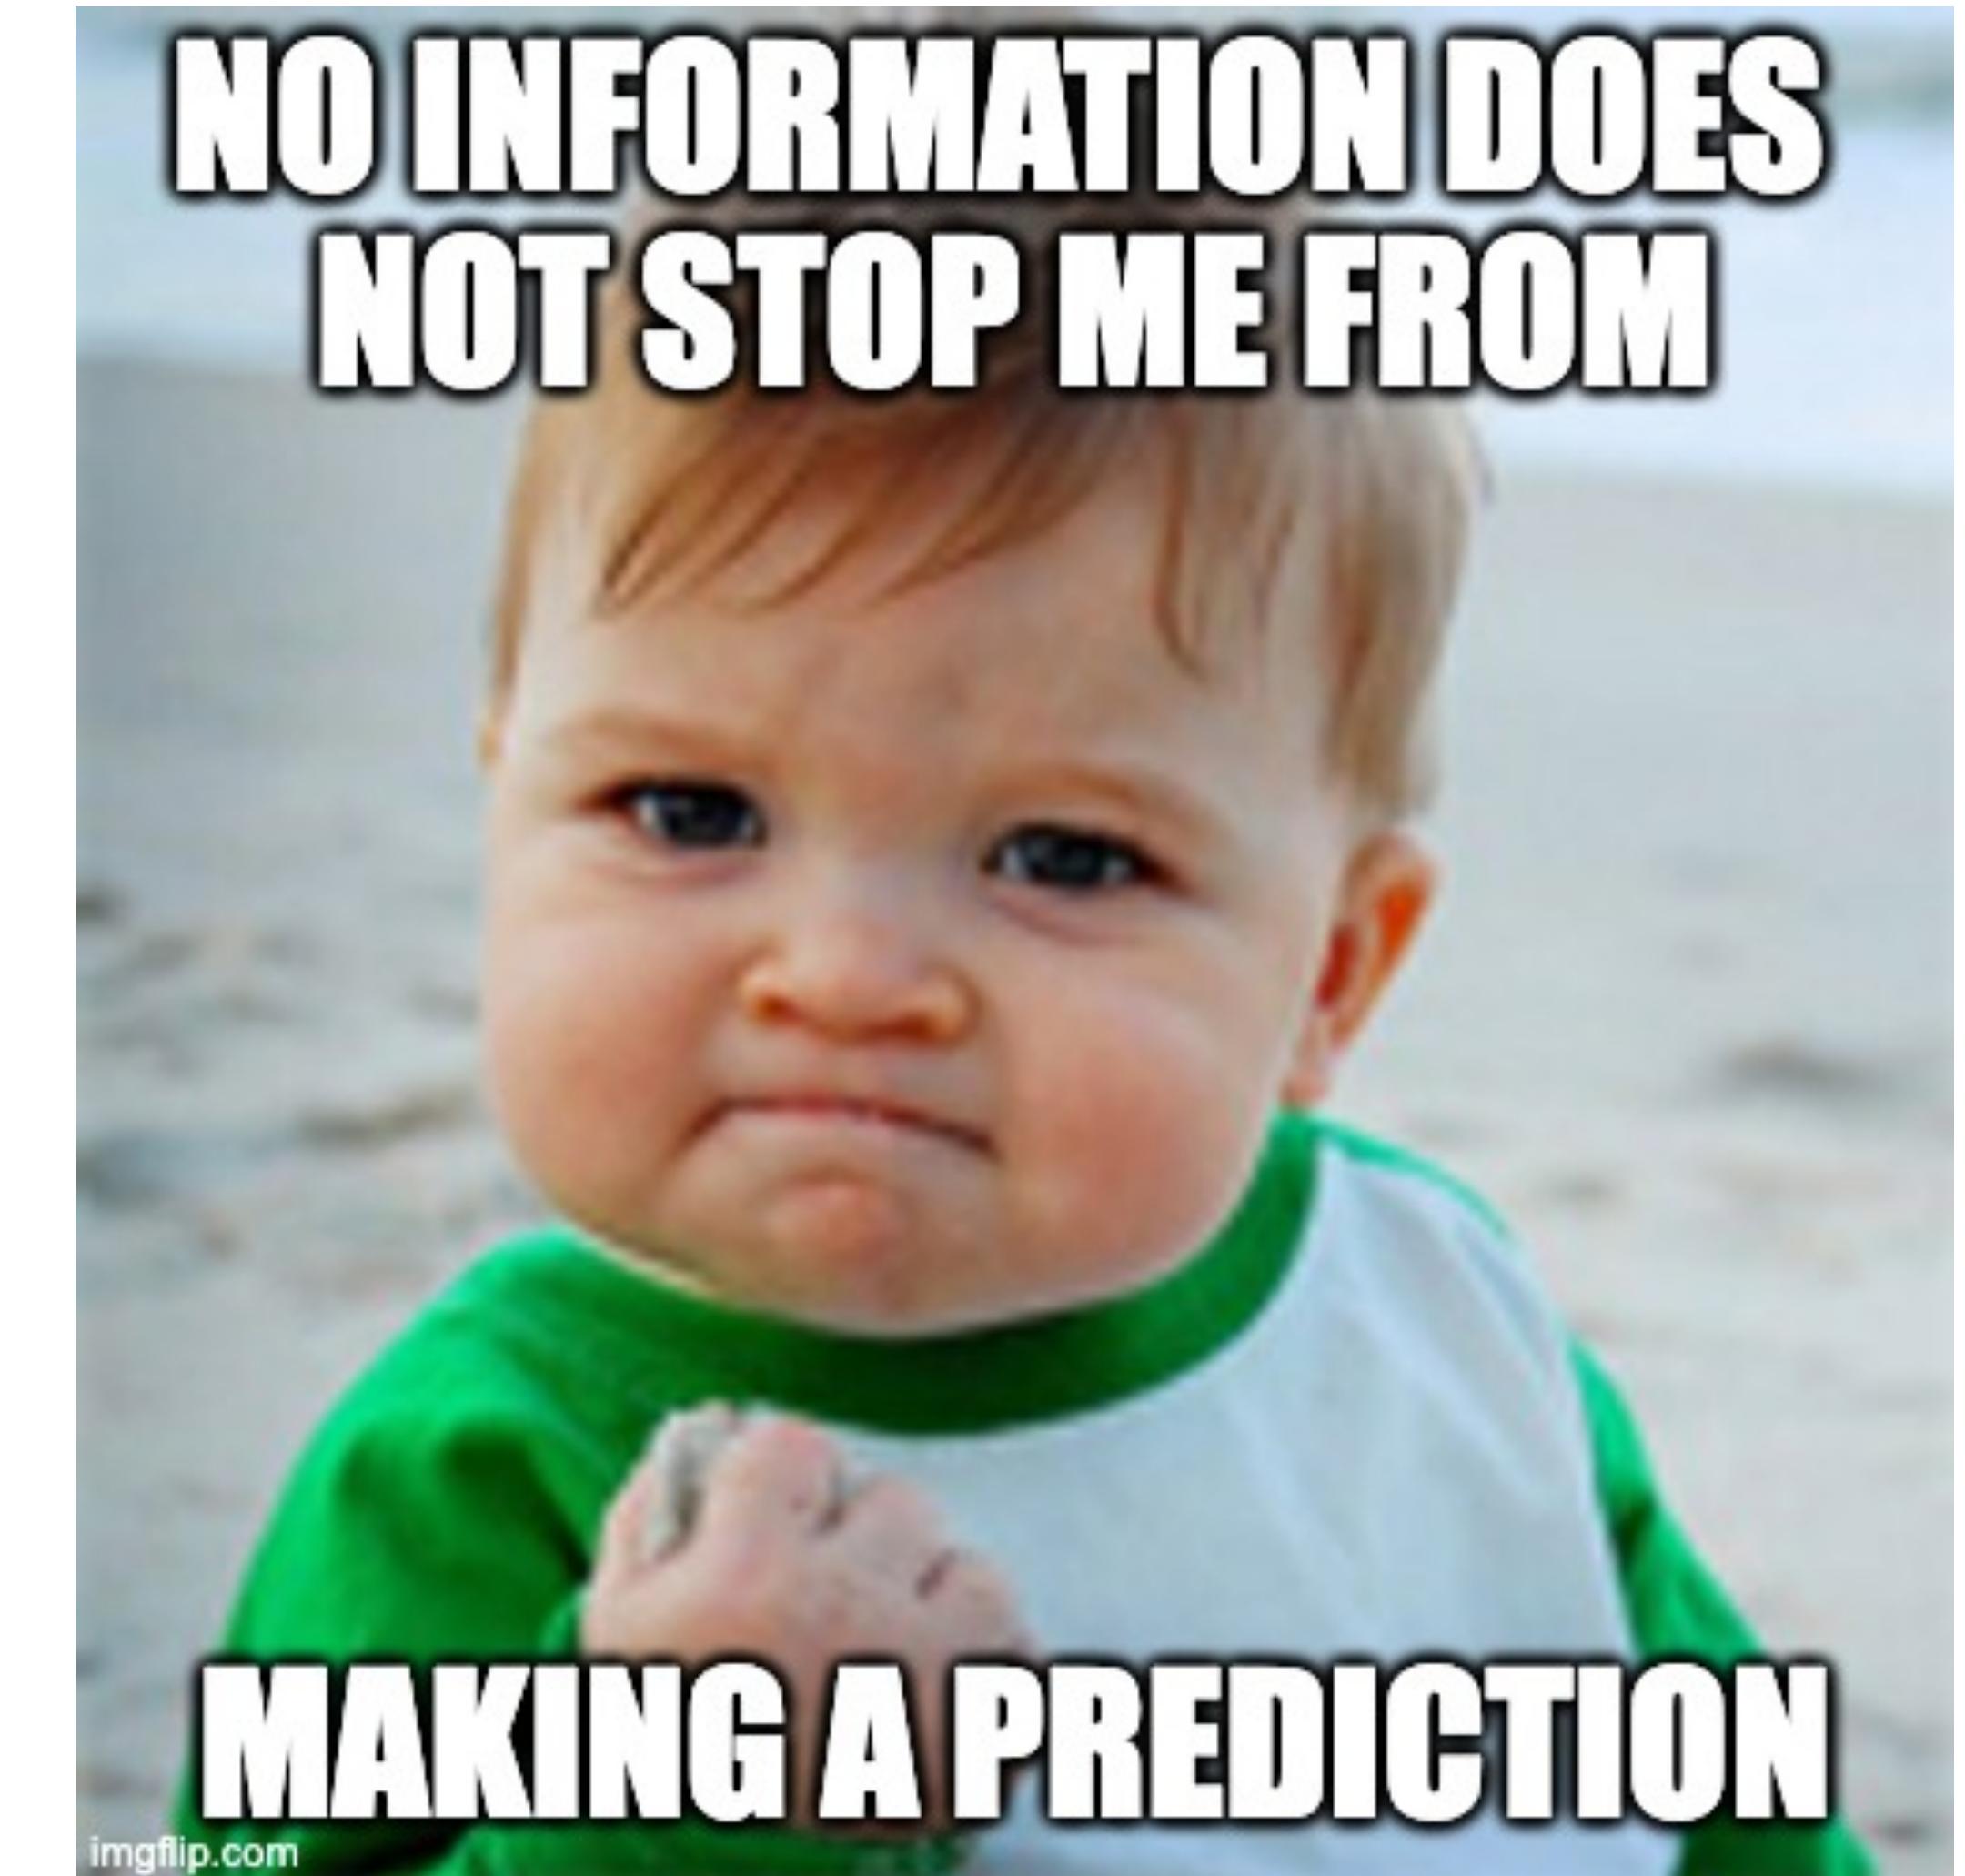
percentage inhabitants
who are unemployed

total population

Predicting murder rate

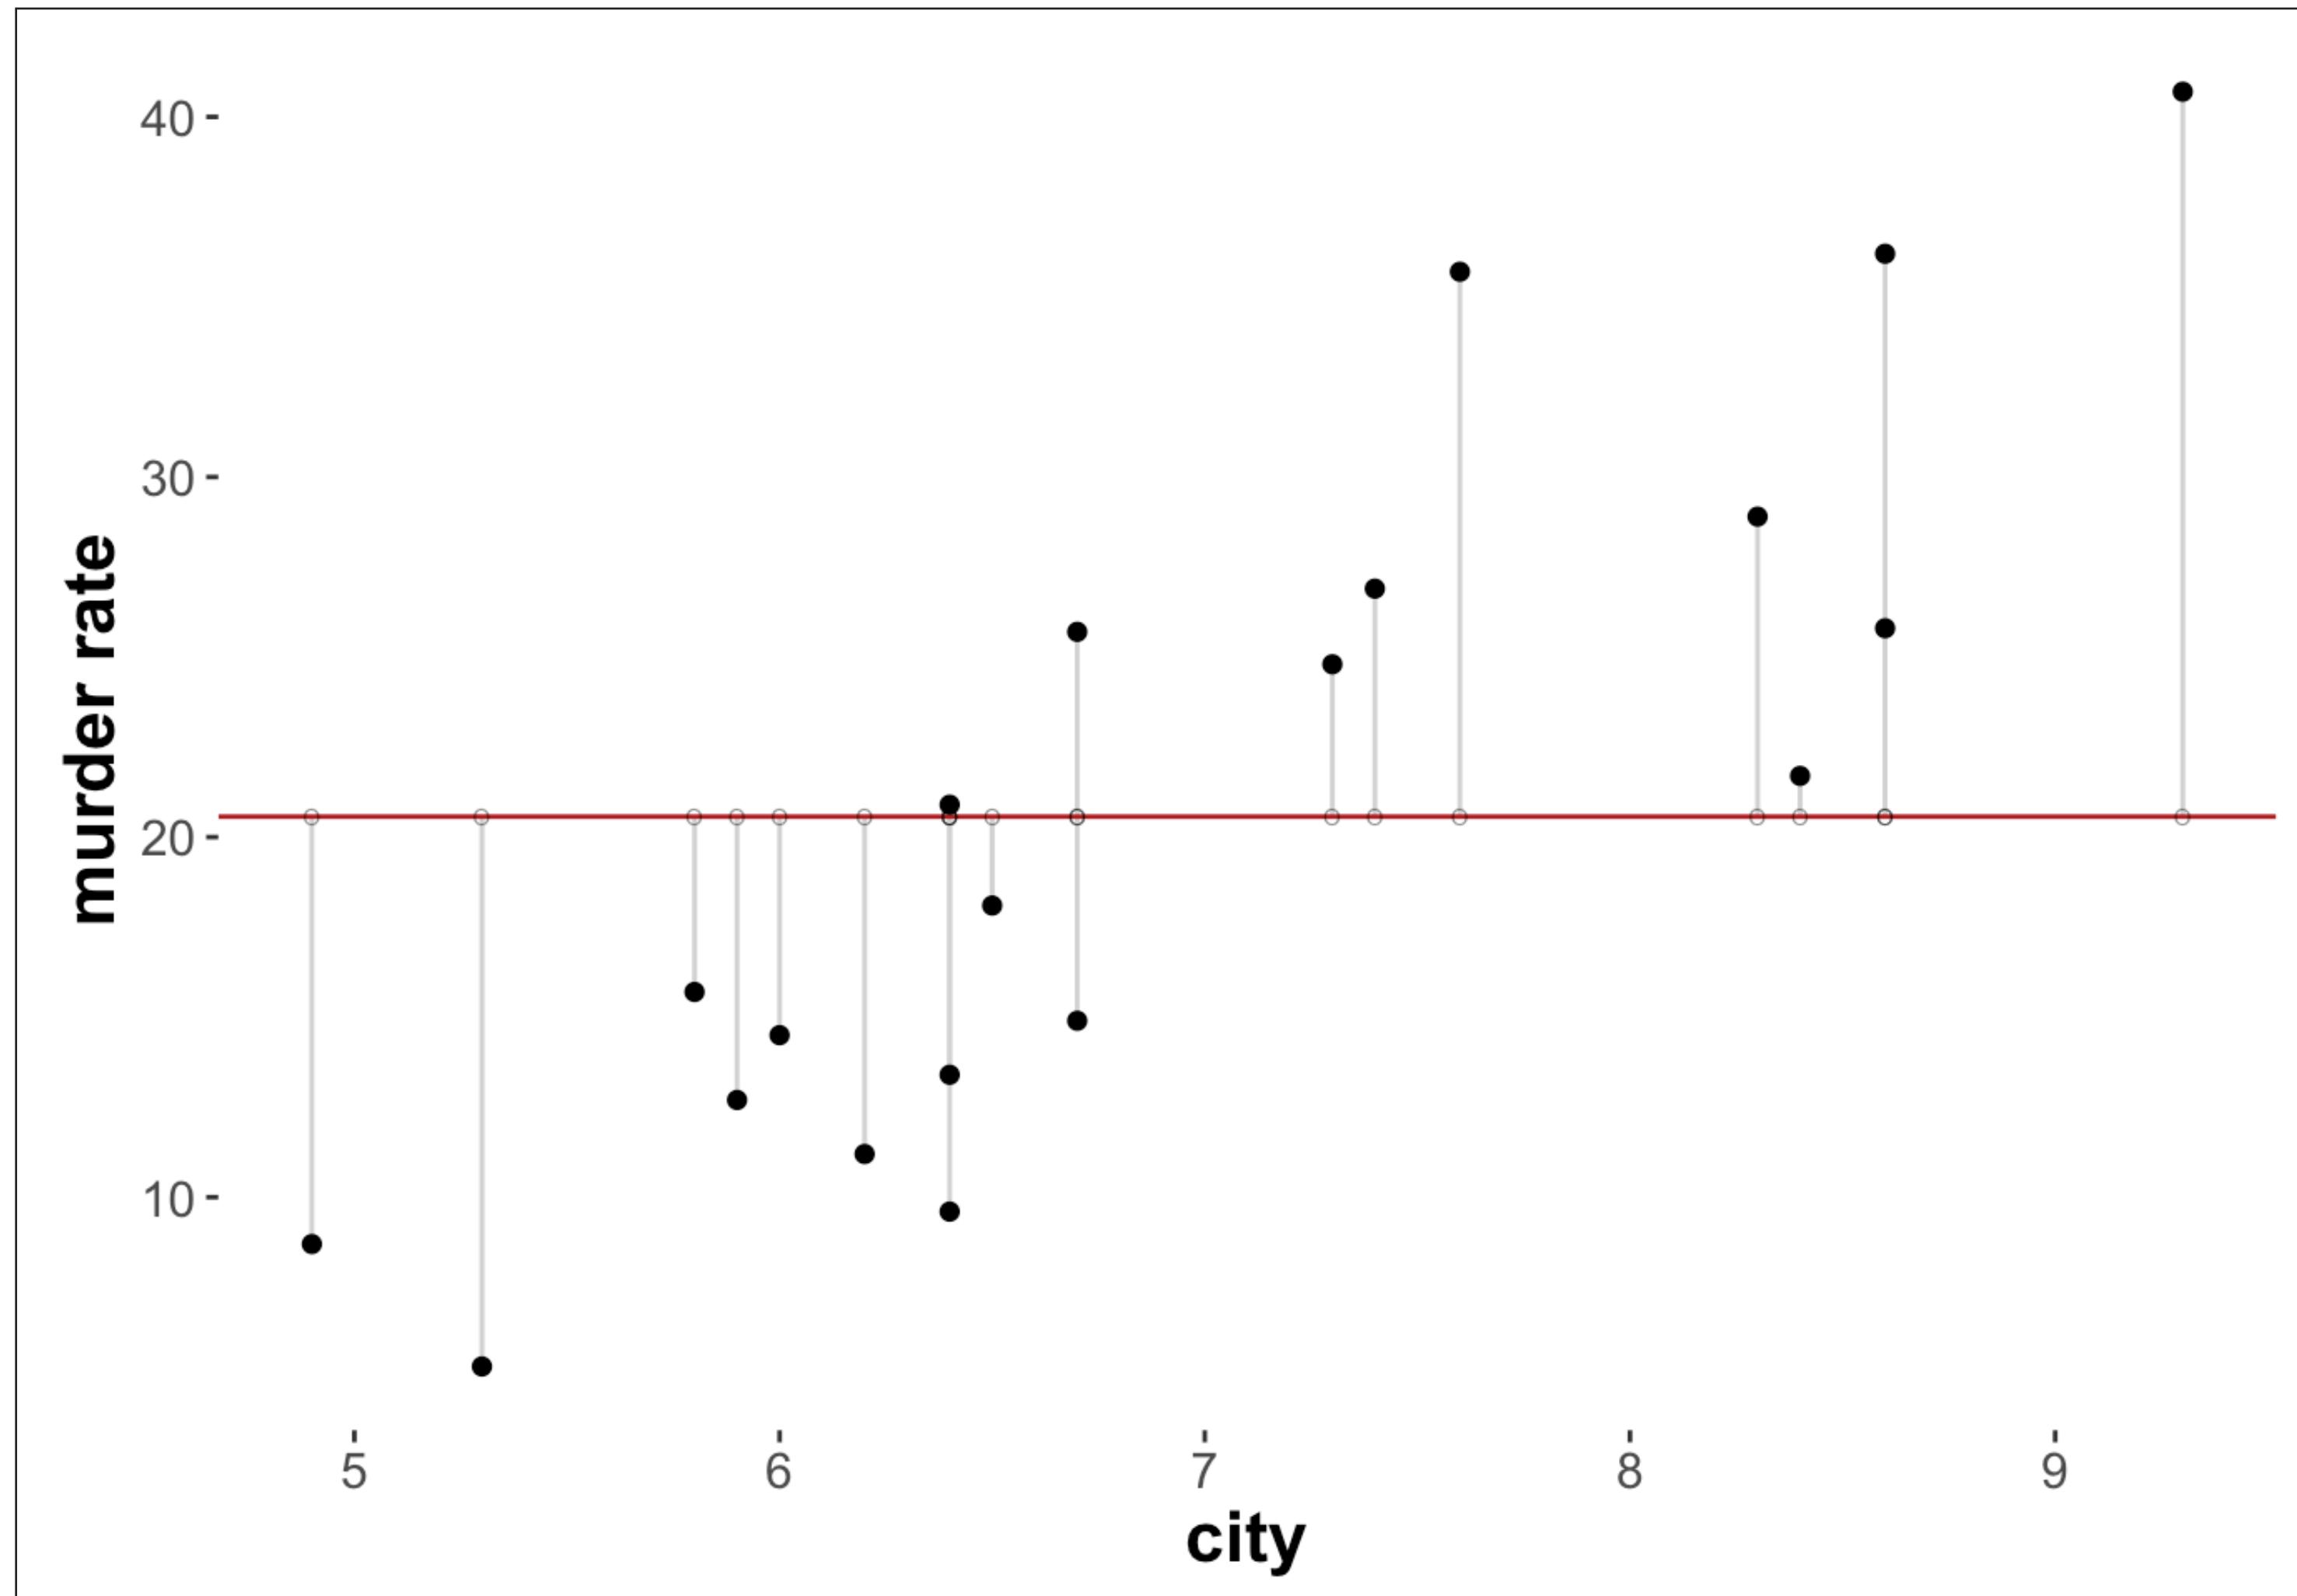
no information at all

```
## # A tibble: 20 x 4
##   murder_rate low_income unemployment population
##       <dbl>      <dbl>        <dbl>      <dbl>
## 1       11.2      16.5        6.2     587000
## 2       13.4      20.5        6.4     643000
## 3       40.7      26.3        9.3     635000
## 4        5.3      16.5        5.3     692000
## 5       24.8      19.2        7.3    1248000
## 6       12.7      16.5        5.9     643000
## 7       20.9      20.2        6.4    1964000
## 8       35.7      21.3        7.6   1531000
## 9        8.7      17.2        4.9     713000
## 10      9.6      14.3        6.4     749000
## 11      14.5      18.1        6     7895000
## 12      26.9      23.1        7.4     762000
## 13      15.7      19.1        5.8    2793000
## 14      36.2      24.7        8.6     741000
## 15      18.1      18.6        6.5     625000
## 16      28.9      24.9        8.3     854000
## 17      14.9      17.9        6.7     716000
## 18      25.8      22.4        8.6     921000
## 19      21.7      20.2        8.4     595000
## 20      25.7      16.9        6.7    3353000
```



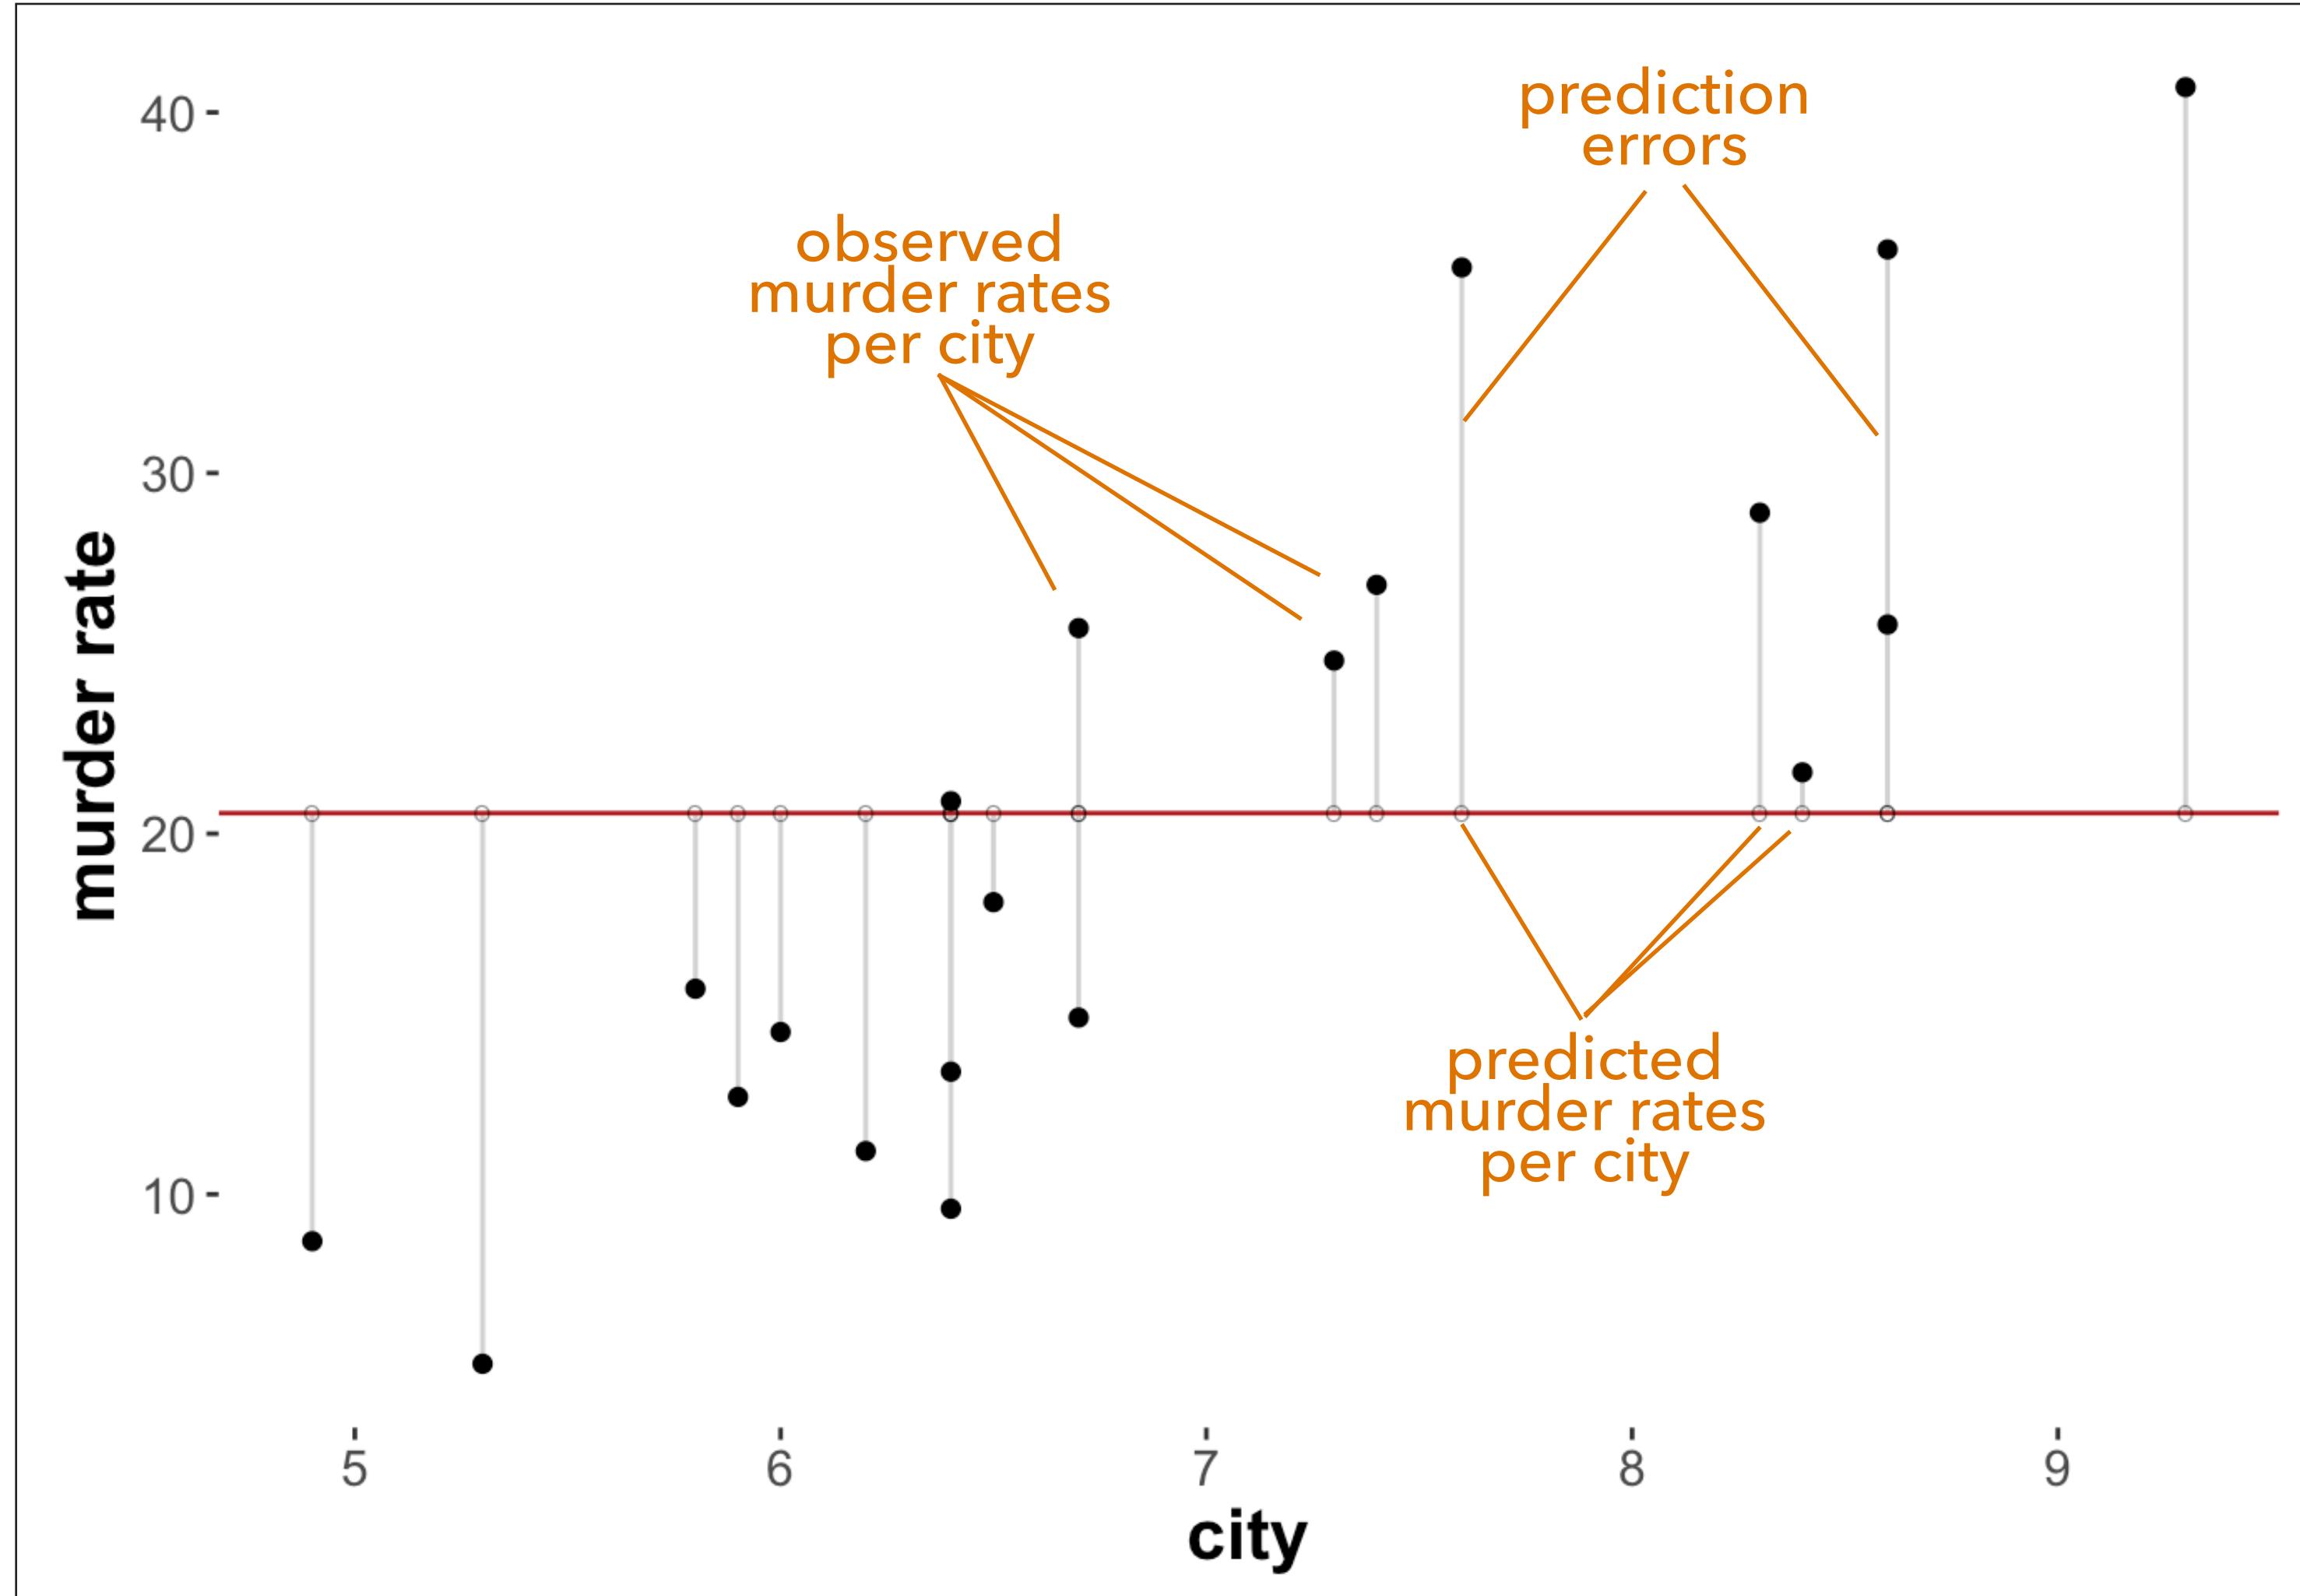
Predicting murder rate

by empirical mean



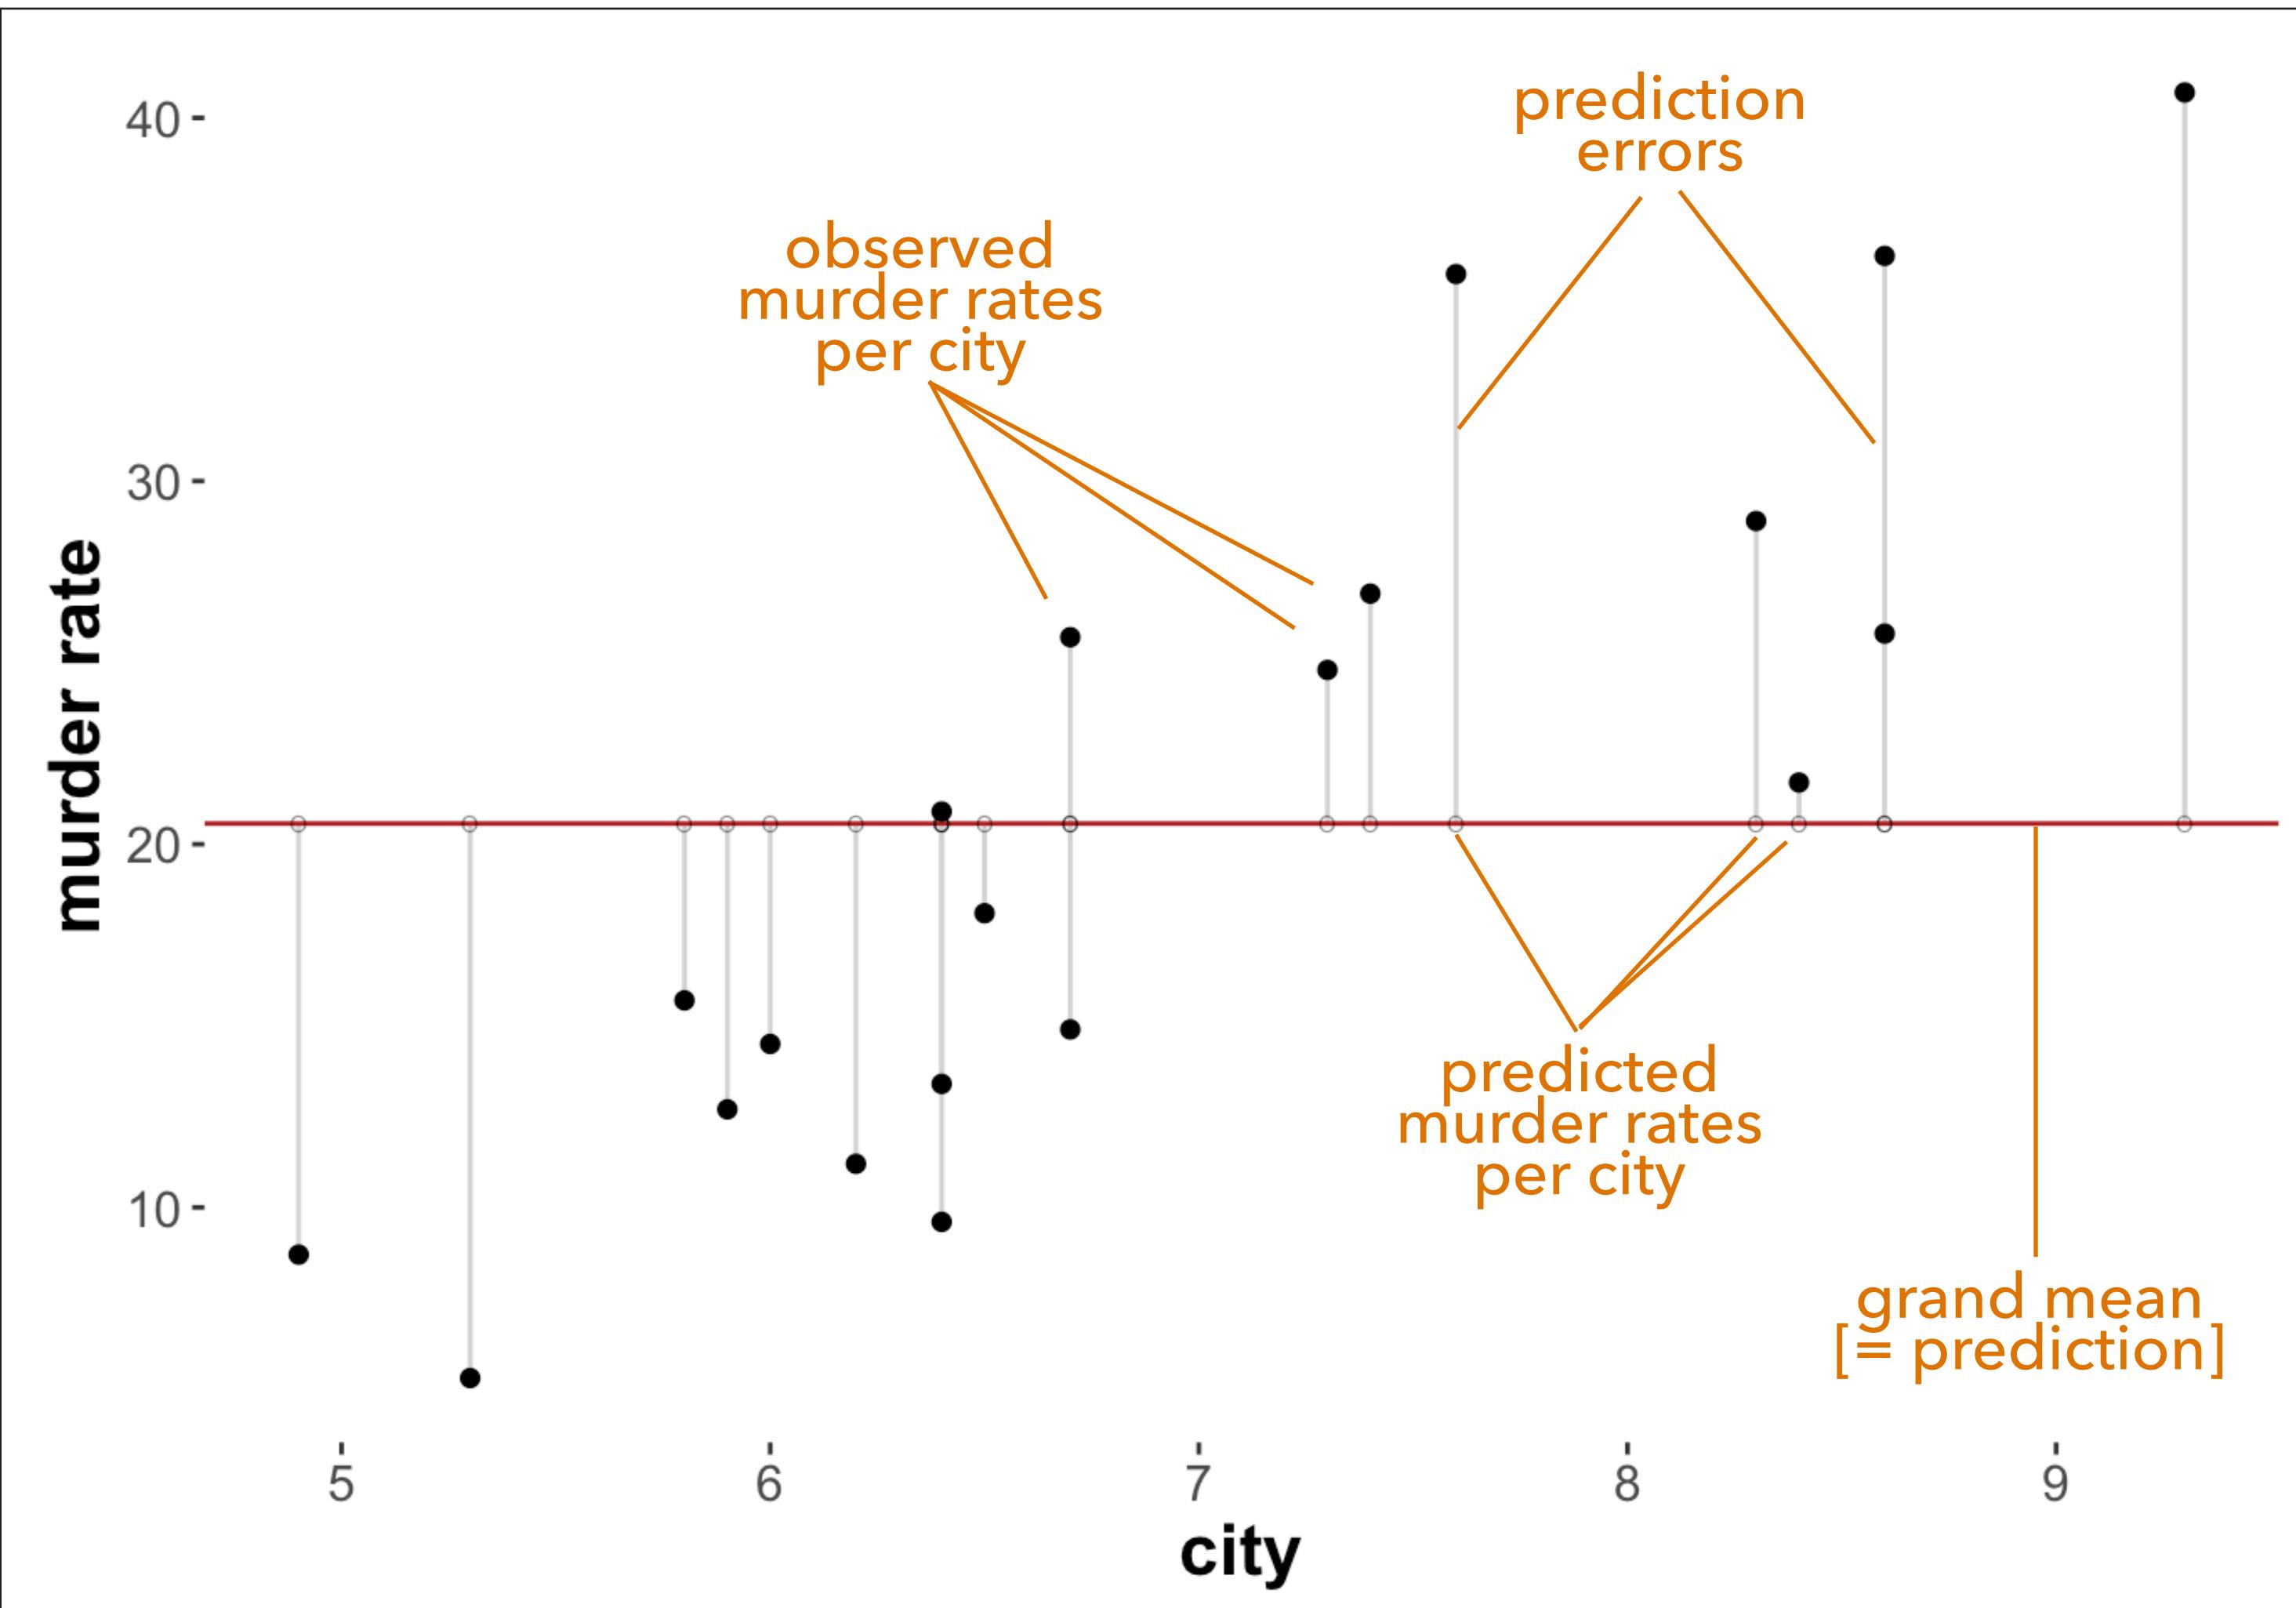
Predicting murder rate

by grand mean



Predicting murder rate

by grand mean



$$TSS = \sum_{i=1}^n (y_i - \bar{y})^2$$

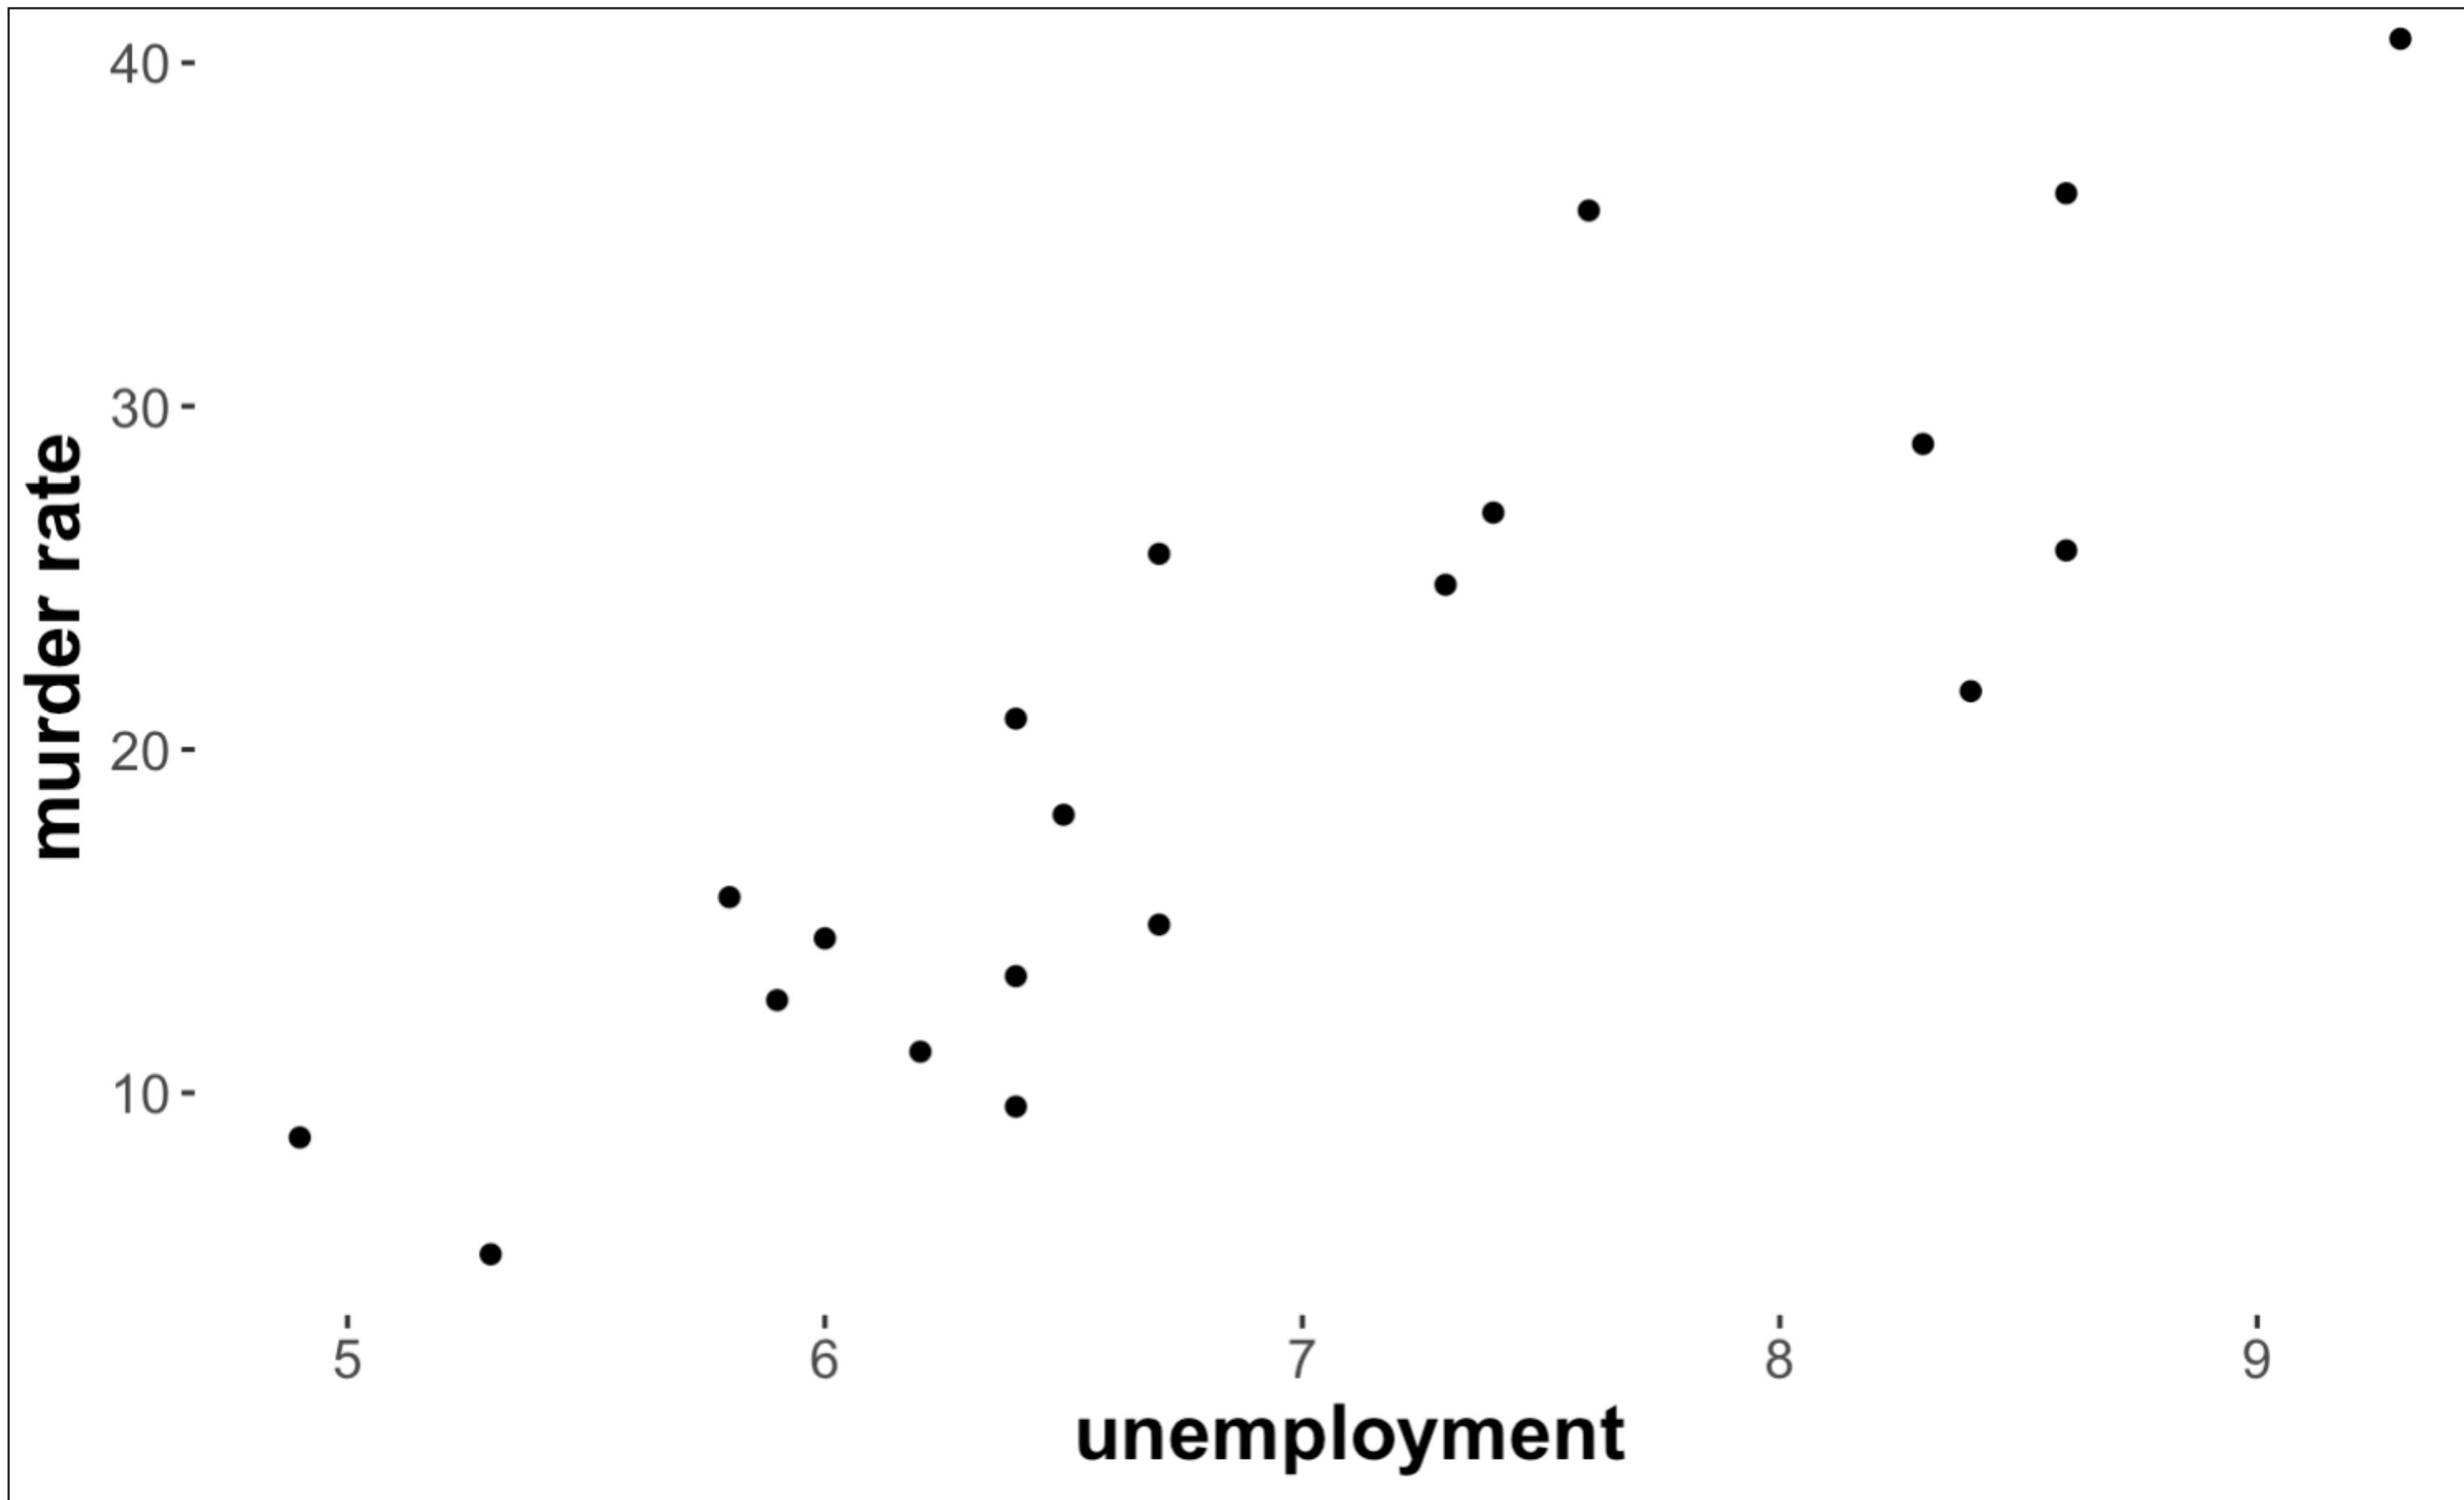
[total sum of squares]

```
y <- murder_data %>% pull(murder_rate)
n <- length(y)
tss_simple <- sum((y - mean(y))^2)
tss_simple
```

```
## [1] 1855.202
```

Predicting murder rate based on unemployment rate

some wild linear guessing



We are to predict the murder rate y_i of a randomly drawn city i . We know that city's unemployment rate, x_i , but nothing more.

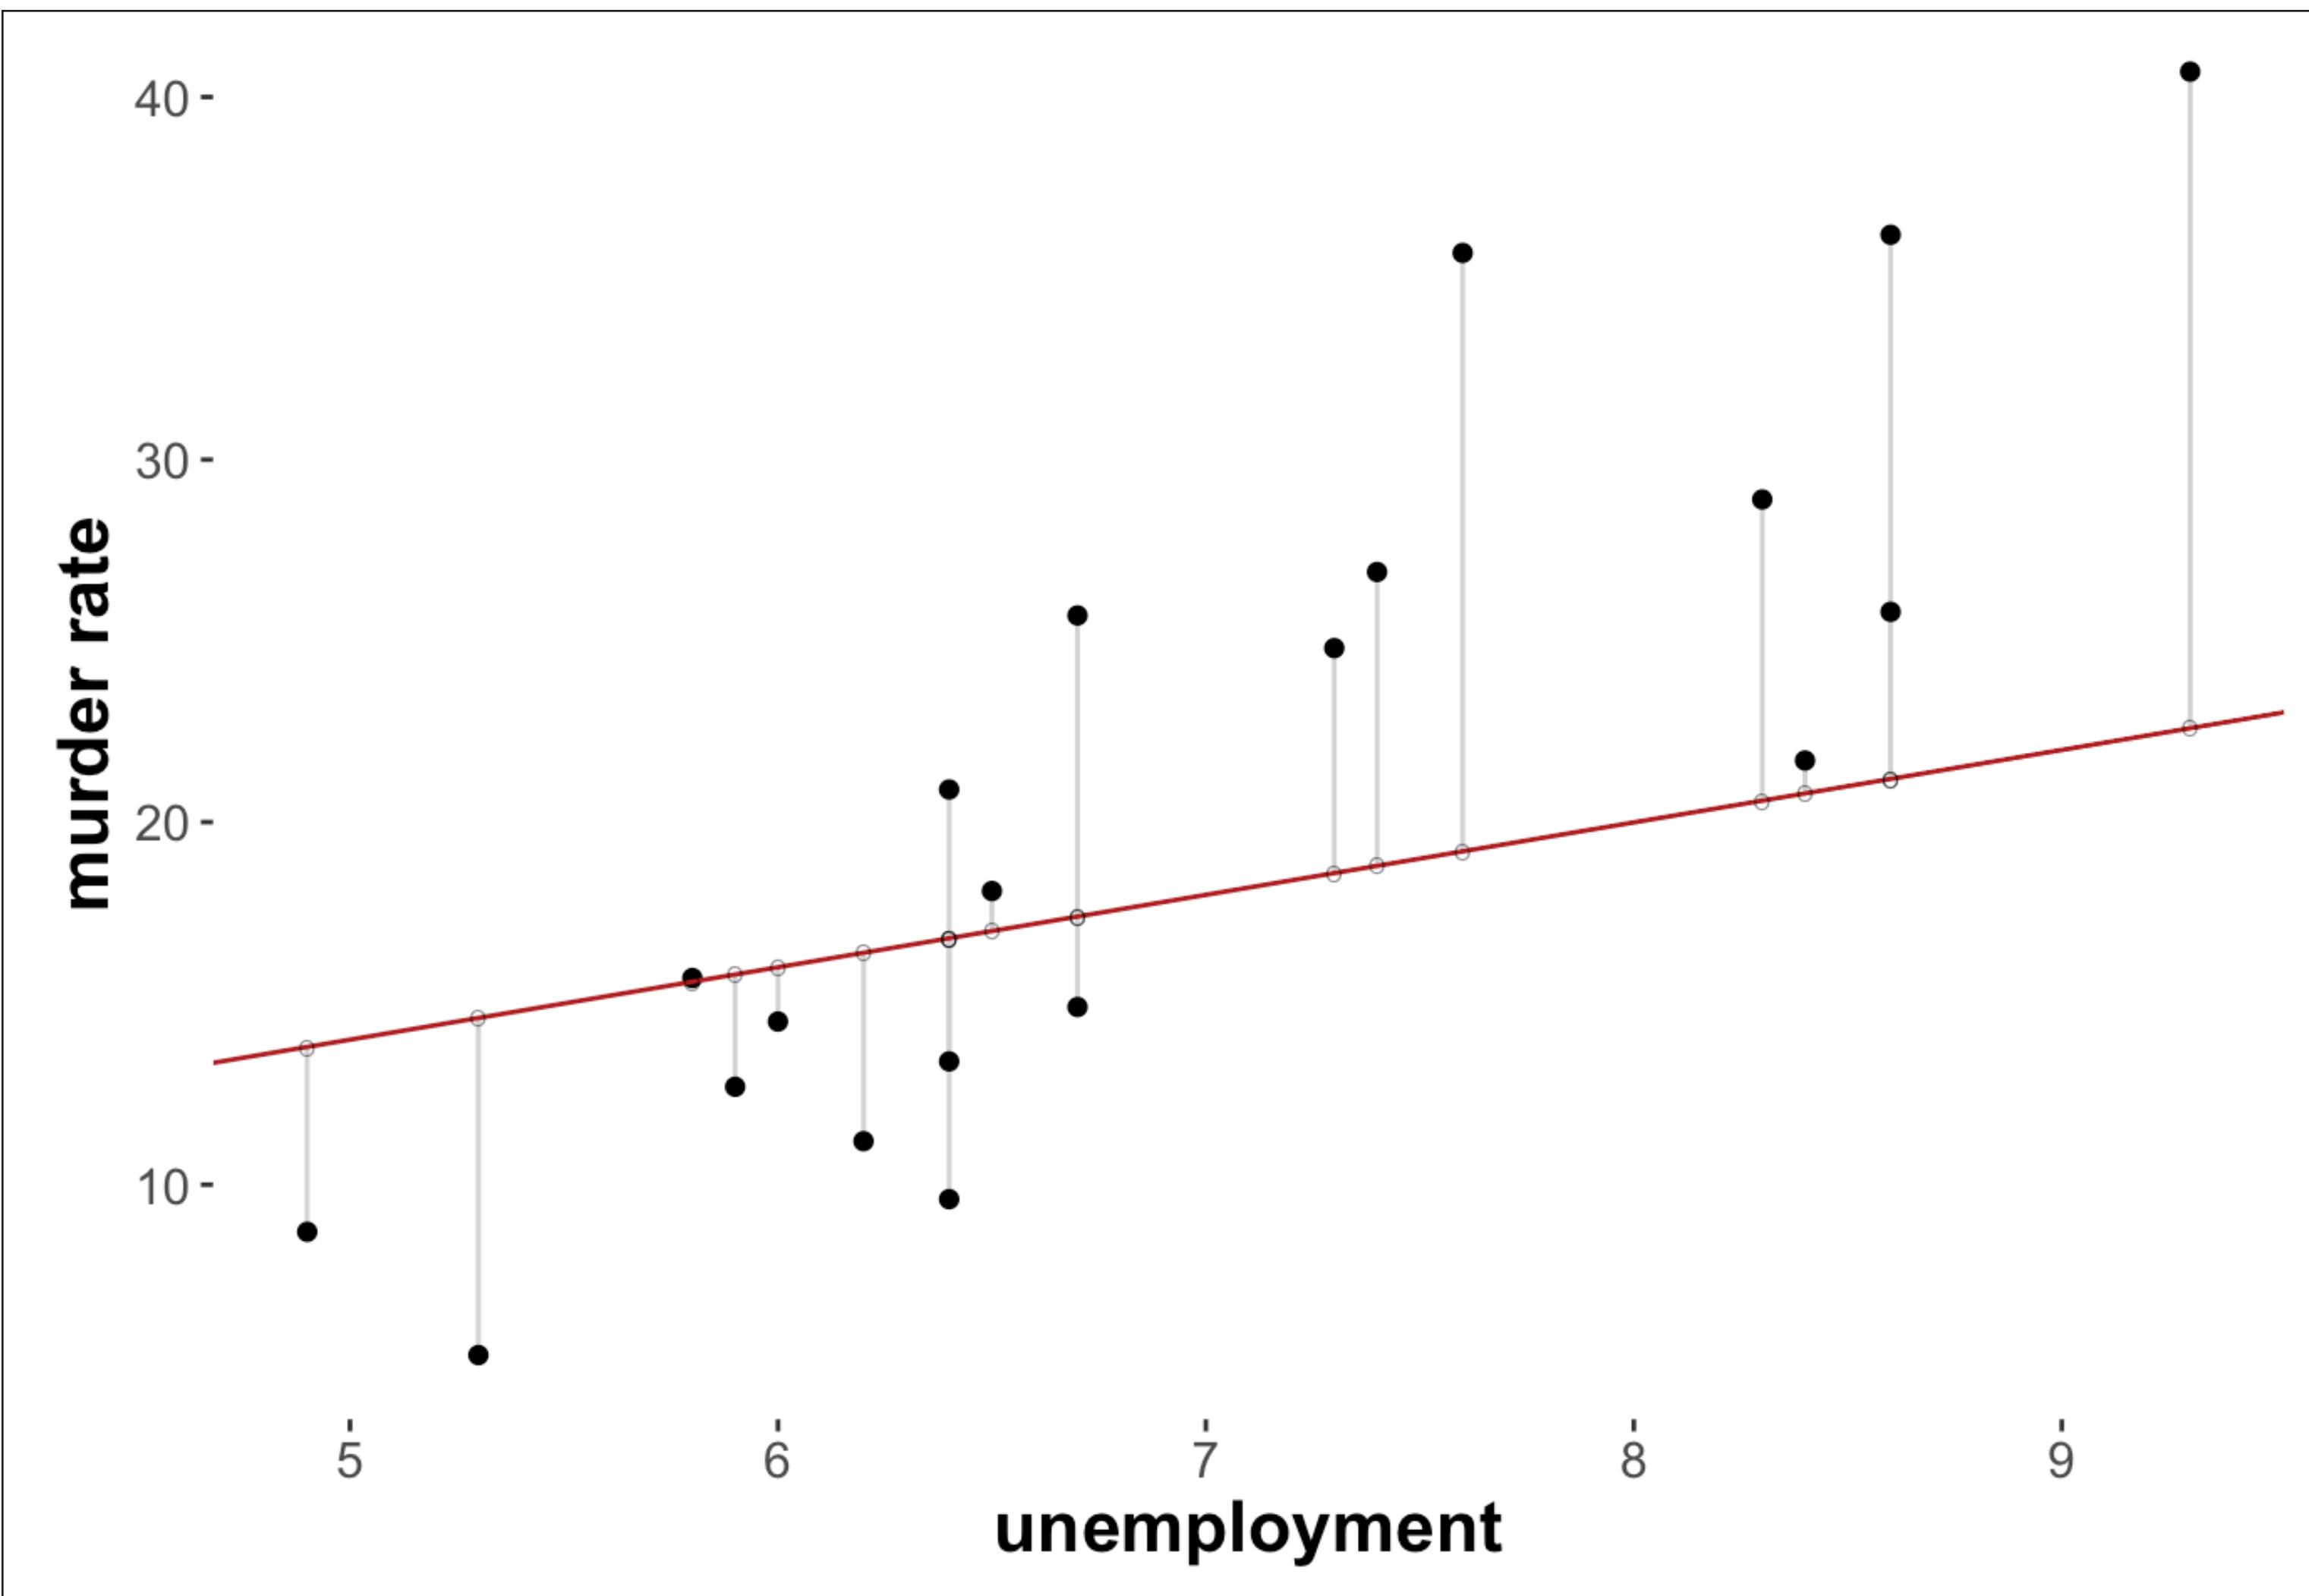
Let's just assume the following linear relationship to make a prediction b/c why not?!?

$$\hat{y}_i = 4 + 2x_i$$

How good is this prediction?

How good is any given prediction?

quantifying distance or likelihood



Distance-based approach

Residual Sum-of-Squares:

$$RSS = \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

- ▶ no predictions about spread around linear predictor

Likelihood-based approach:

Normal likelihood:

$$LH = \prod_{i=1}^n \mathcal{N}(y_i | \mu = \hat{y}_i, \sigma)$$

- ▶ fully predictive

Likelihood-based simple linear regression

- ▶ likelihood:

$$y_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \beta_0 + x_1 \cdot \beta_1$$

- ▶ differential likelihood:

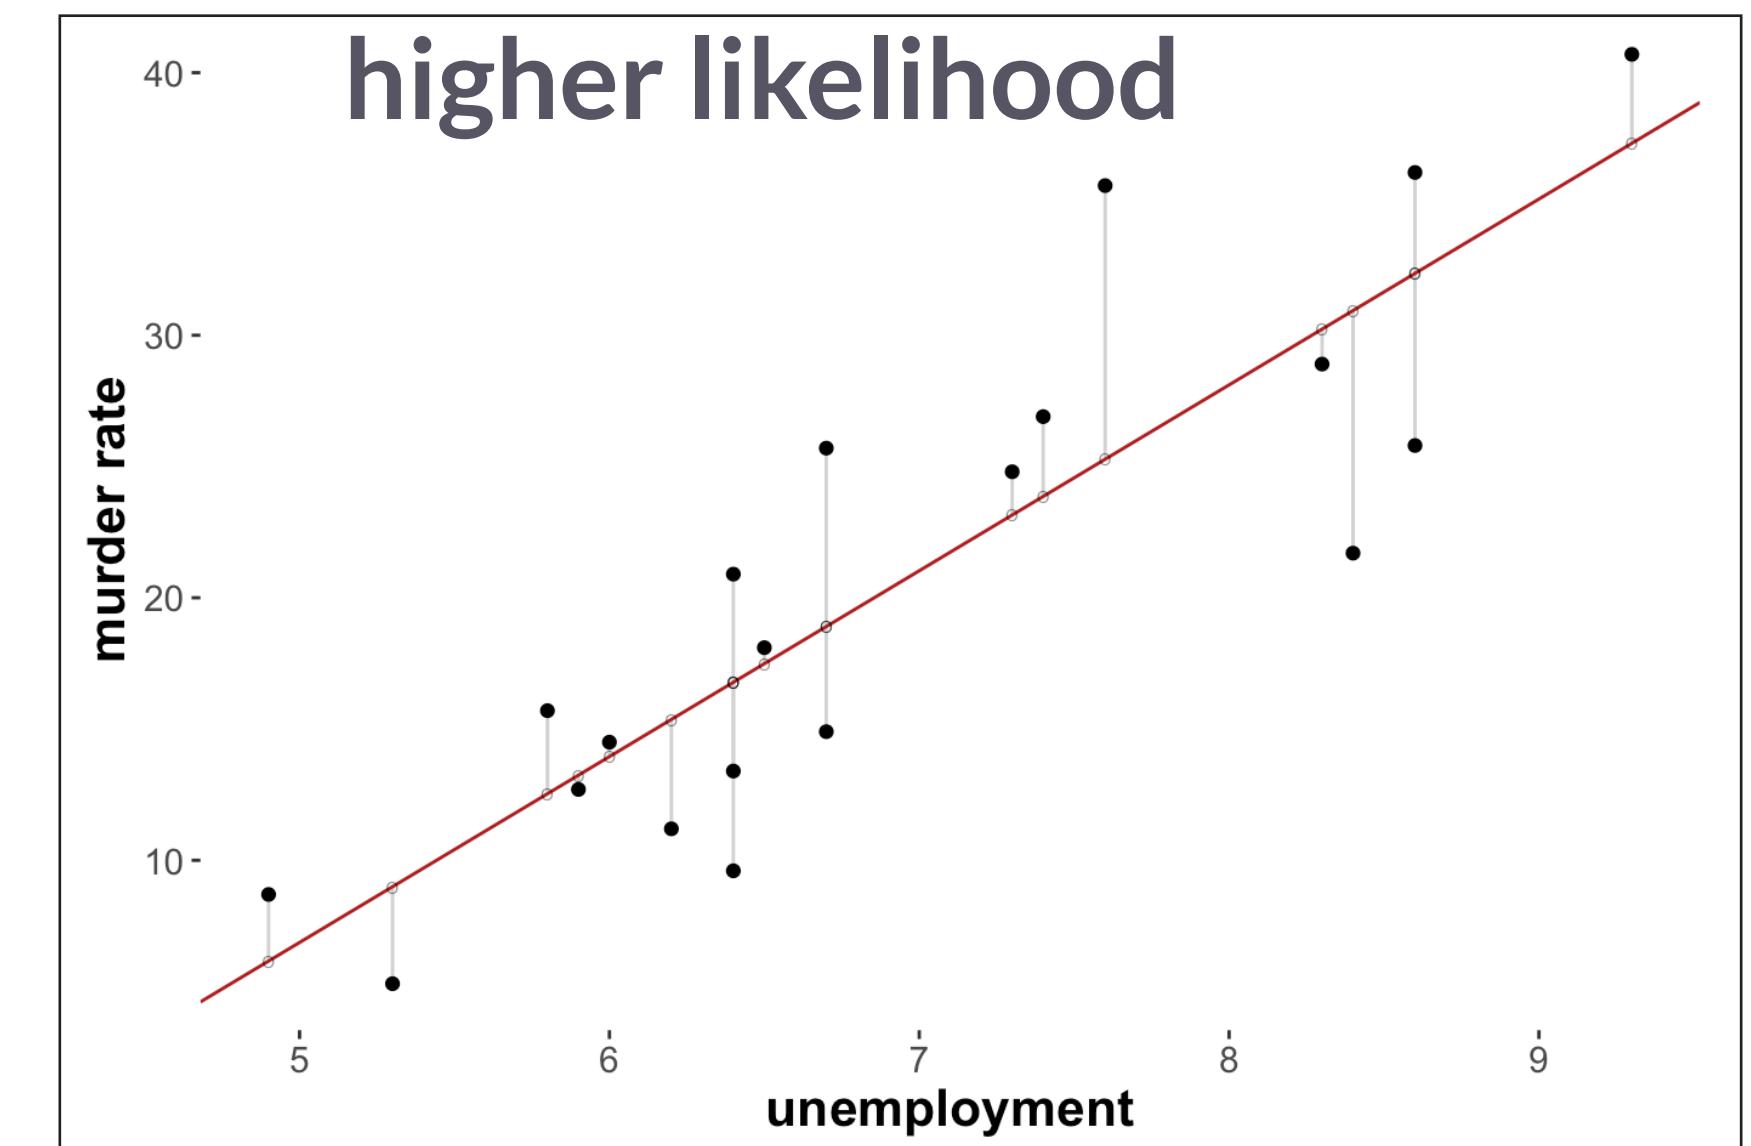
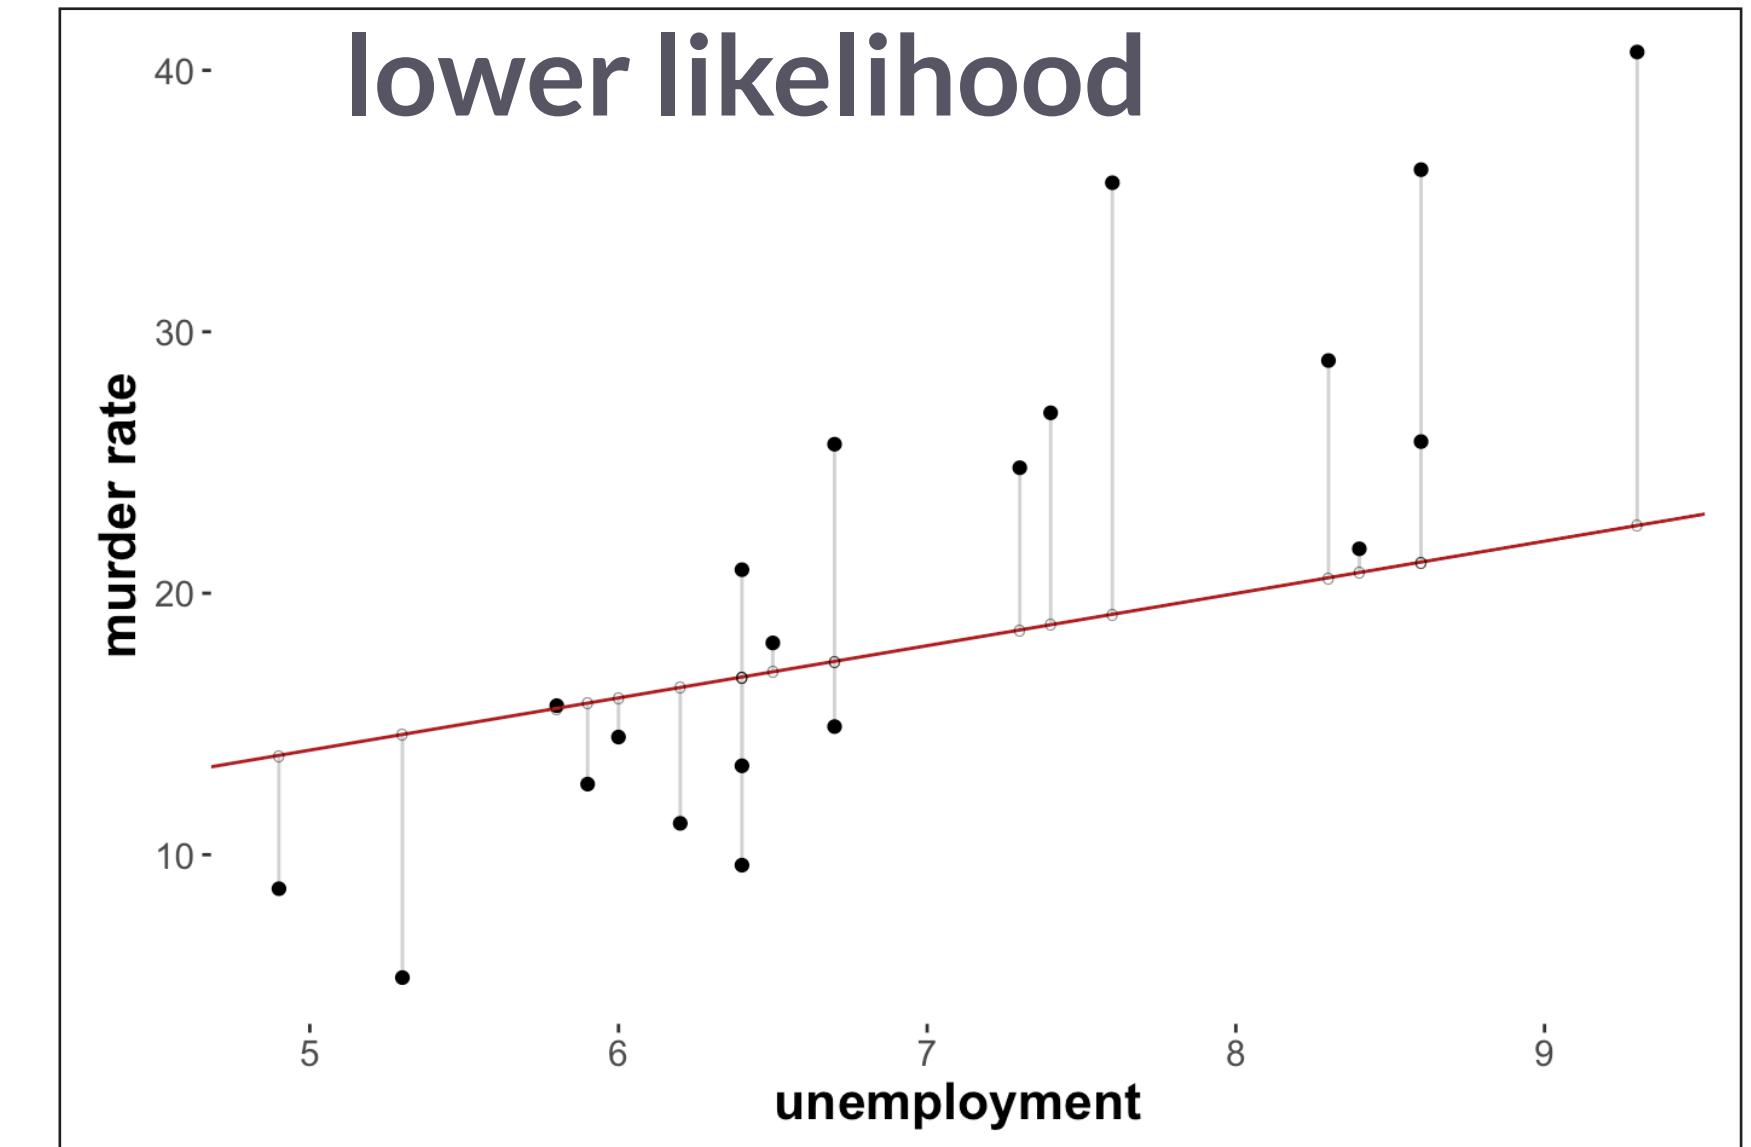
- parameter triples $\langle \beta_0, \beta_1, \sigma \rangle$ can be better or worse
- higher vs. lower likelihood $P(D | \beta_0, \beta_1, \sigma)$

- ▶ maximum-likelihood solution:

$$\arg \max_{\beta_0, \beta_1, \sigma} P(D | \beta_0, \beta_1, \sigma)$$

- standard (frequentist) solution
 - MLE corresponds to MAP for “flat” priors
- ▶ Bayesian approach: full posterior distribution

$$P(\beta_0, \beta_1, \sigma | D) \propto P(D | \beta_0, \beta_1, \sigma) P(\beta_0, \beta_1, \sigma)$$



Bayesian linear regression in R

using BRMS and Stan

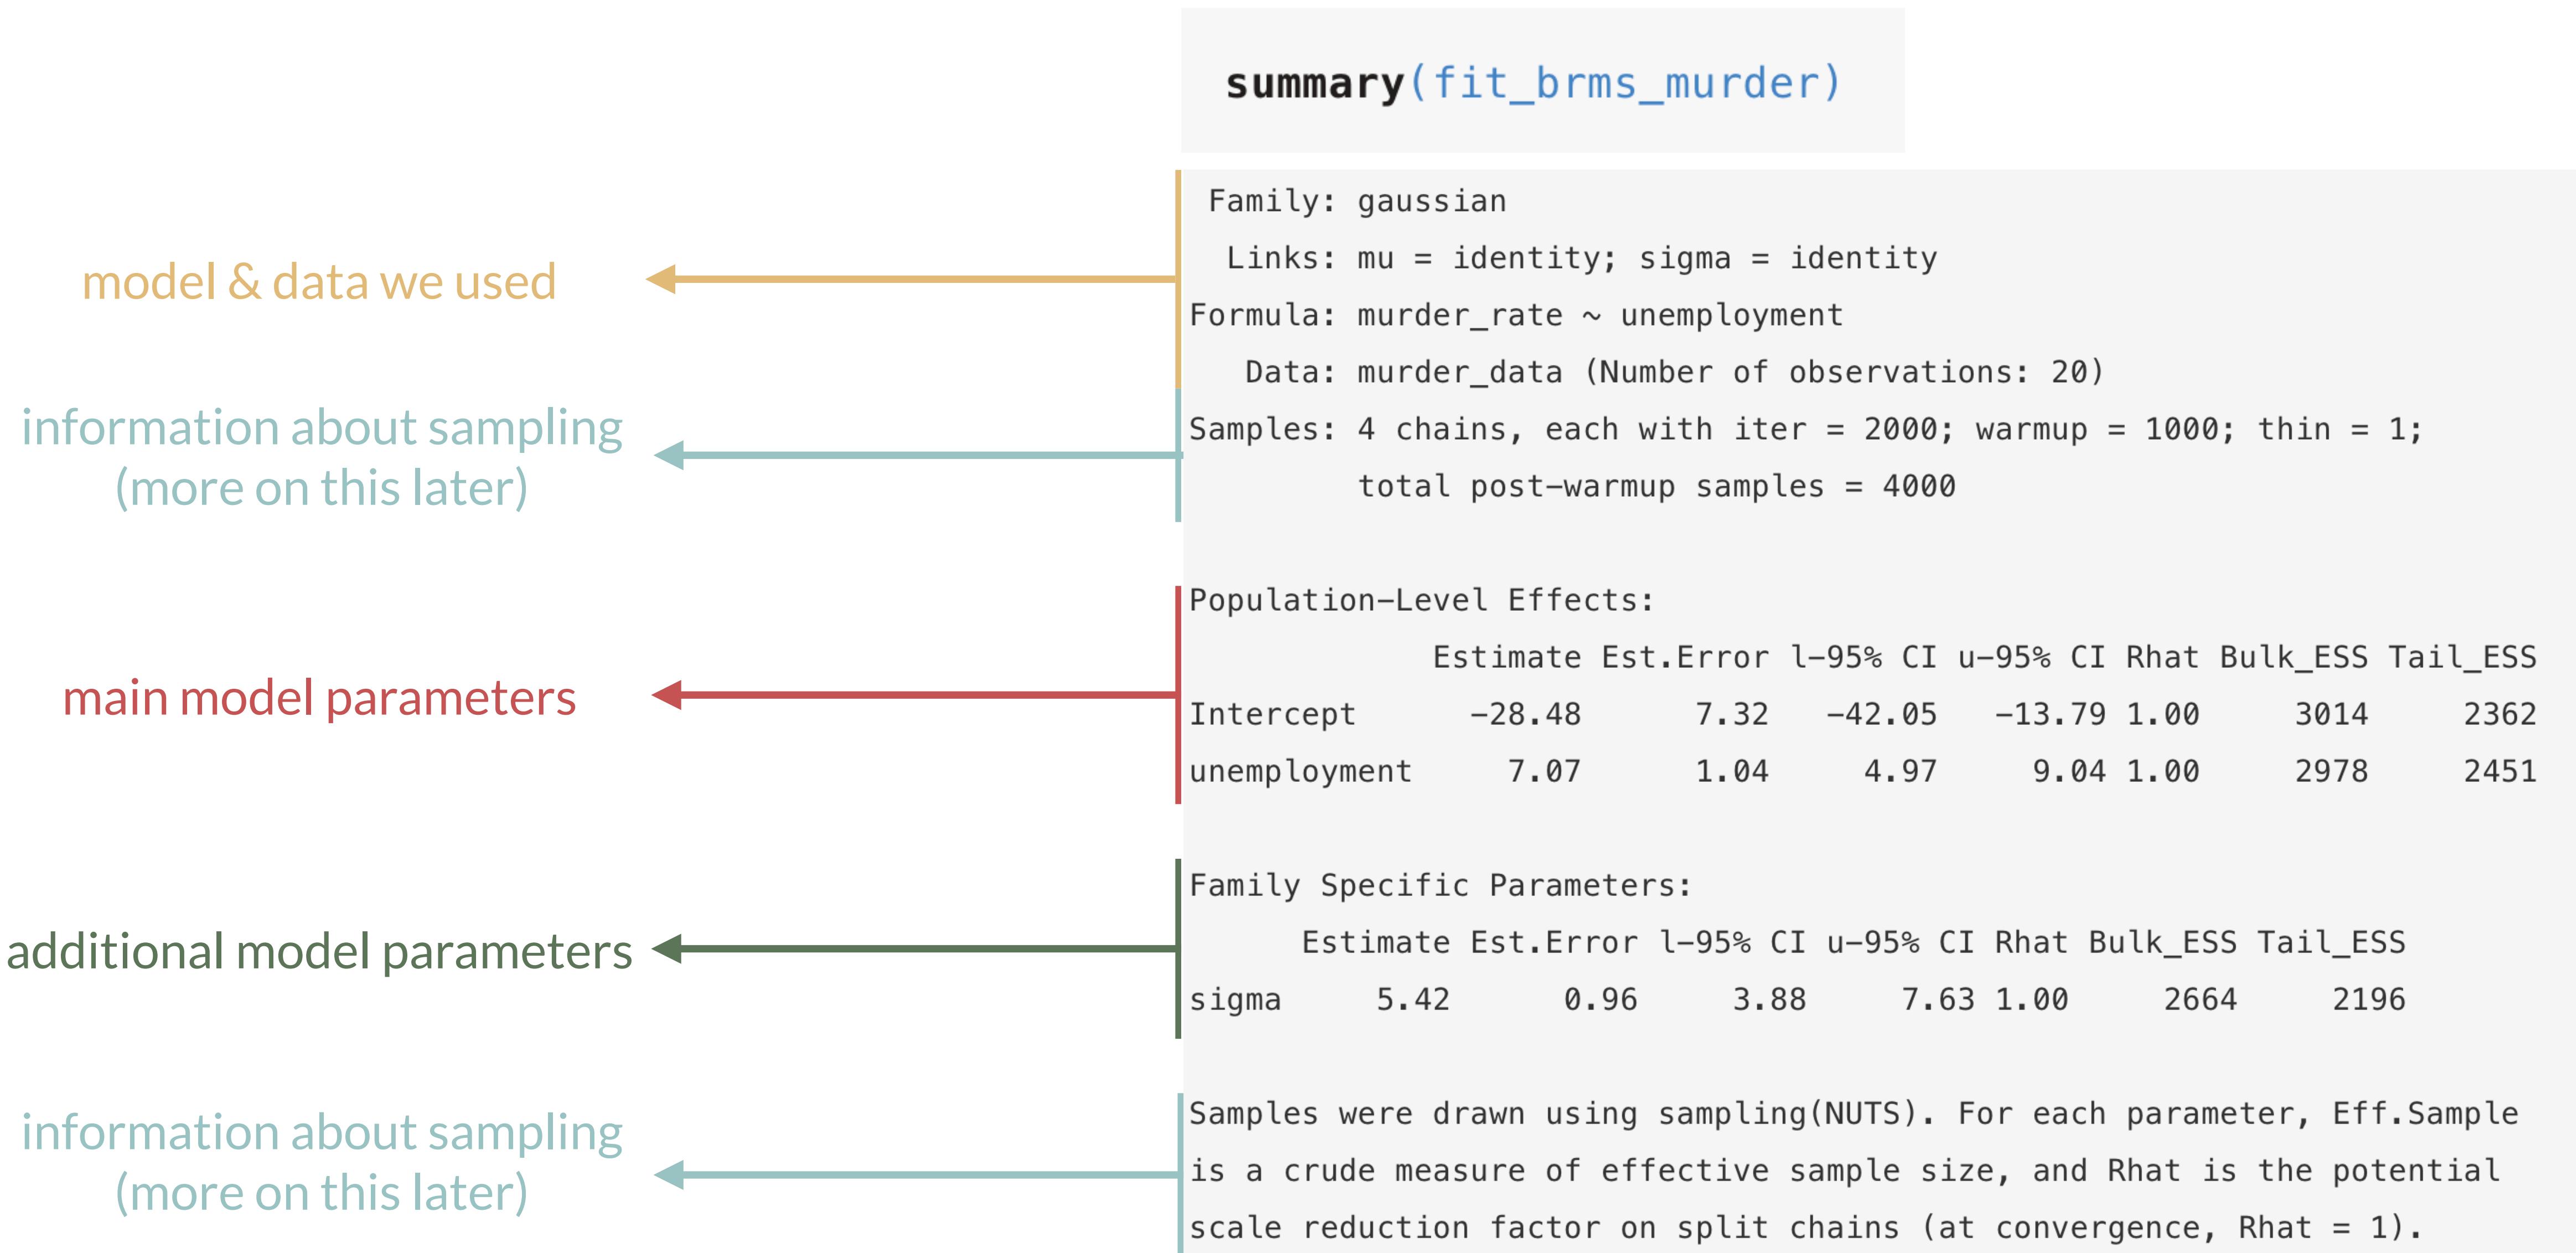
- ▶ R package BRMS provides high-level interface for Bayesian linear regression
- ▶ models are specified with R's formula syntax
- ▶ returns samples from the posterior distribution
 - alternatives: MAPs, variational inference
- ▶ builds on probabilistic programming language Stan
 - powerful, cutting-edge tool for Bayesian computation
 - strong, non-commercial development team
 - many interfaces: stand-alone, R, Python, Julia, ...

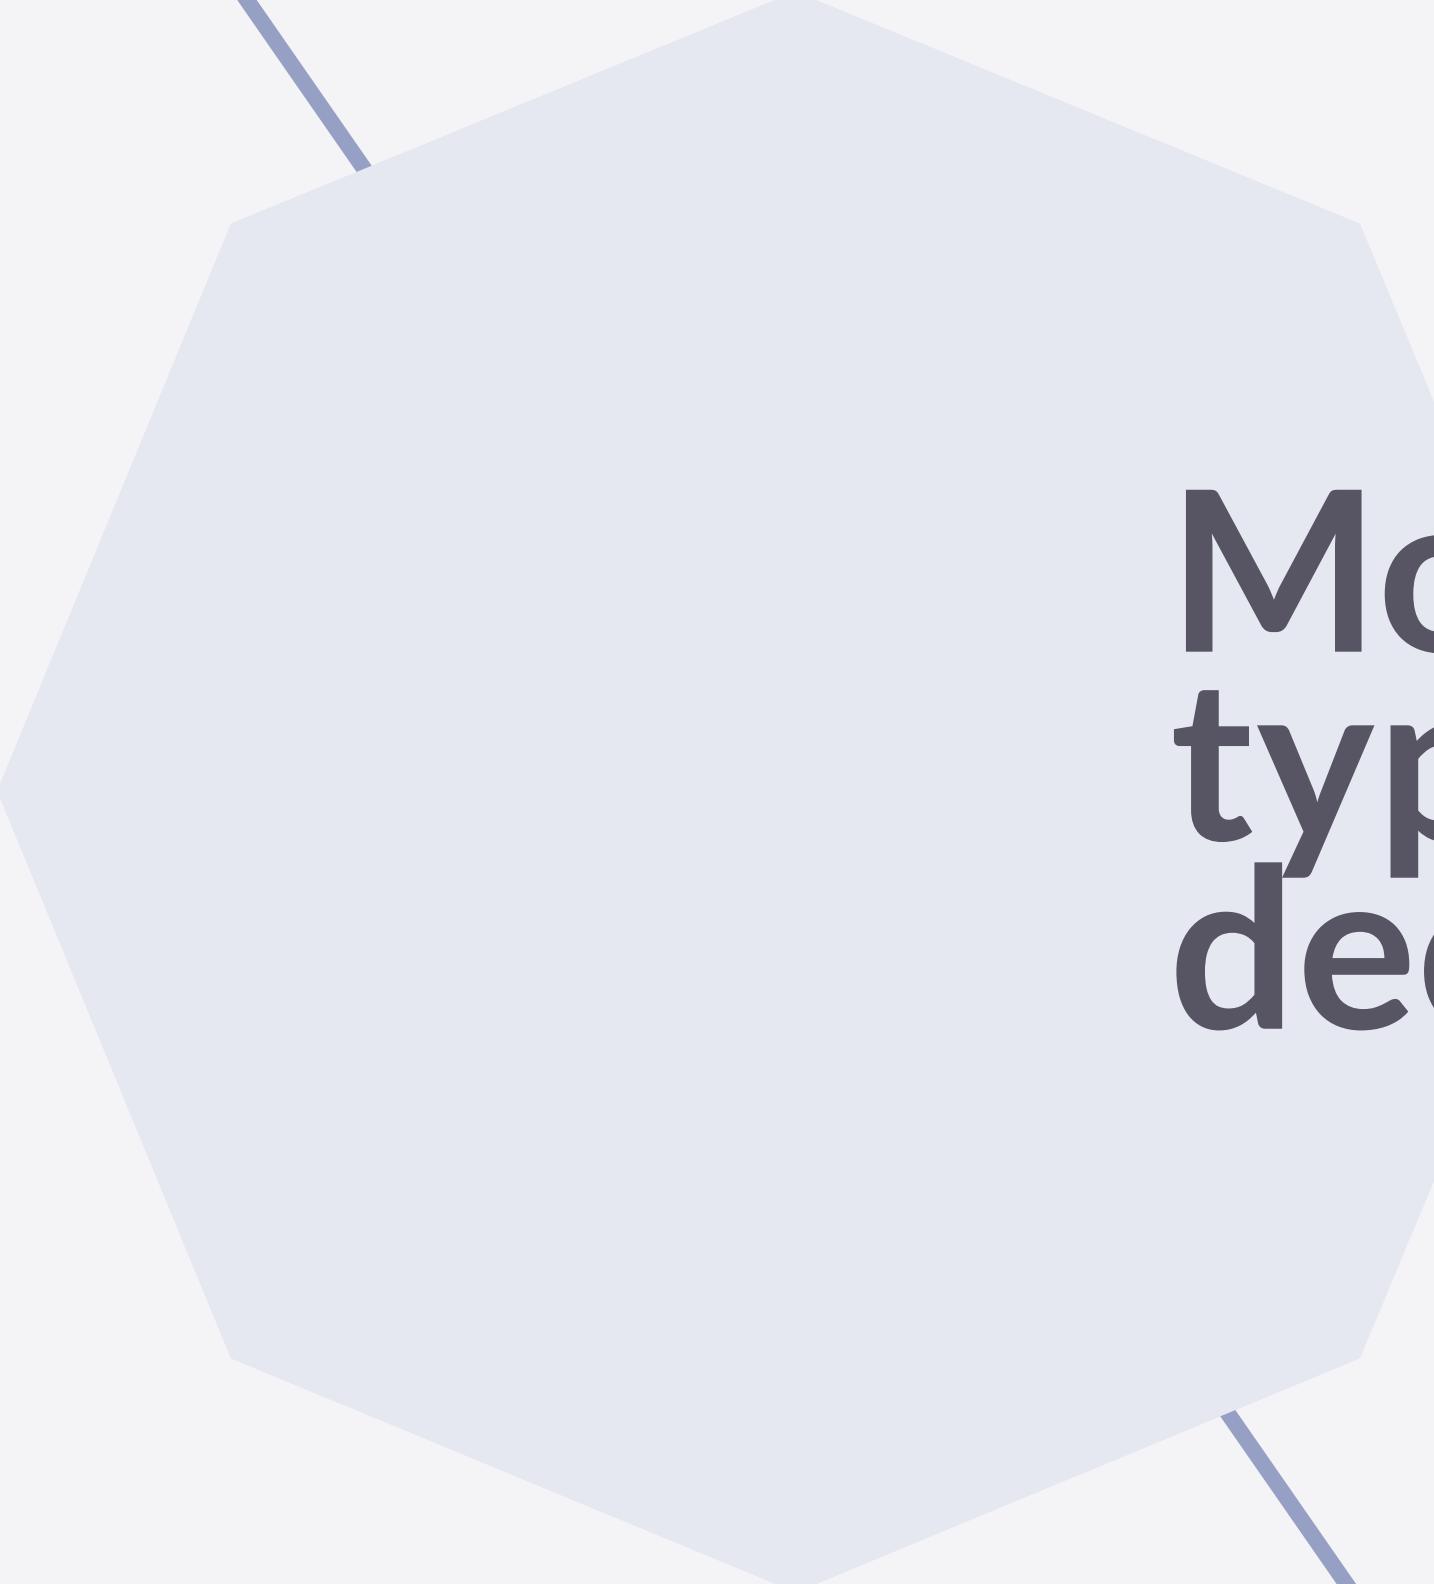
```
fit_brms_murder <- brm(  
  # specify what to explain in terms of what  
  # using the formula syntax  
  formula = murder_rate ~ unemployment,  
  # which data to use  
  data = murder_data  
)
```



Stan

Navigating BRMS output





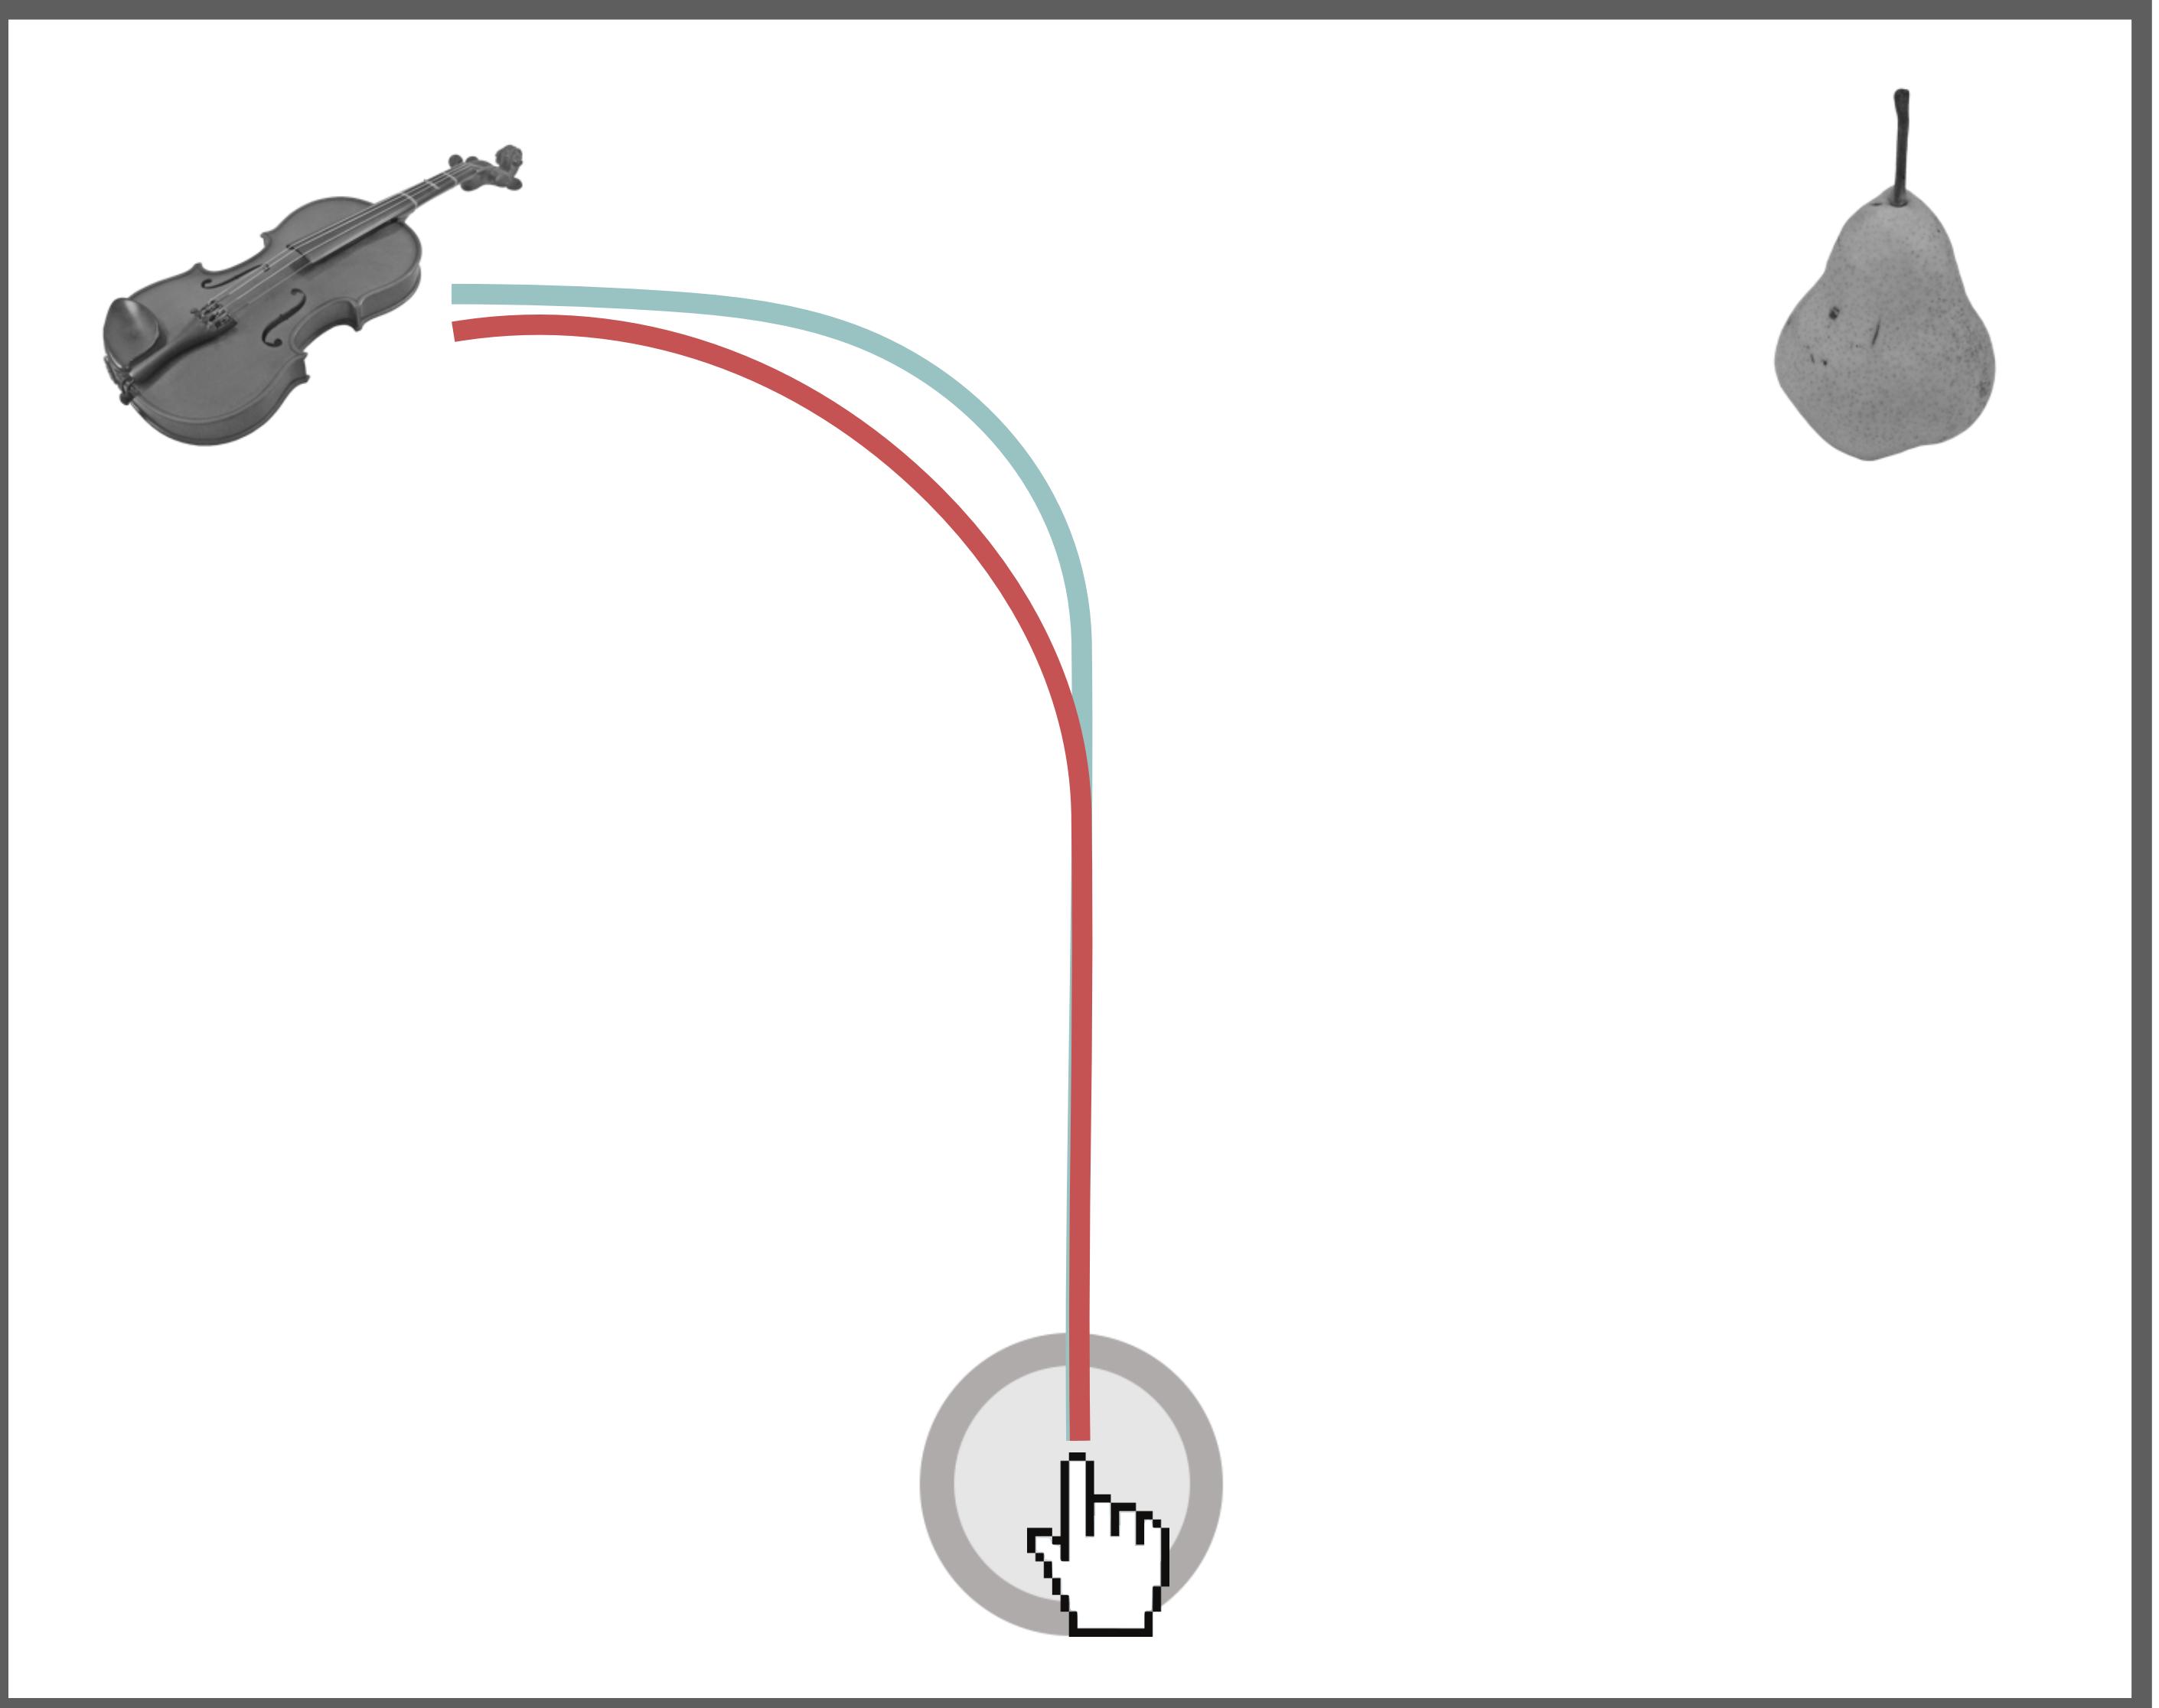
Mouse-tracking data on
typicality in category
decisions

Mouse-tracking

Hand-movement during decision making



- ▶ general idea: motor-execution provides information about the ongoing decision process
 - uncertainty
 - gradual evidence accumulation
 - change-of-mind
 - time-point of decision
 - ...
- ▶ many subtle design decisions
 - click vs touch
 - move horizontally or vertically
 - ...

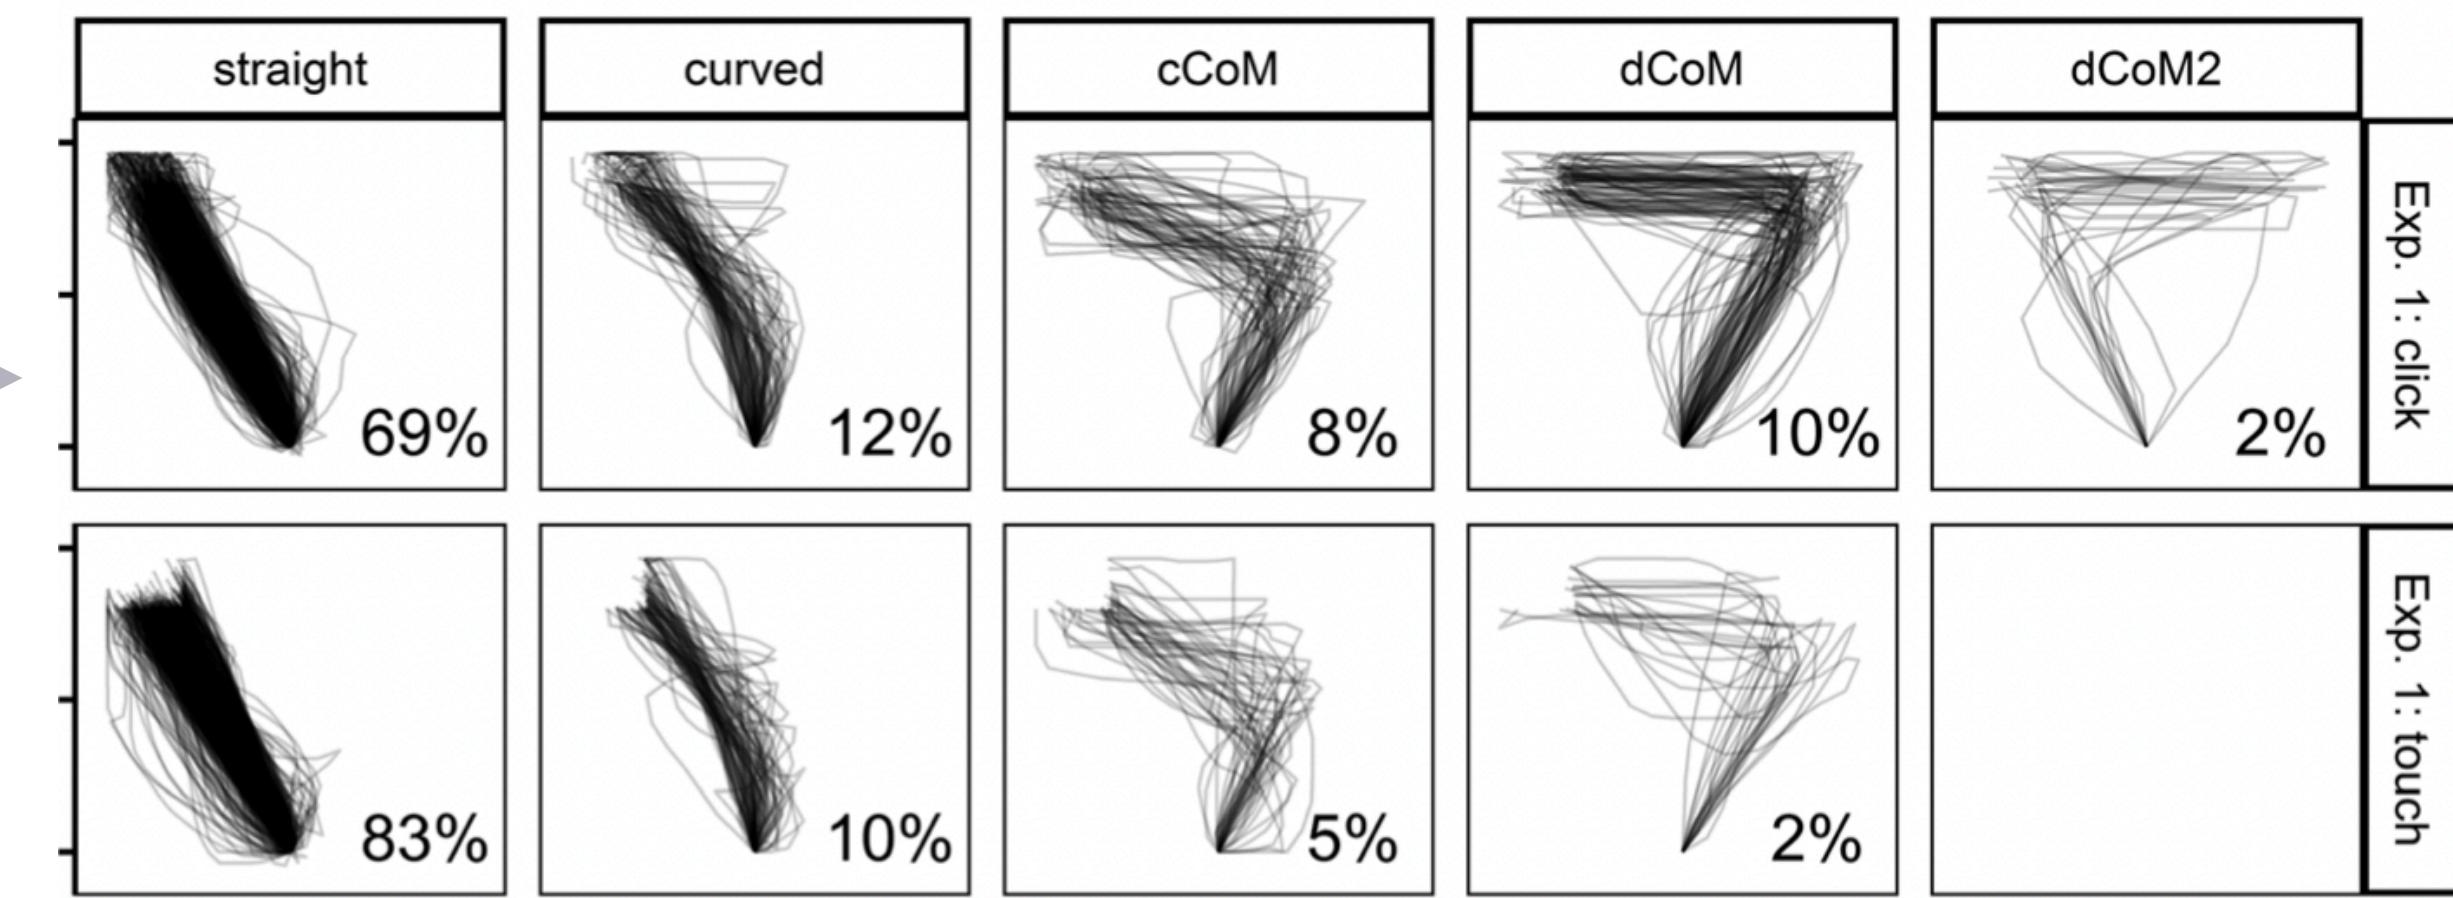
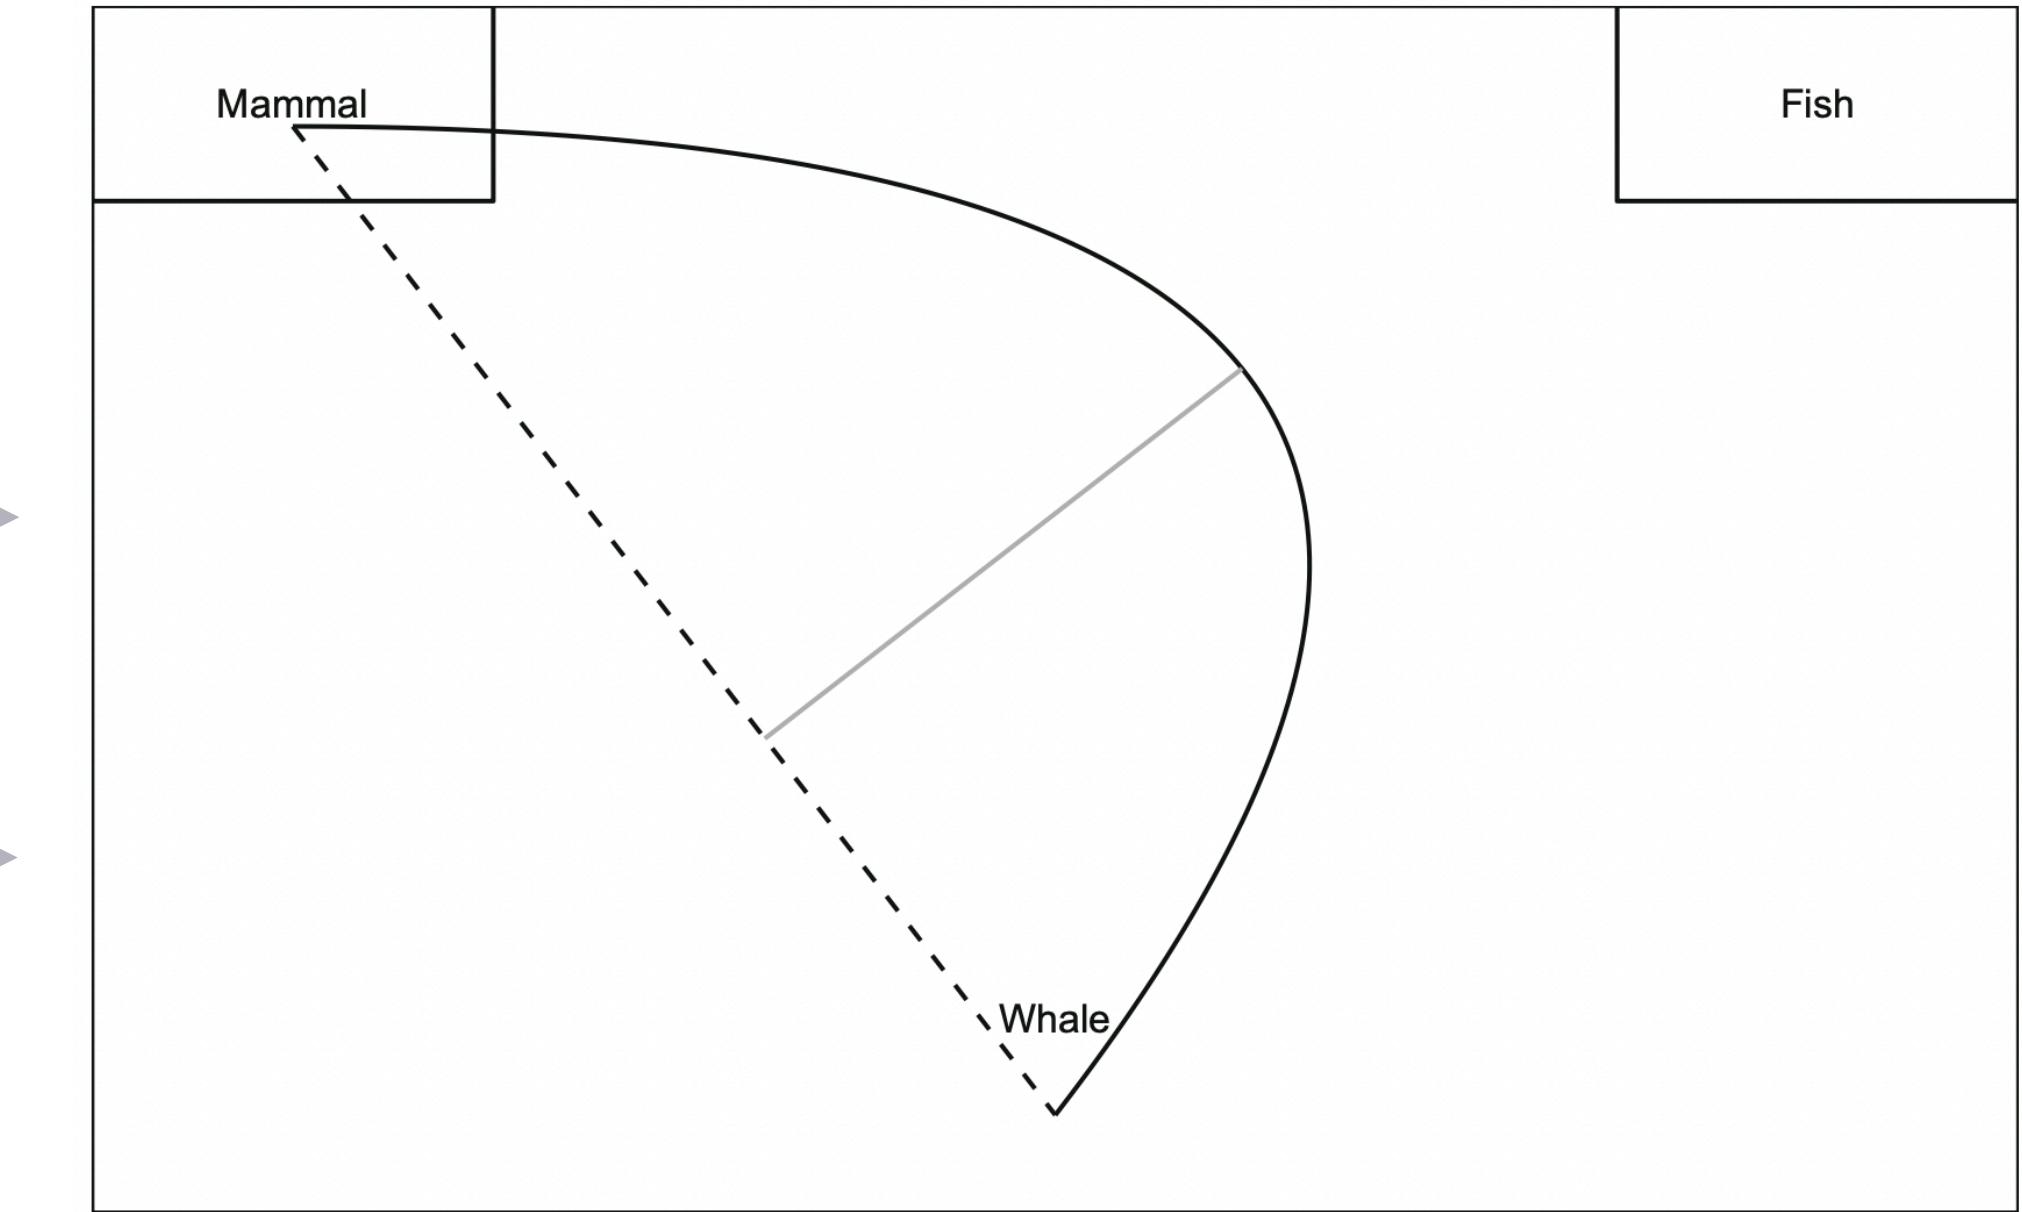
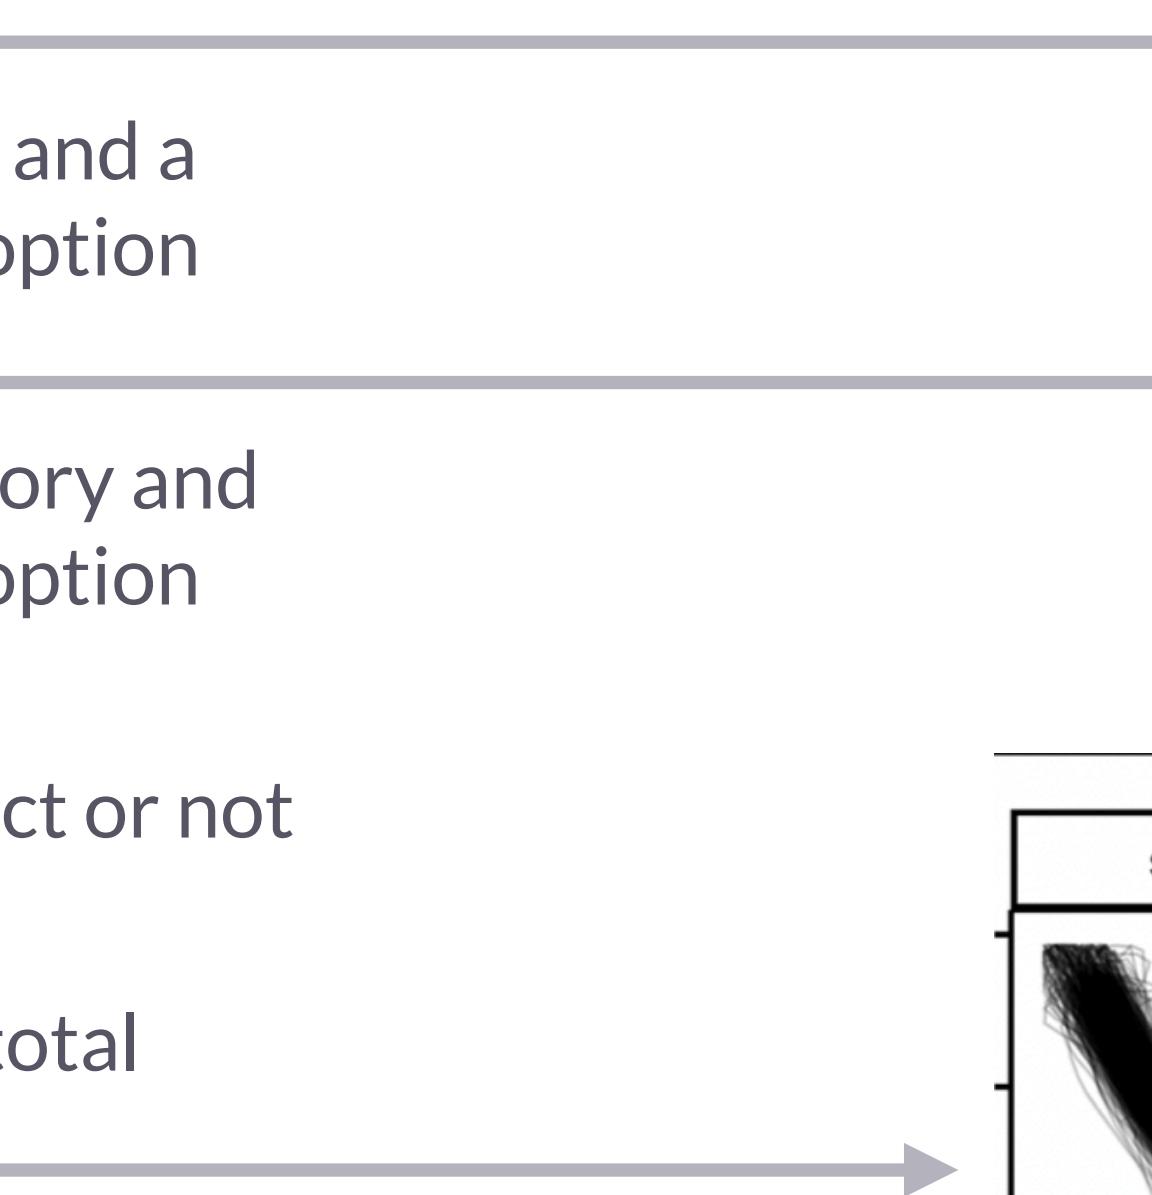


Mouse-tracking

common measures of mouse-trajectories



- ▶ raw data are lists of triples
 - (time, x-position, y-position)
- ▶ commonly used measures
 - area-under the curve (AUC)
 - area between the mouse trajectory and a straight line from start to selected option
 - maximal deviation (MAD)
 - maximum distance between trajectory and straight line from start to selected option
 - correctness
 - whether choice of option was correct or not
 - reaction time (RT)
 - how long did the movement last in total
 - type of trajectory
 - result of clustering analysis based on shape of the trajectories (usually some 3-5 categories)
 - x-flips
 - number of times the trajectory crossed the vertical middle line (at x = 0)



Running example

category recognition for typical vs atypical exemplars

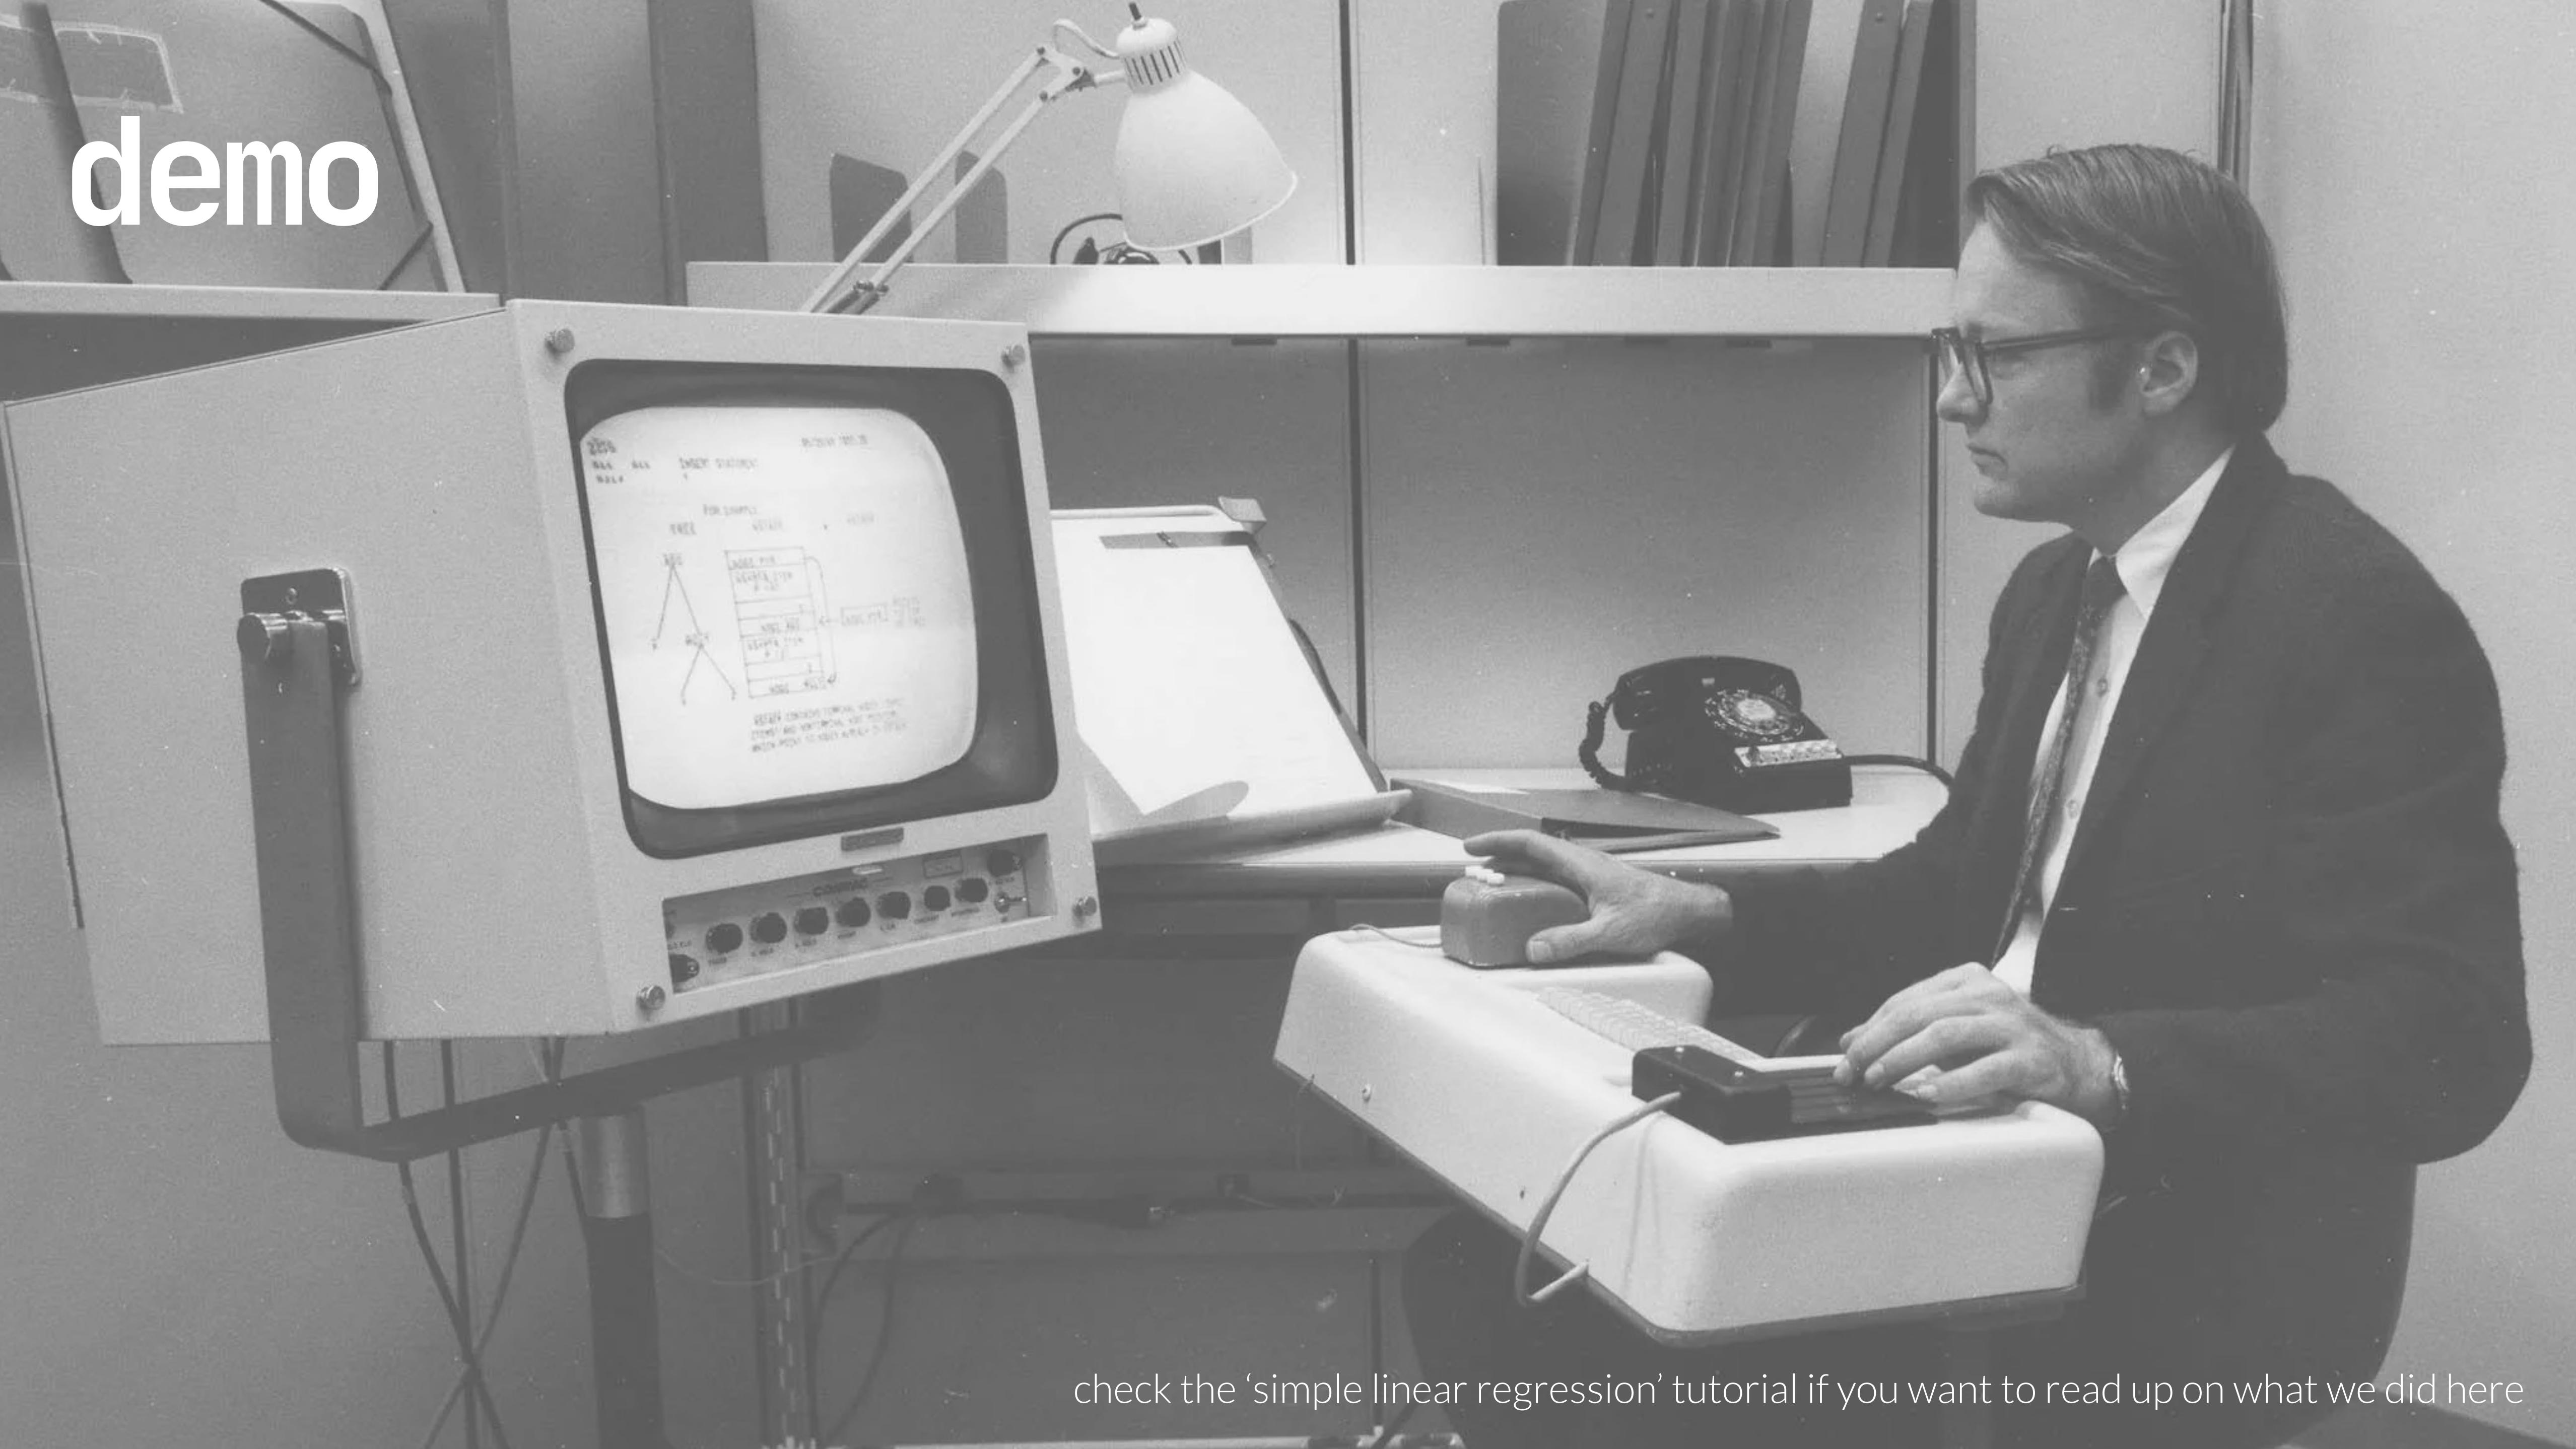


- ▶ materials & procedure
 - participants read an animal name (e.g. 'dolphin')
 - they choose the true category the animal belongs to (e.g., 'fish' or 'mammal')
 - some trigger words are typical others atypical representatives of the true category
- ▶ methodological investigation:
 - two groups: **click vs touch** to select category
- ▶ **hypothesis:** typical exemplars are easier to categorize than atypical ones
 - fewer mistakes
 - smaller RTs, AUC, MAD
 - less x-flips
 - less "change-of-mind" curve types
- ▶ **research question (methods):** any differences between click & touch selection?

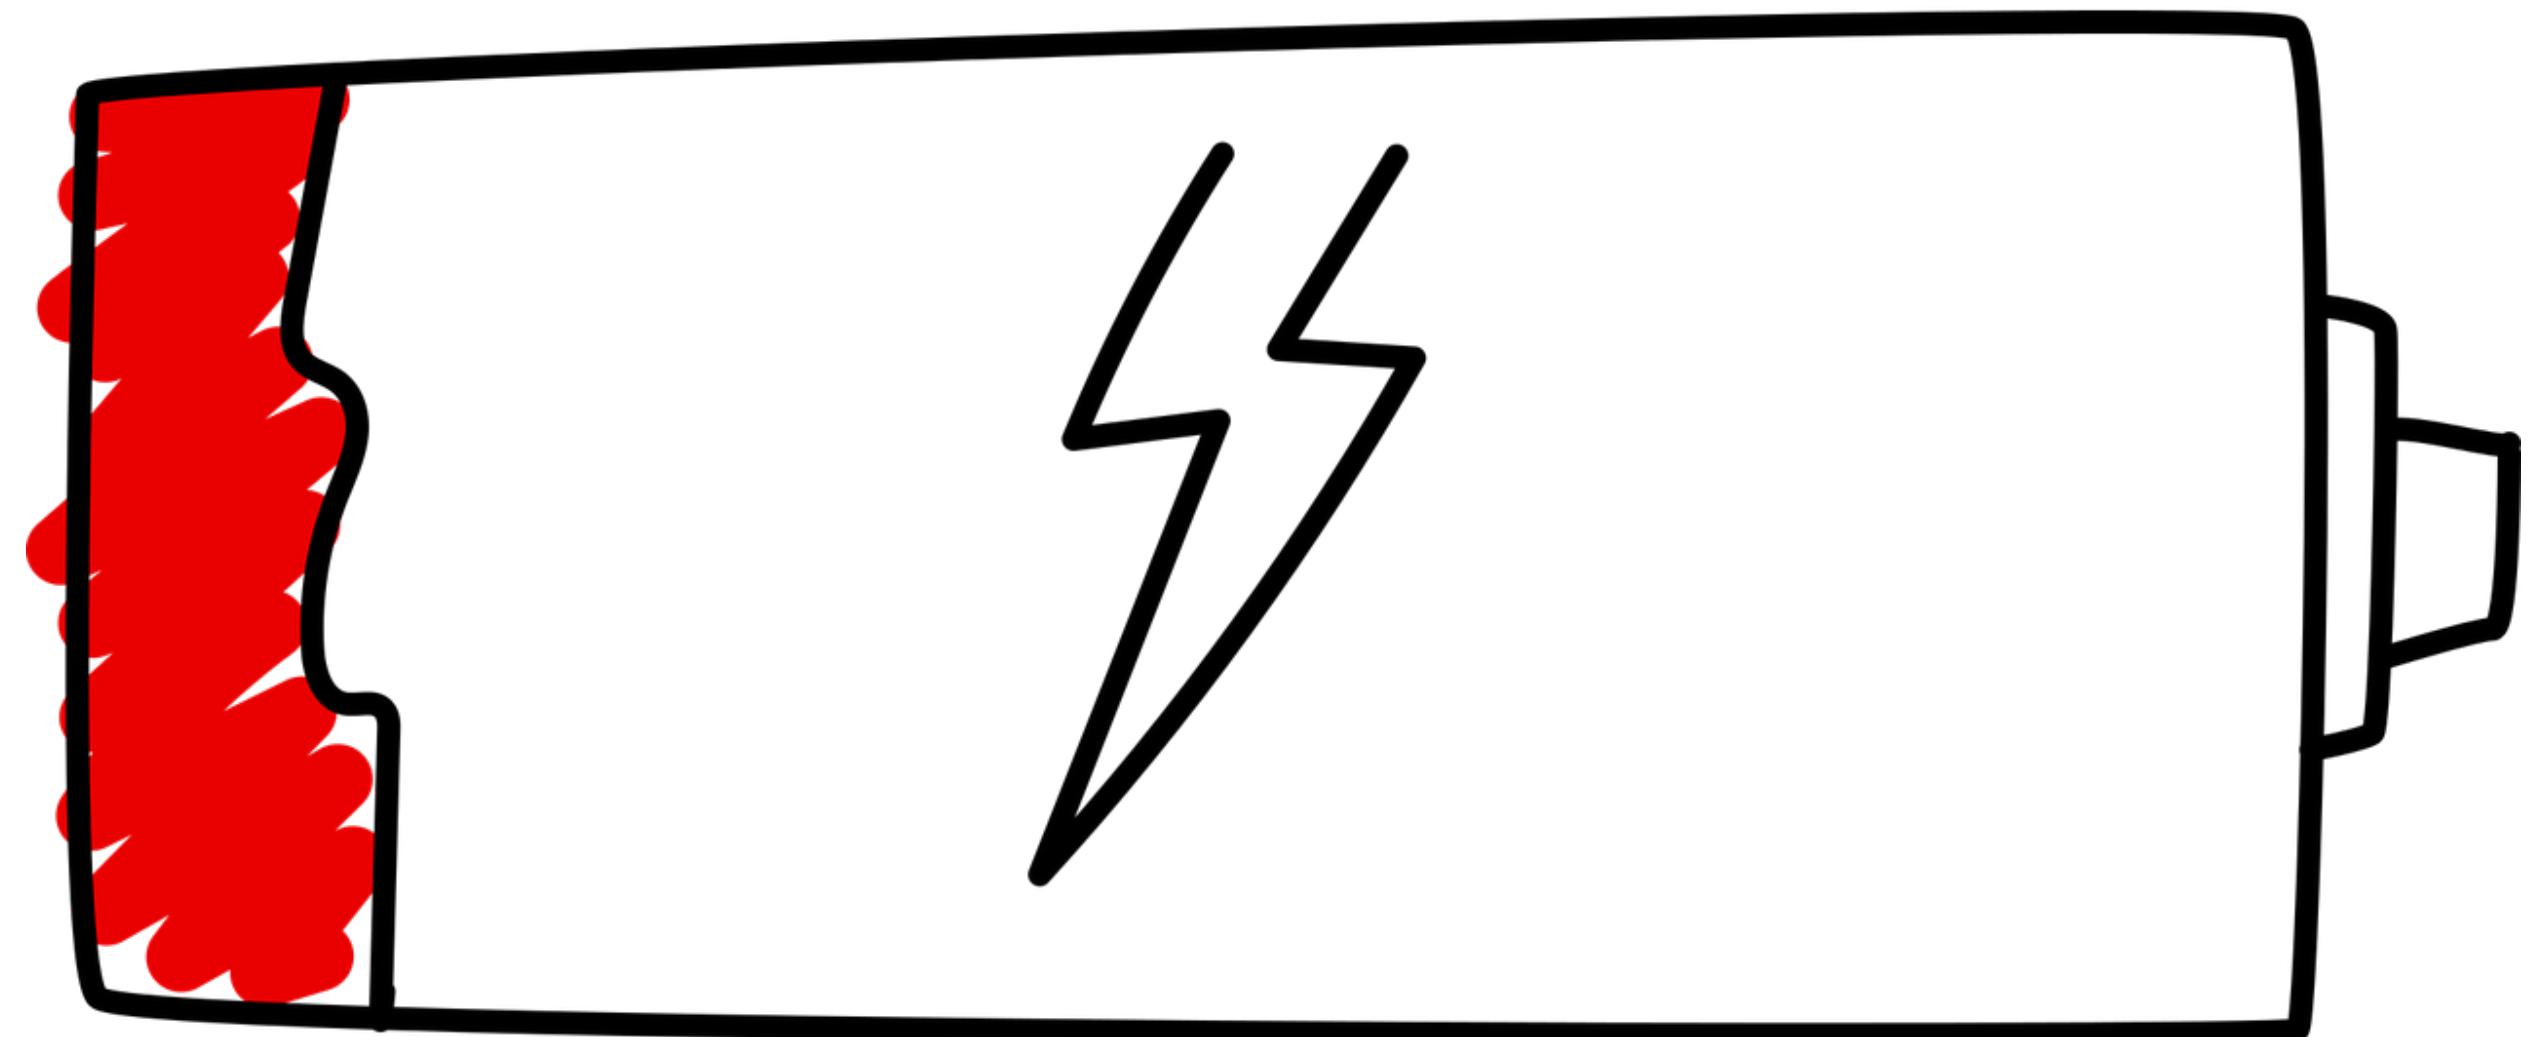
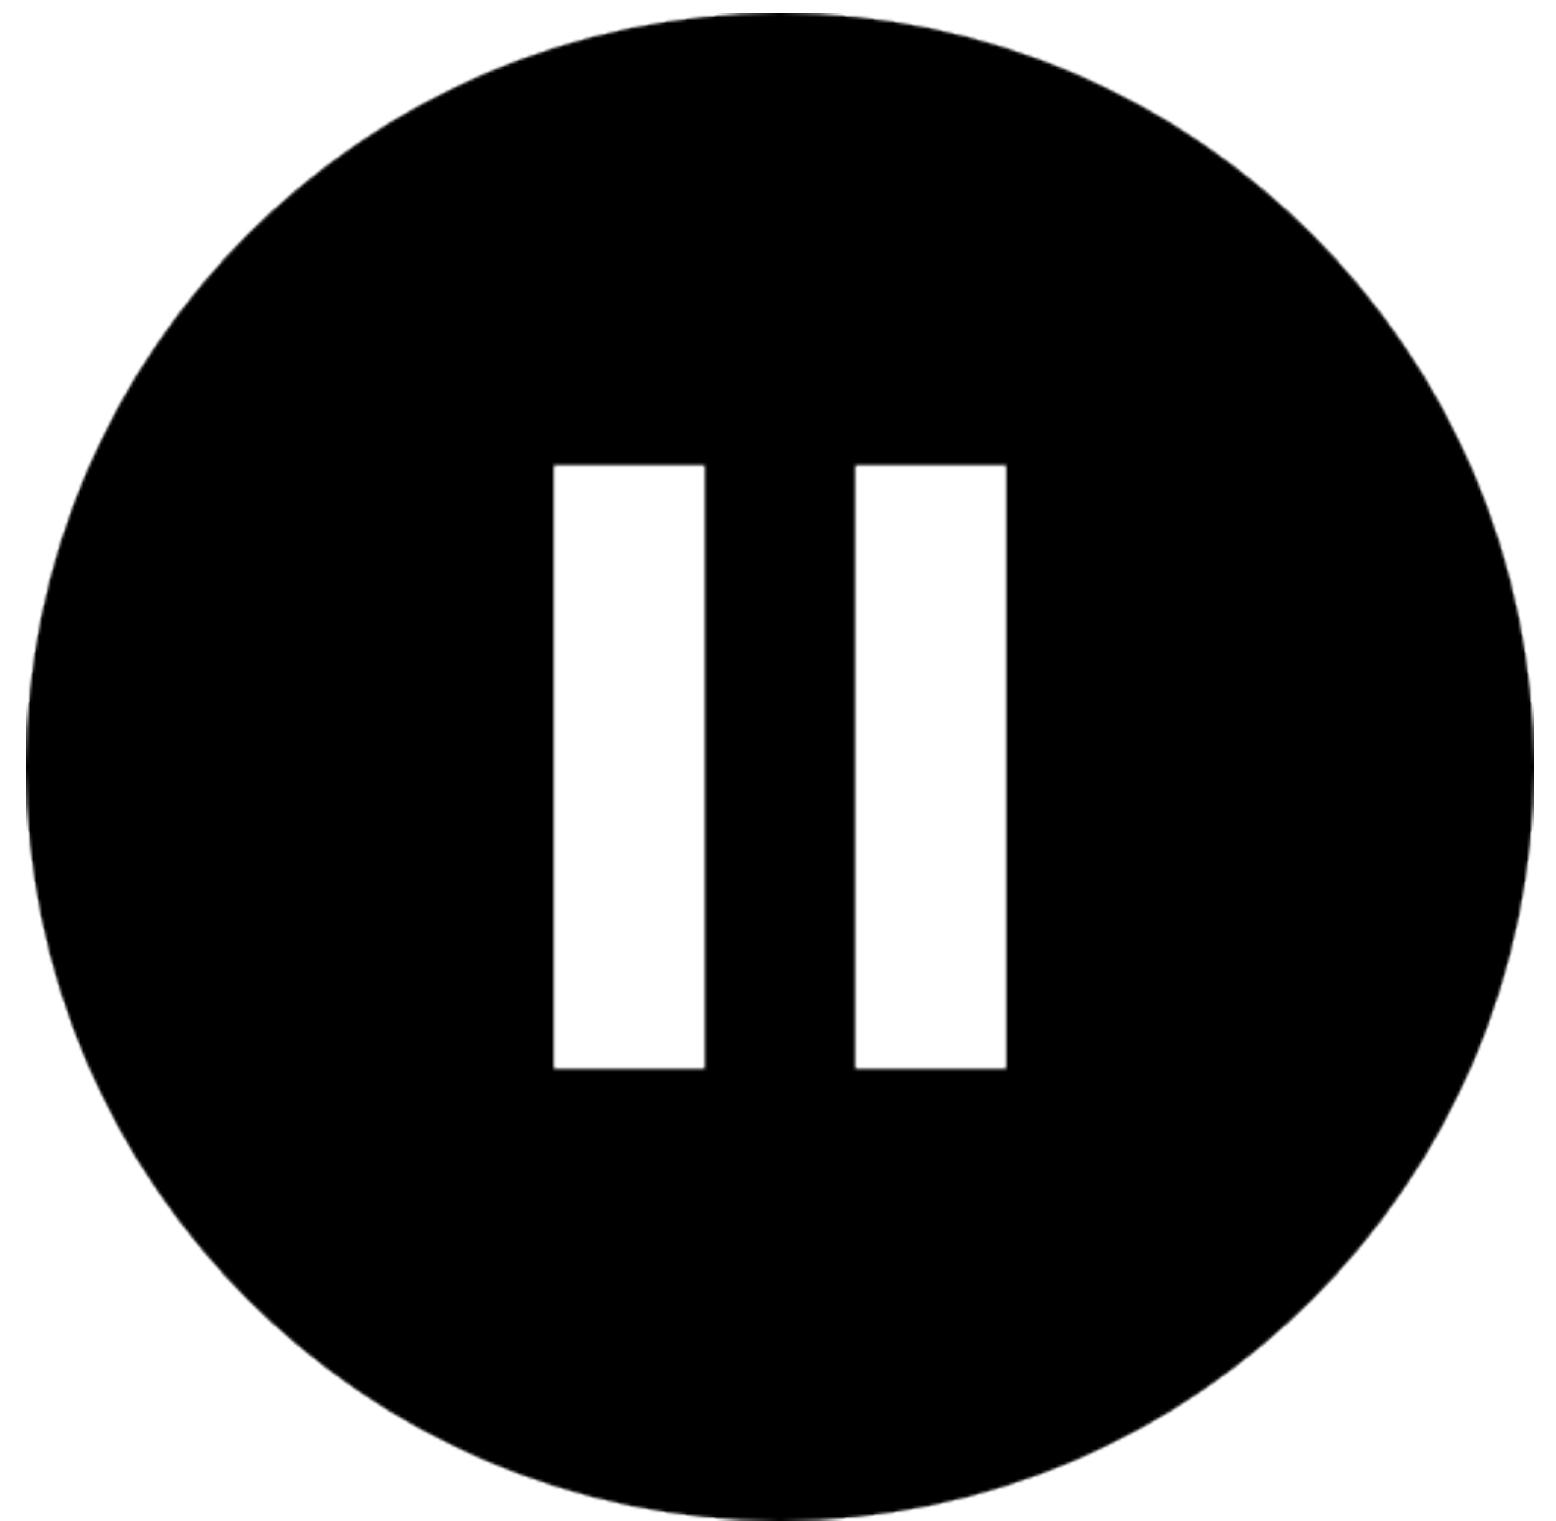
variables used in the data set

- `trial_id` = unique id for individual trials
- `MAD` = maximal deviation into competitor space
- `AUC` = area under the curve
- `xpos_flips` = the amount of horizontal direction changes
- `RT` = reaction time in ms
- `prototype_label` = different categories of prototypical movement strategies
- `subject_id` = unique id for individual participants
- `group` = groups differ in the response design (click vs. touch)
- `condition` = category membership (Typical vs. Atypical)
- `exemplar` = the concrete animal
- `category_left` = the category displayed on the left
- `category_right` = the category displayed on the right
- `category_correct` = the category that is correct
- `response` = the selected category
- `correct` = whether or not the `response` matches `category_correct`

demo



check the 'simple linear regression' tutorial if you want to read up on what we did here





Bayesian parameter estimation

Computing posterior distributions

problem of computational complexity

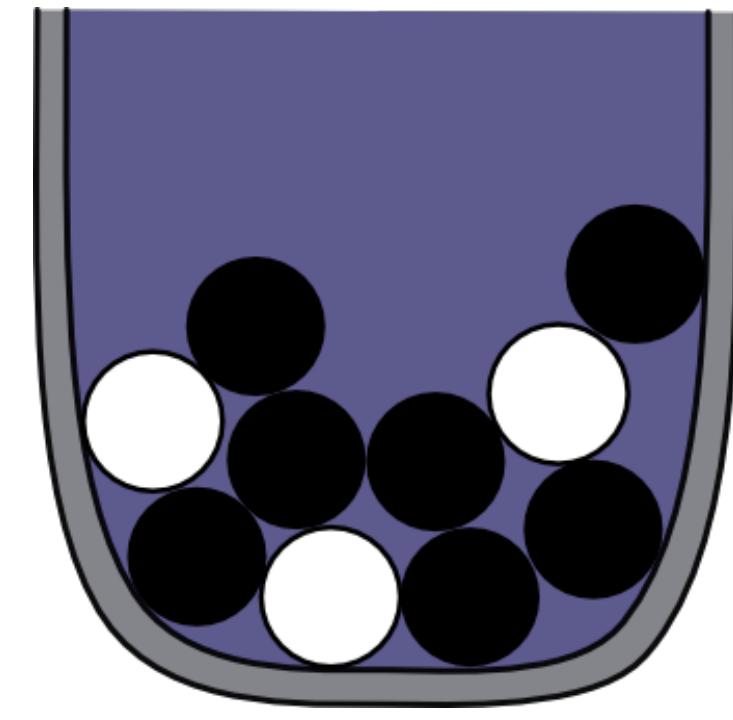
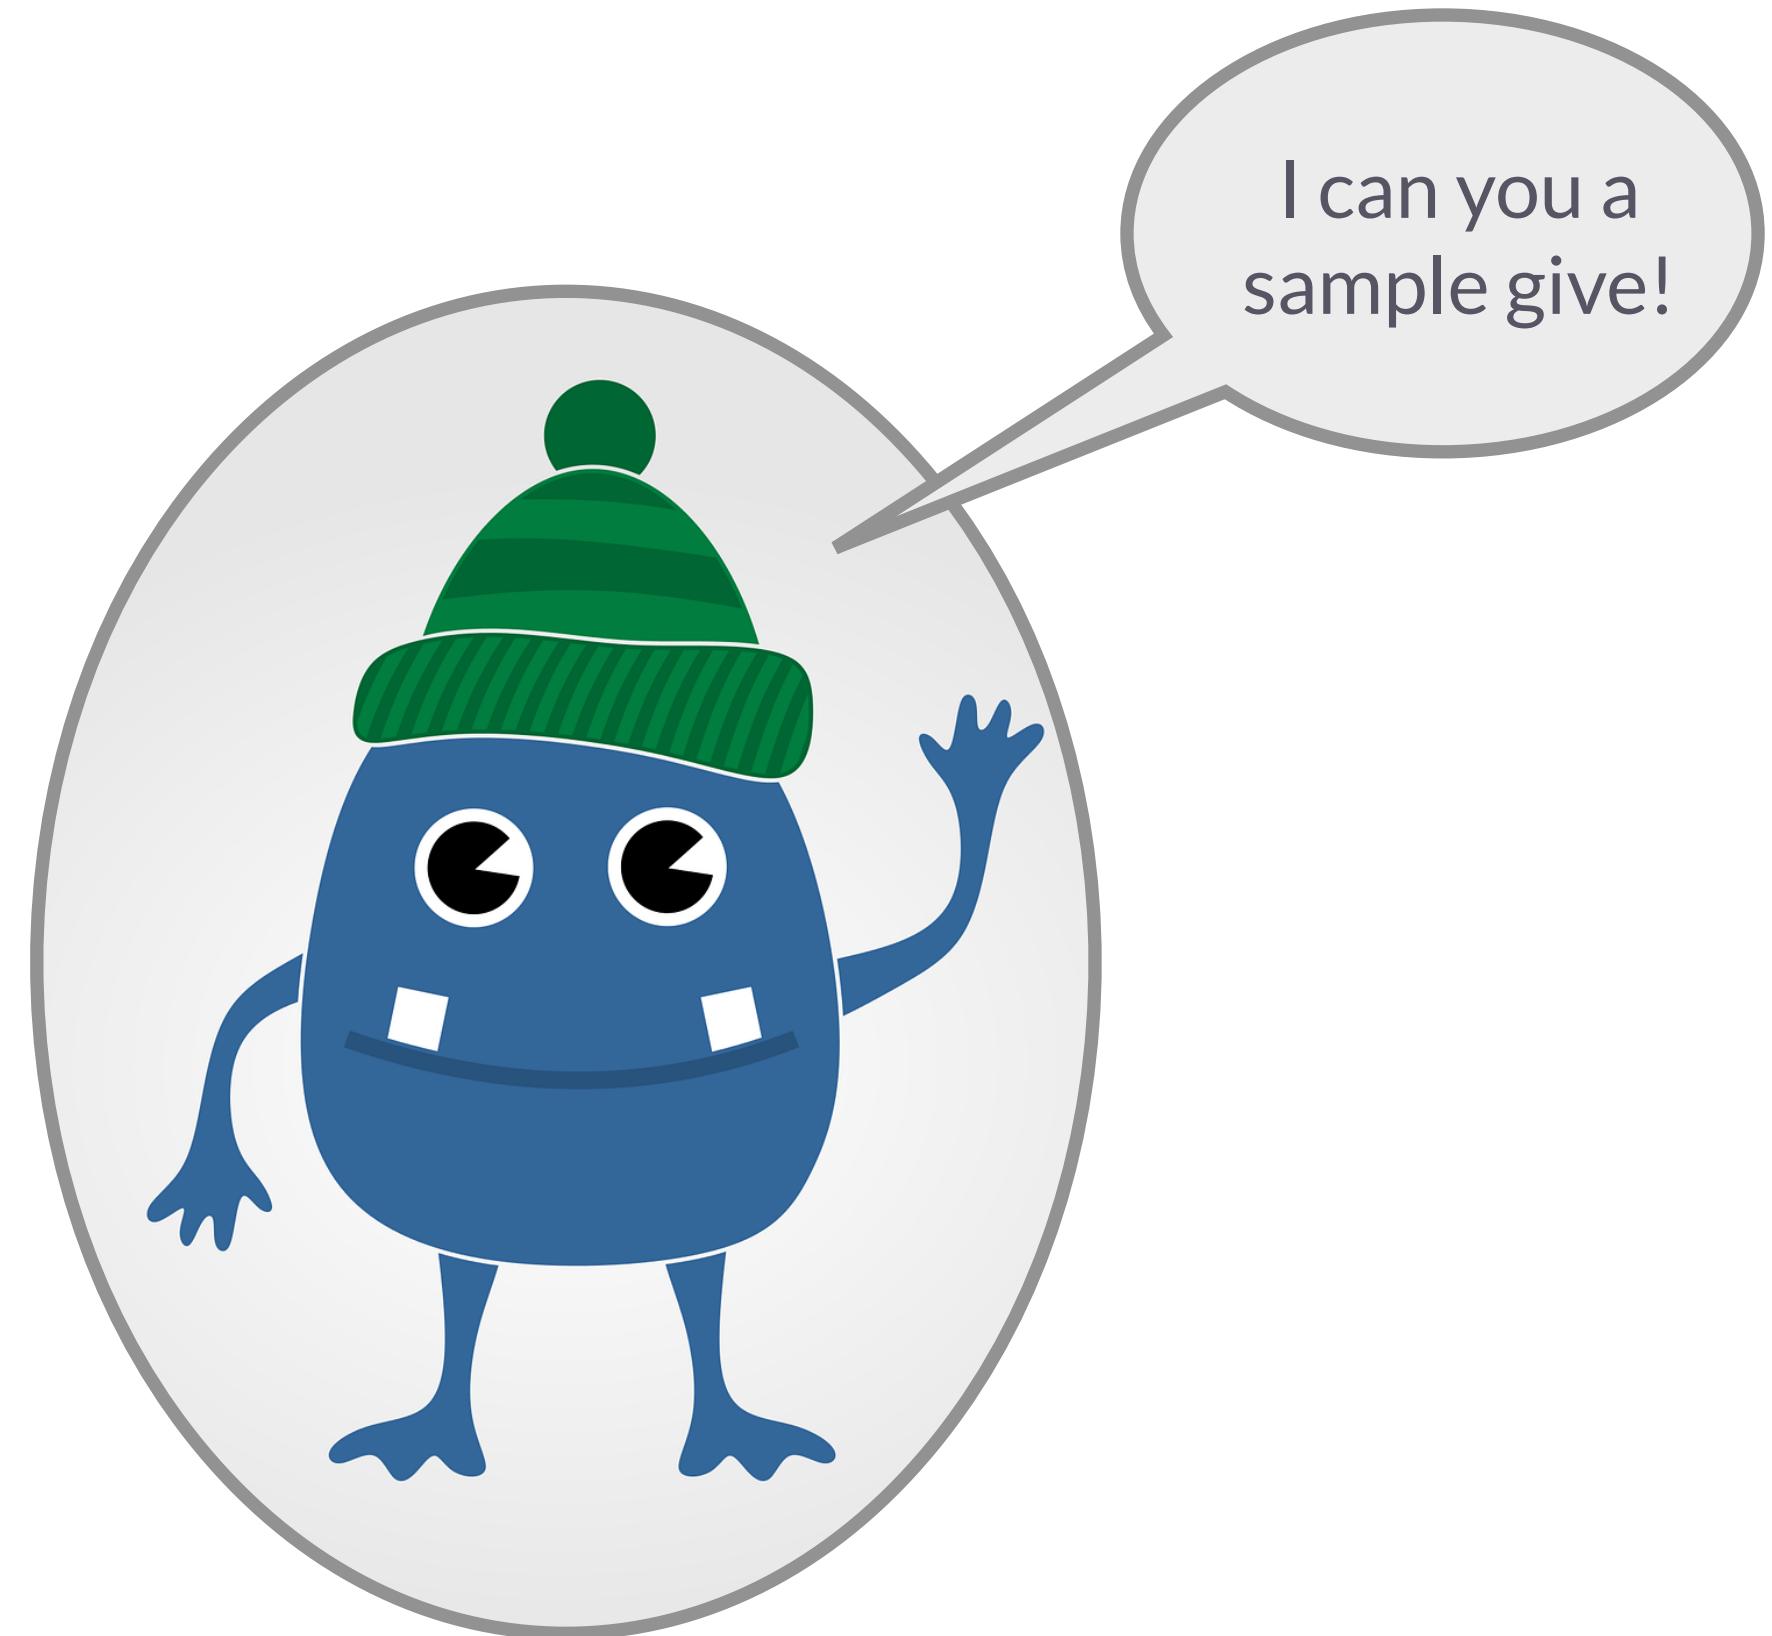
$$P(\theta | D) = \frac{P(D | \theta) \pi(\theta)}{\int P(D | \theta) \pi(\theta) d\theta}$$

The equation shows the posterior distribution $P(\theta | D)$ as a fraction. The numerator is $P(D | \theta) \pi(\theta)$, where $P(D | \theta)$ is labeled "fast & easy" and $\pi(\theta)$ is labeled "fast & easy". The denominator is $\int P(D | \theta) \pi(\theta) d\theta$, which is labeled "possibly intractable" with a skull and crossbones symbol.

Approximating distributions via sampling

our go-to solution for approximating posterior distributions beyond conjugacy

- ▶ we can approximate any probability distribution by either:
 - a large set of representative samples; or
 - an oracle that returns a sample if needed.



Temporal development of the proportion of draws from an urn



Excursion: Posteriors from conjugacy

closed-form posteriors from clever choice of priors

- ▶ prior $P(\theta)$ is a **conjugate prior** for likelihood $P(D | \theta)$ iff prior $P(\theta)$ and posterior $P(\theta | D)$ are the same kind of probability distribution, e.g.:
 - prior: $\theta \sim \text{Beta}(1,1)$
 - posterior: $\theta | D \sim \text{Beta}(8,18)$
- ▶ **claim:** the beta distribution is a conjugate prior for the binomial likelihood function
 - proof:

$$P(\theta | k, N) \propto \text{Binomial}(k; N, \theta) \text{ Beta}(\theta | a, b)$$

$$P(\theta | k, N) \propto \theta^k (1 - \theta)^{N-k} \theta^{a-1} (1 - \theta)^{b-1}$$

$$P(\theta | k, N) \propto \theta^{k+a-1} (1 - \theta)^{N-k+b-1}$$

$$P(\theta | k, N) = \text{Beta}(\theta | k + a, N - k + b)$$

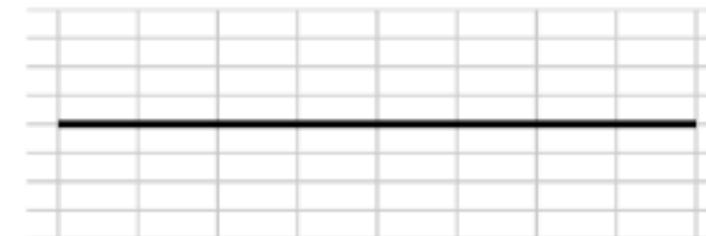


Excursion: Sequential updating

for the beta-binomial model

- ▶ sequence of updating does not matter
 - any order of single-observation updates
 - any ‘chunking’: whole data set, different subsets in whatever sequence (as long as disjoined)
- ▶ “today’s posterior is tomorrow’s prior”

$a=1$ $b=1$



$a=1$ $b=2$



$a=1$ $b=3$



$a=2$ $b=1$



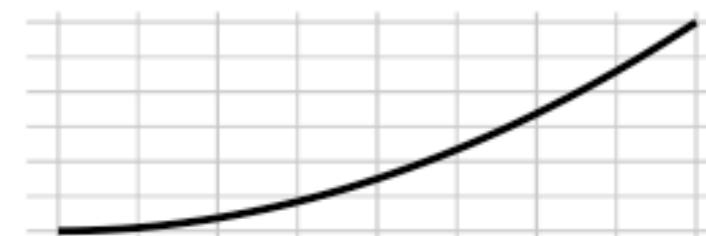
$a=2$ $b=2$



$a=2$ $b=3$



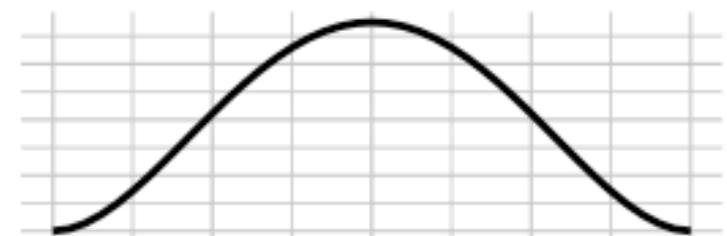
$a=3$ $b=1$



$a=3$ $b=2$



$a=3$ $b=3$



read more [here](#)

Excursion: Sequential updating

general proof

- ▶ **claim:** if $\{D_1, D_2\}$ is a partition of D , then $P(\theta | D) \propto P(\theta | D_1) P(D_2 | \theta)$

- ▶ **sketch of proof:**

$$\begin{aligned} P(\theta | D) &= \frac{P(\theta) P(D | \theta)}{\int P(\theta') P(D | \theta') d\theta'} \\ &= \frac{P(\theta) P(D_1 | \theta) P(D_2 | \theta)}{\int P(\theta') P(D_1 | \theta') P(D_2 | \theta') d\theta'} \\ &= \frac{P(\theta) P(D_1 | \theta) P(D_2 | \theta)}{\frac{k}{k} \int P(\theta') P(D_1 | \theta') P(D_2 | \theta') d\theta'} \\ &= \frac{\frac{P(\theta) P(D_1 | \theta)}{k} P(D_2 | \theta)}{\int \frac{P(\theta') P(D_1 | \theta')}{k} P(D_2 | \theta') d\theta'} \\ &= \frac{P(\theta | D_1) P(D_2 | \theta)}{\int P(\theta' | D_1) P(D_2 | \theta') d\theta'} \end{aligned}$$

[from multiplicativity of likelihood]

[for random positive k]

[rules of integration; basic calculus]

[Bayes rule with $k = \int P(\theta) P(D_1 | \theta) d\theta$]

read more [here](#)

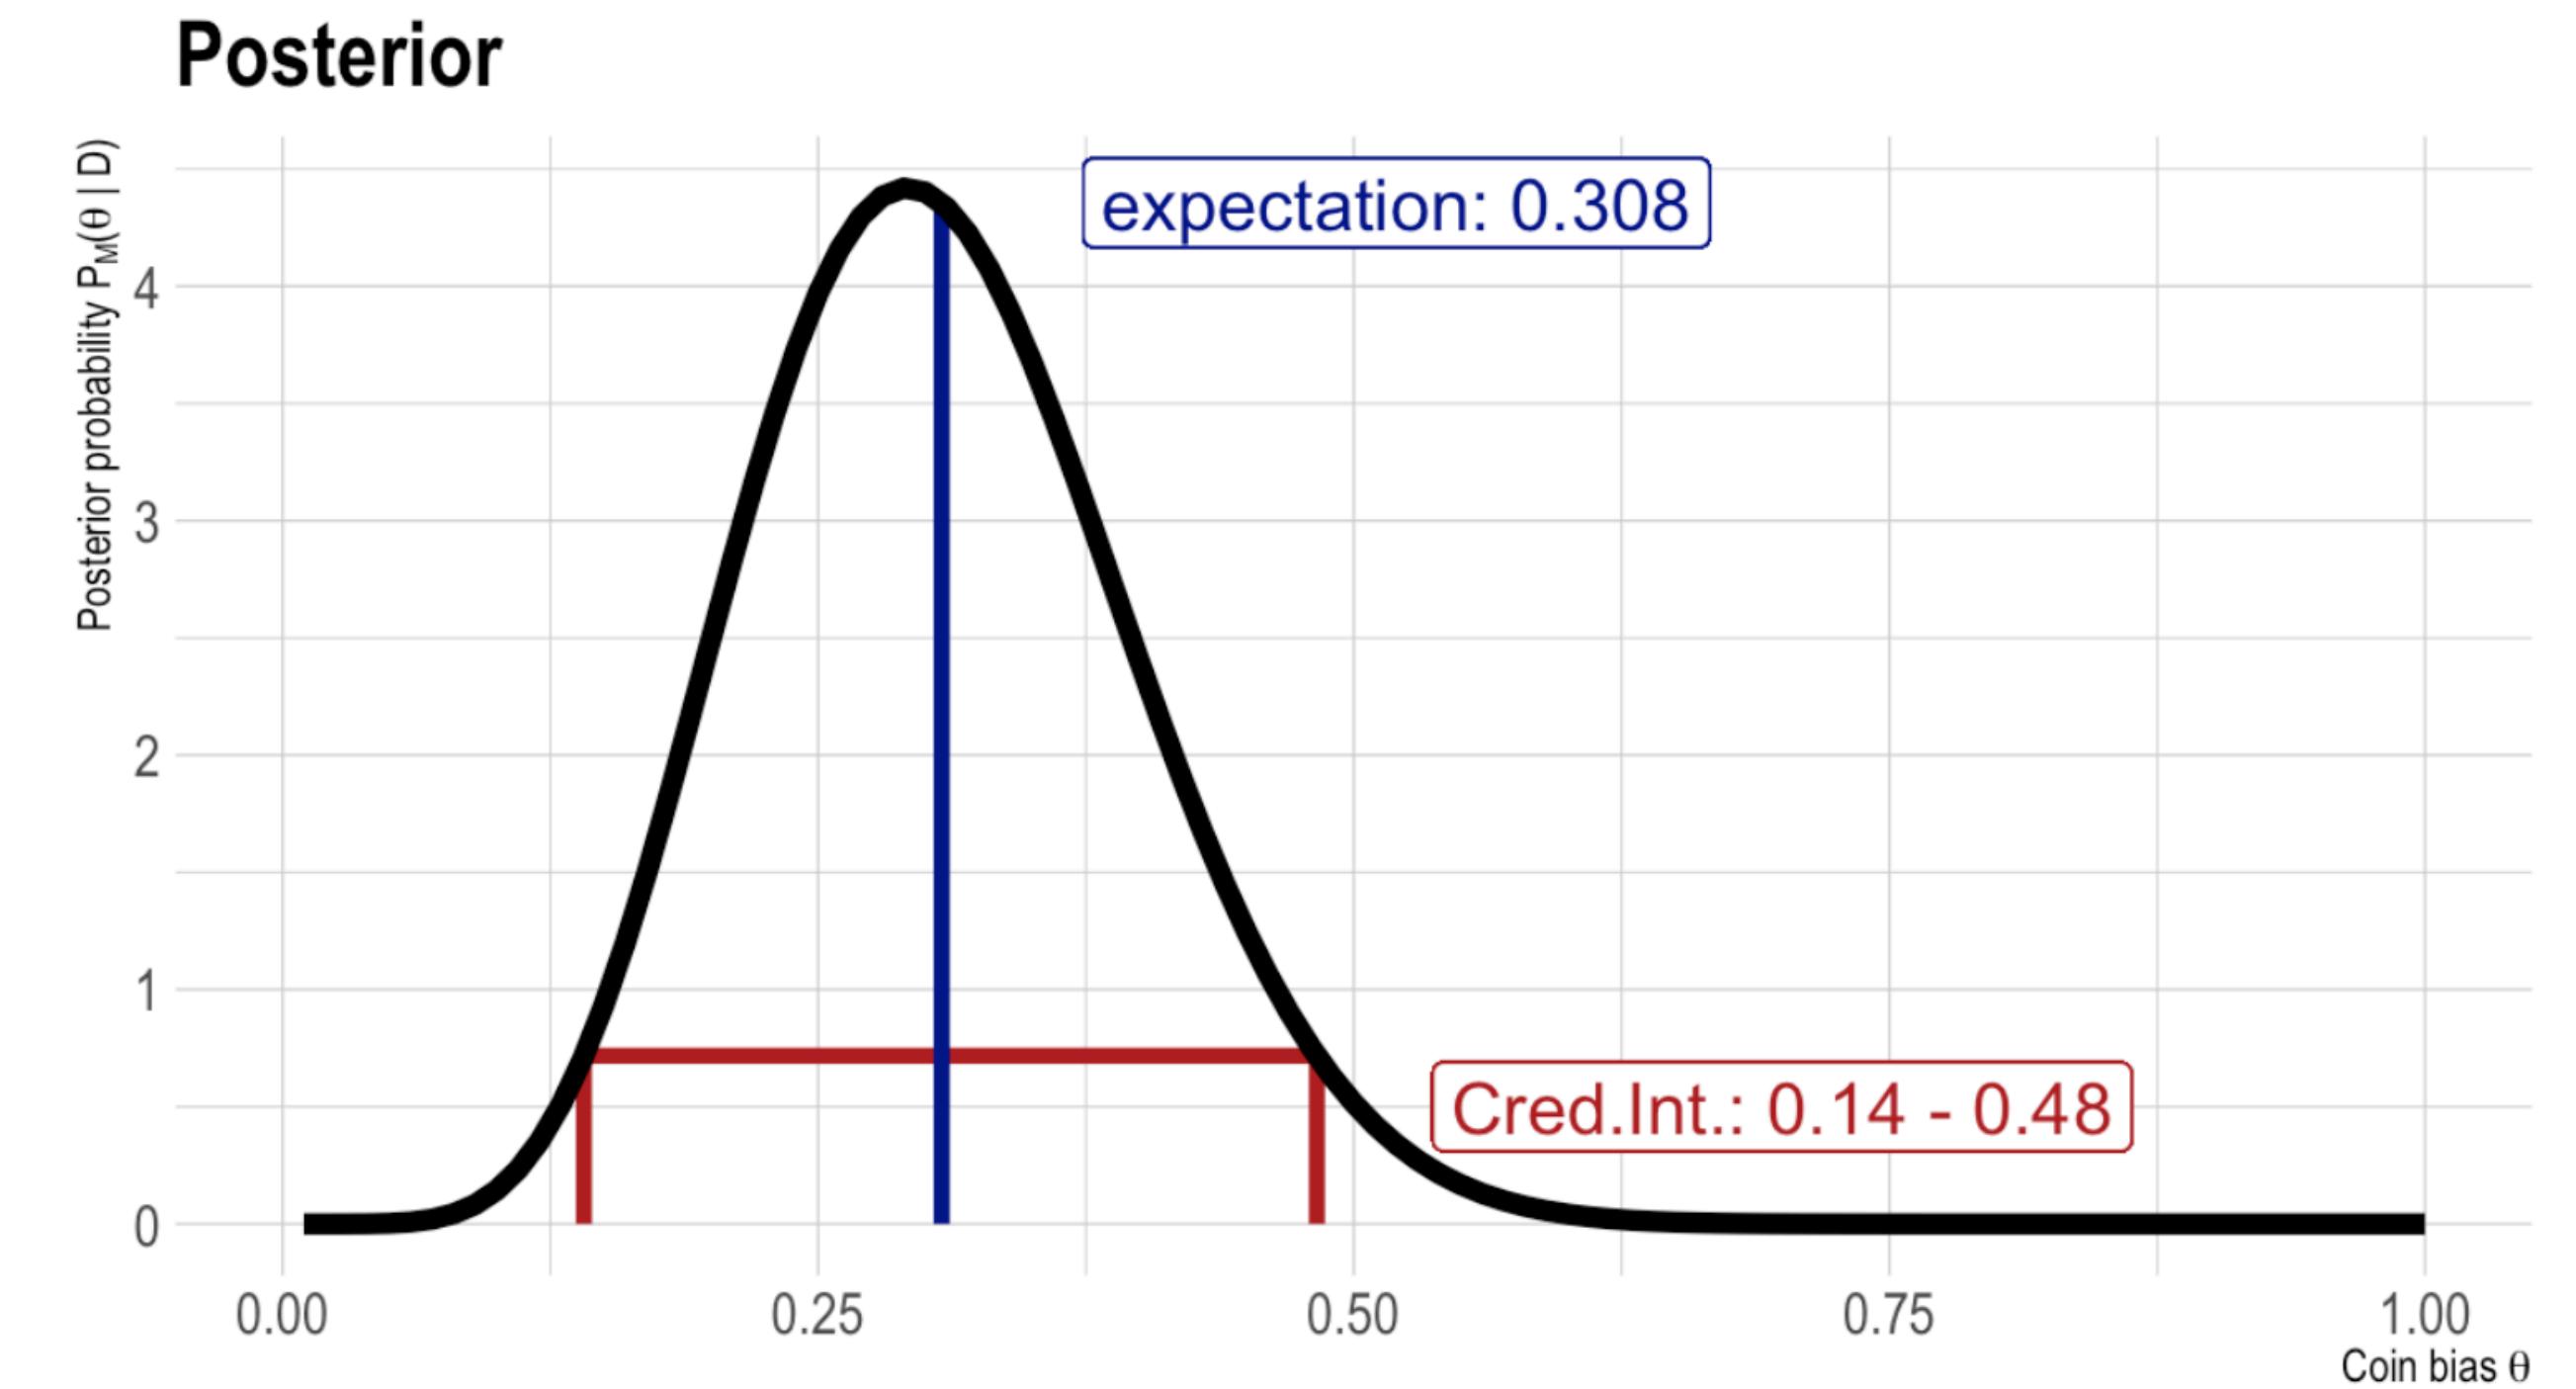
Parameter estimation

point- and interval-valued estimates

- ▶ Bayes' rule for parameter estimation:

$$P(\theta | D) = \frac{P(D | \theta) P(\theta)}{\int P(D | \theta) P(\theta) d\theta}$$

- ▶ common point estimates (“best” values):
 - maximum likelihood estimate (MLE)
 - maximum a posteriori (MAP)
 - **posterior mean / expected value**
- ▶ common interval estimates (range of “good” values):
 - confidence intervals
 - inner quantile ranges
 - **credible intervals**



read more [here](#)

Point-valued estimates

MLE, MAP and (posterior) expected value

► MLE:

$$\arg \max_{\theta} P(D | \theta)$$

- doesn't take prior into account (not Bayesian)
- not necessarily unique

► MAP:

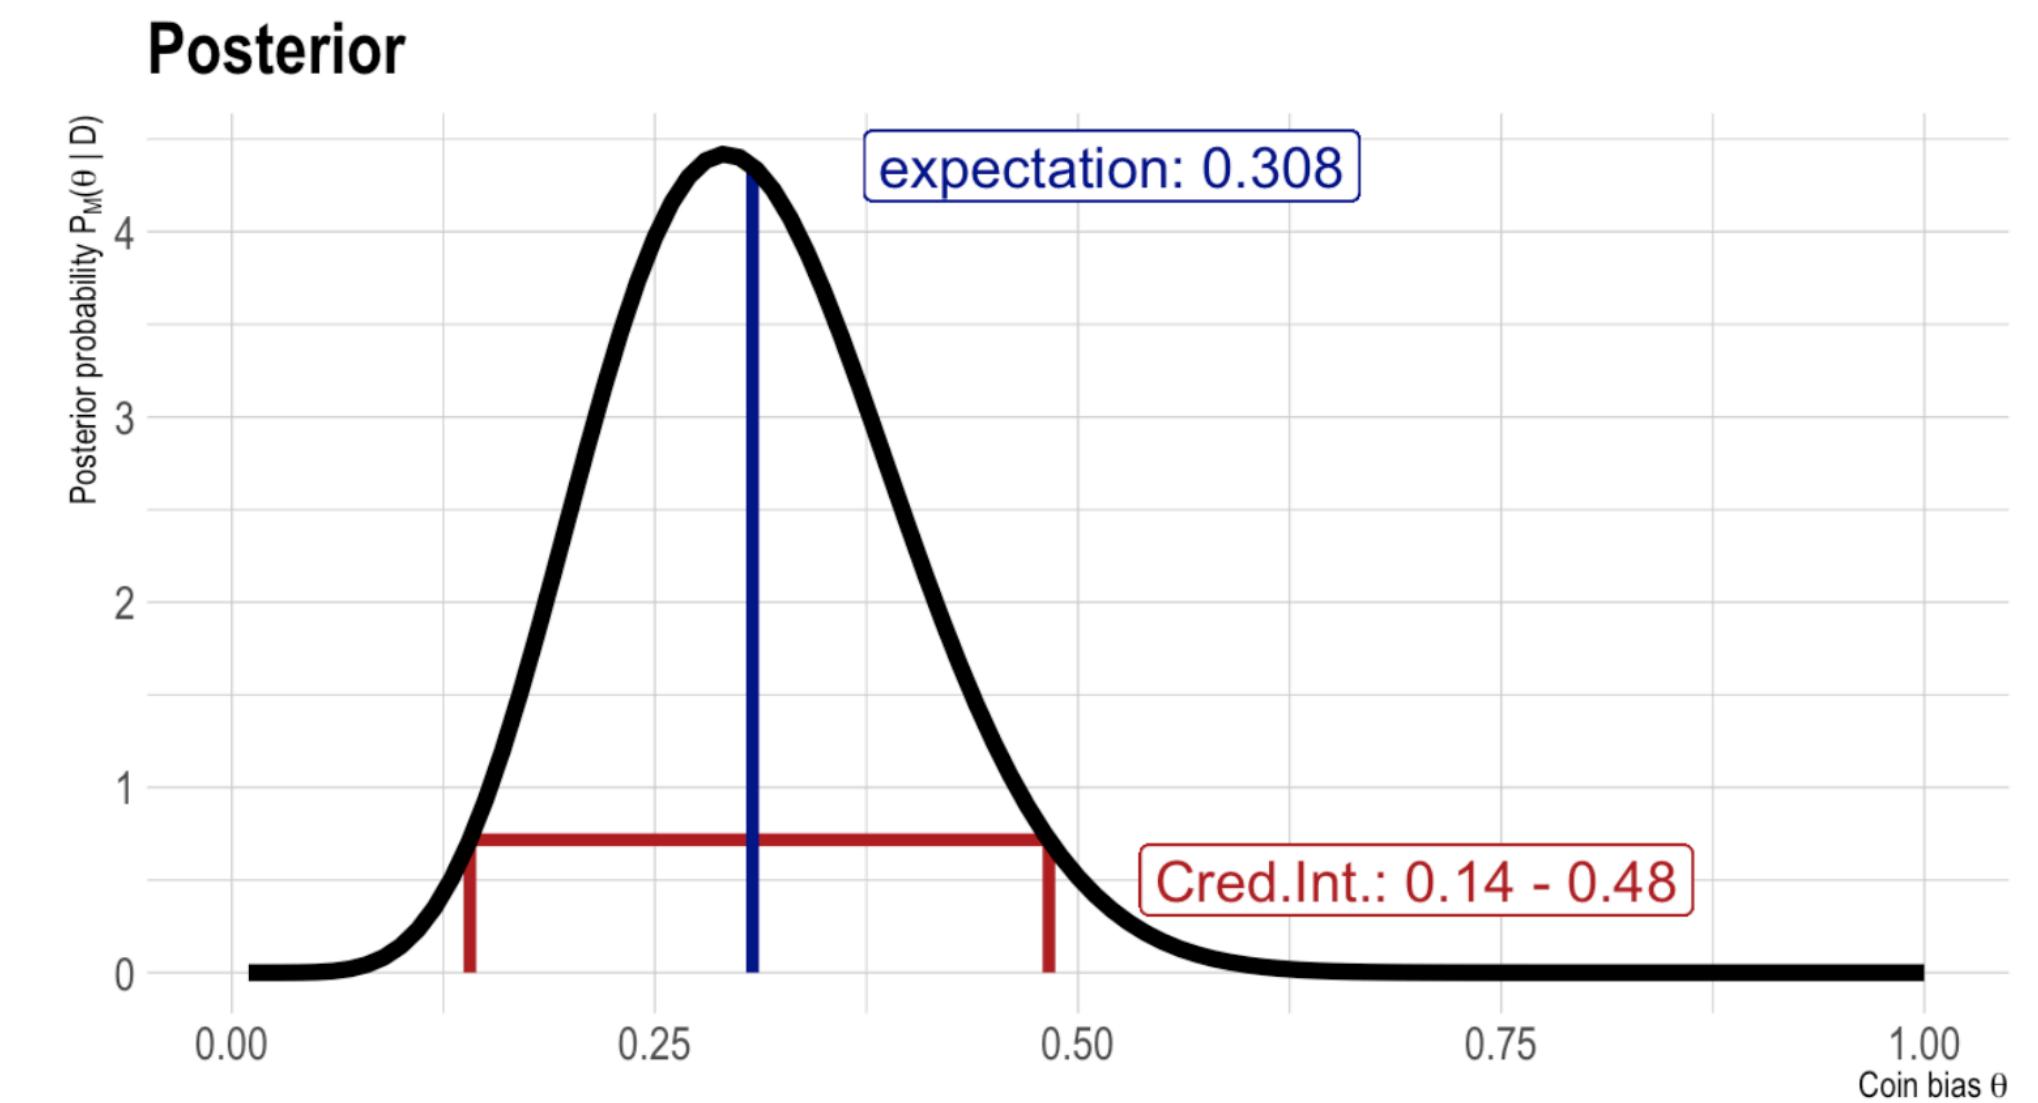
$$\arg \max_{\theta} P(\theta | D)$$

- local / does not consider full distribution (not fully Bayesian)
- increasingly uninformative in larger parameter spaces
- not necessarily unique

► posterior mean / expected valued

$$\mathbb{E}_{P(\theta|D)} = \int \theta P(\theta | D) d\theta$$

- holistic / depends on full distribution ("genuinely Bayesian")
- always unique (for proper priors/posteriors)



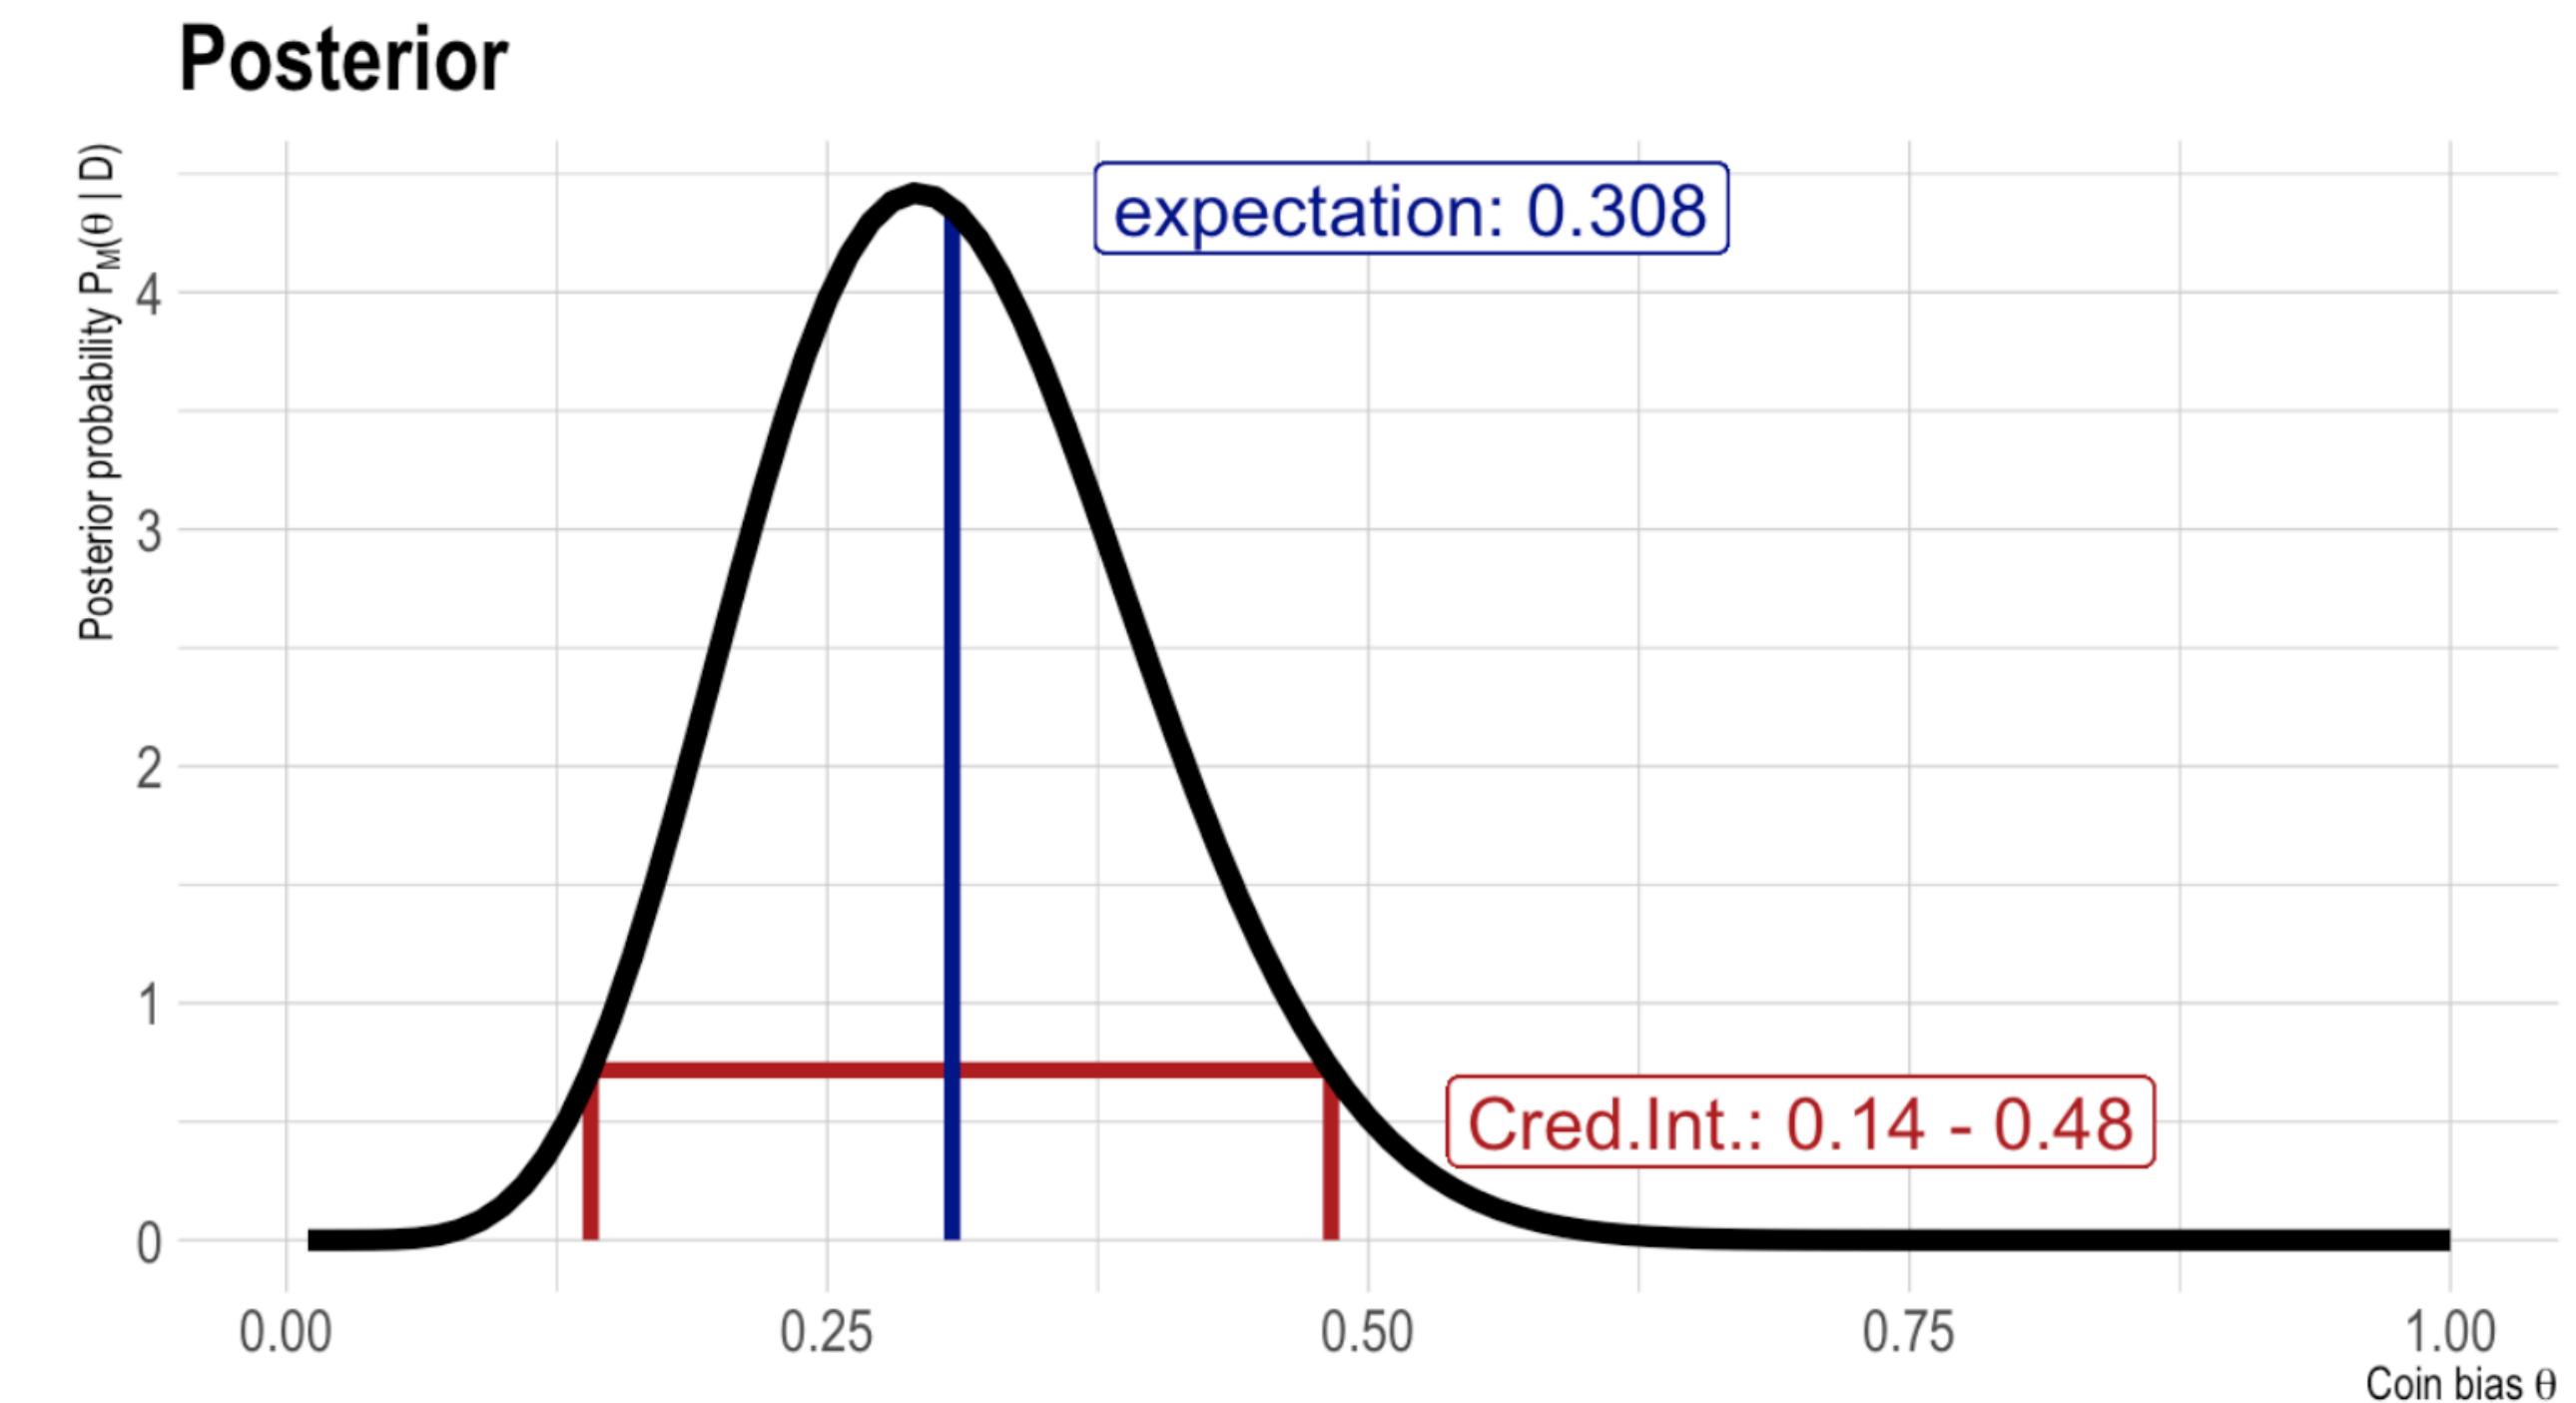
read more [here](#)

Bayesian credible intervals

a.k.a.: highest density interval ...

An interval I is a $\gamma\%$ credible interval, if:

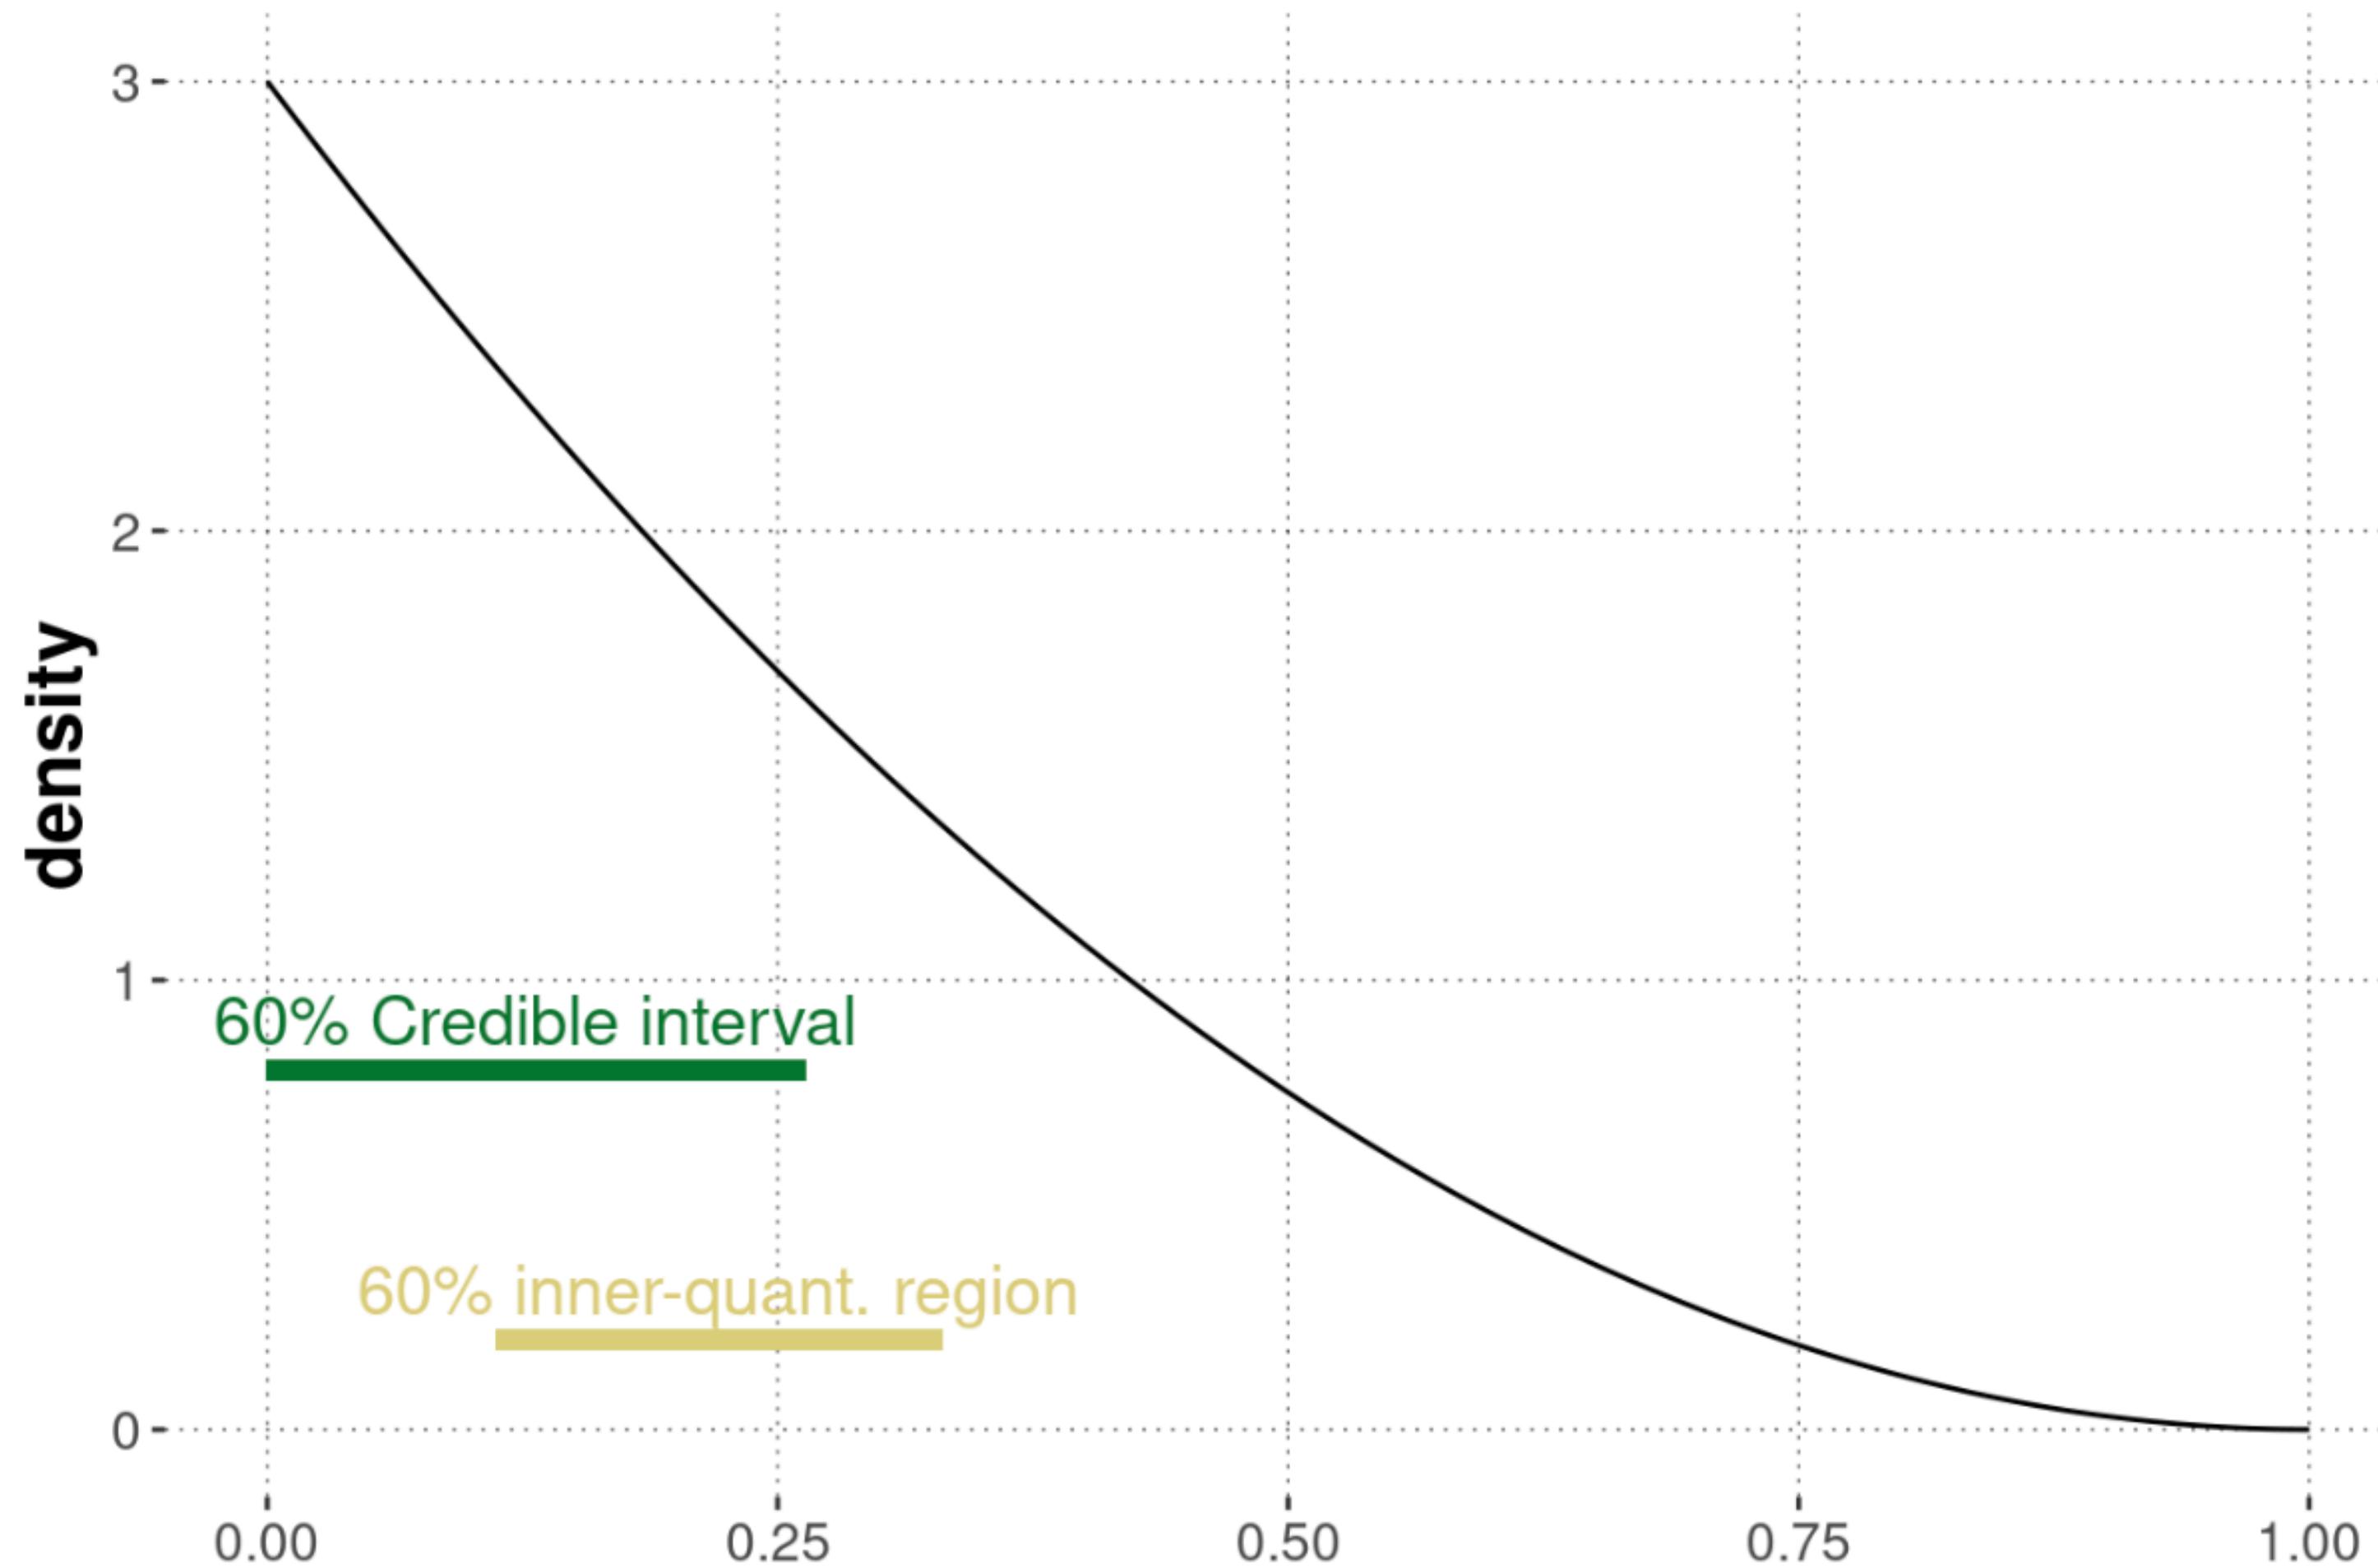
1. the probability that $\theta \in I$ is $\frac{\gamma}{100}$, and
2. no value outside of I is more likely than any point inside of I .



read more [here](#)

Inner quantile regions

!!!are not credible intervals!!!



read more [here](#)

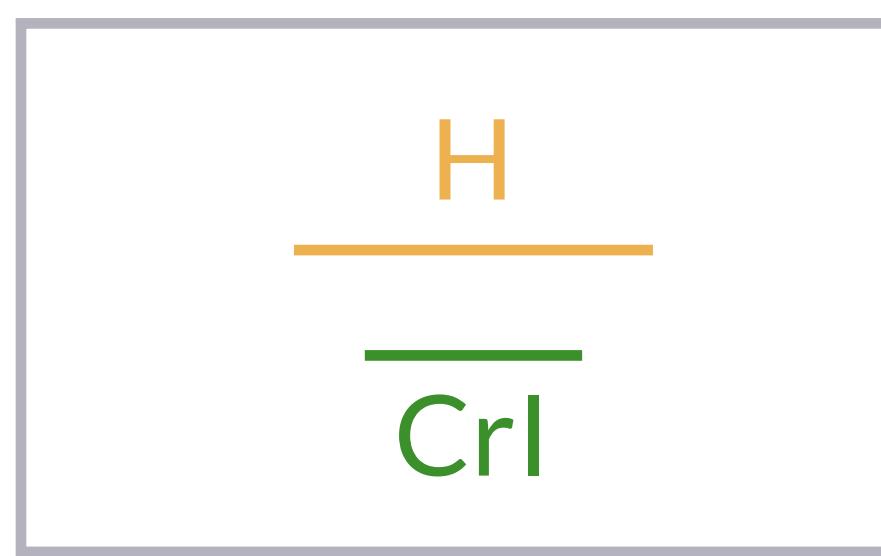


hypotheses &
evidence

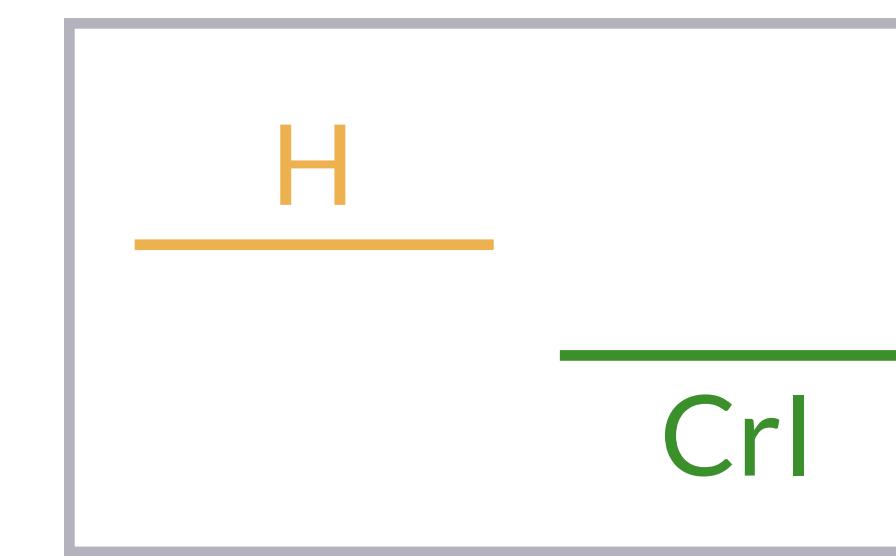
Bayesian hypothesis testing /w posterior credible intervals

!!! caveat: it is controversial whether this is the best (Bayesian) approach to hypothesis testing !!!

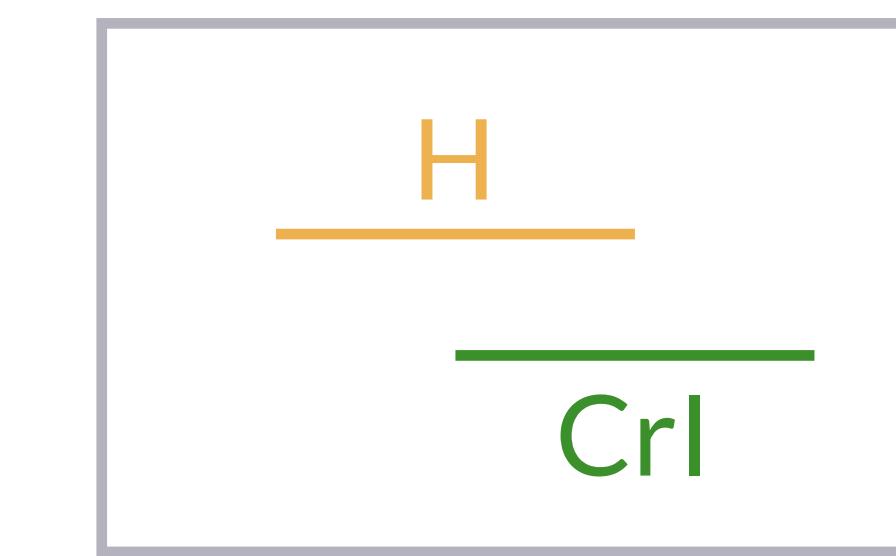
- ▶ consider an interval-based hypothesis: $\theta \in I$
 - e.g., inequality-based: “slope is positive” $\beta_1 = \theta > 0$
 - e.g. a **region of practical equivalence [ROPE]**: an ϵ -region around some θ^* : $I = [\theta^* - \epsilon, \theta^* + \epsilon]$
- ▶ if $[l; u]$ is a posterior credible interval for θ , we consider this:
 - **reason to accept** hypothesis I if $[l; u]$ is contained entirely in I ;
 - **reason to reject** hypothesis I if $[l; u]$ and I have no overlap;
 - **withhold judgement** otherwise.
- ▶ this approach is “categorical” (accept, reject, suspend) and not quantitative



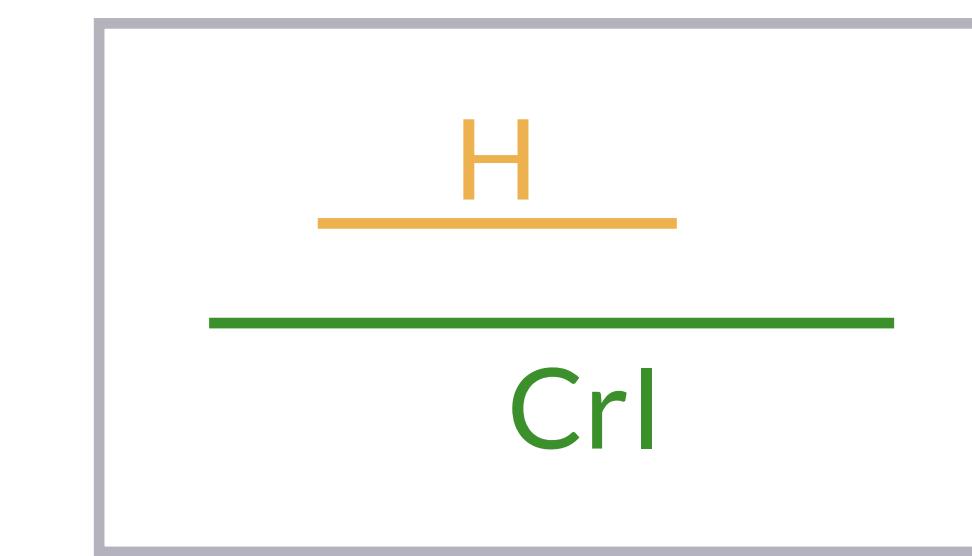
yes



no



maybe



maybe

read more [here](#)

Posterior plausibility of interval-based hypotheses

this is NOT a testing approach, just one way of quantifying support

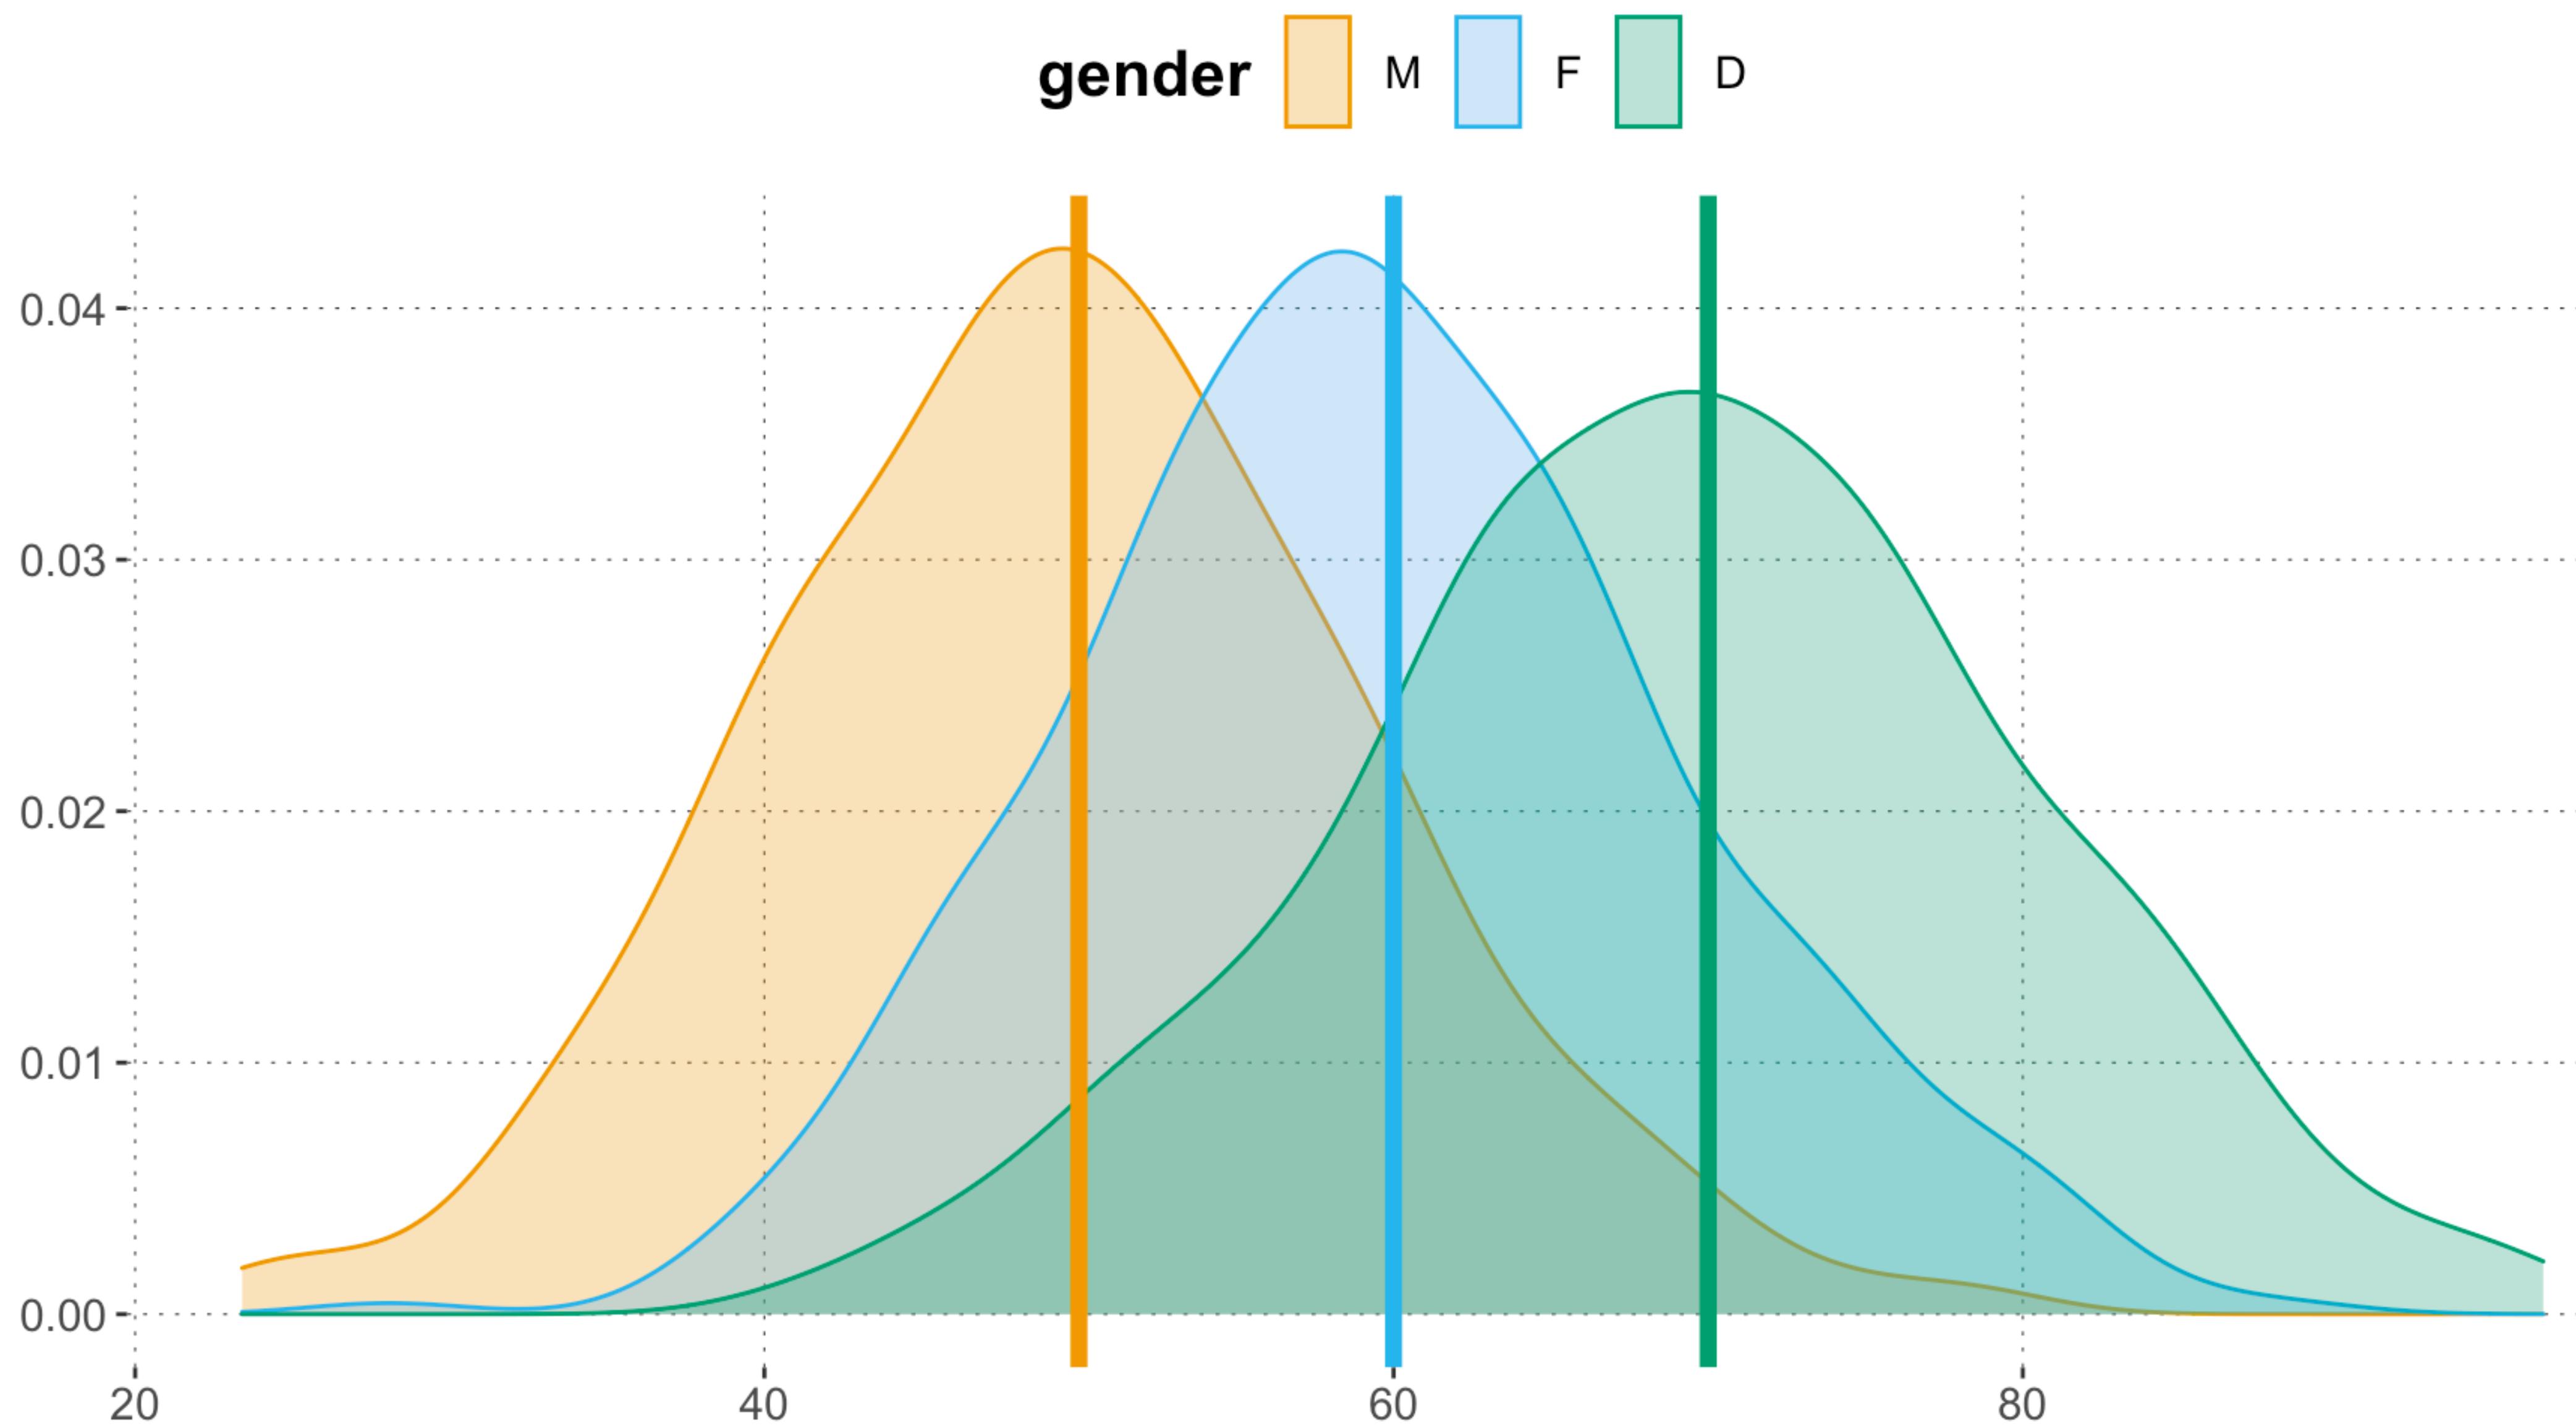
- ▶ consider an interval-based hypothesis $\theta \in I$ as before
- ▶ the **posterior plausibility** of I given a model M and the data D is just the posterior probability: $P(\theta \in I | D)$
- ▶ **not** a notion of observational evidence:
 - if prior is high for I and data is uninformative, posterior plausibility can be high
- ▶ good-enough first heuristic when priors are “unbiased” regarding I
- ▶ more on hypothesis testing later (if you want)



**contrast
coding**

Some fake data

three-way categorical variable



Finding numbers for categories

metric predictors

y	x_0	x_1	x_2
42	1	4	163
19	1	7	128
38	1	2	99
:	:	:	:

categorical predictor

y	gender
51	M
59	F
73	D
:	:

Treatment coding

comparing against a reference category

y	gender	x_0	x_1	x_2
51	M	1	0	0
59	F	1	1	0
73	D	1	0	1
⋮	⋮	⋮	⋮	⋮

$$y_i = \sum_{j=0}^k \beta_j x_{ij} + \epsilon_i$$



$$\hat{\mu}_M = \beta_0 1 + \beta_1 0 + \beta_2 0$$

$$\hat{\mu}_F = \beta_0 1 + \beta_1 1 + \beta_2 0$$

$$\hat{\mu}_D = \beta_0 1 + \beta_1 0 + \beta_2 1$$

hypotheses

$$\beta_0 = \hat{\mu}_M$$

$$\beta_1 = \hat{\mu}_F - \hat{\mu}_M$$

$$\beta_2 = \hat{\mu}_D - \hat{\mu}_M$$



Treatment coding

comparing against a reference category

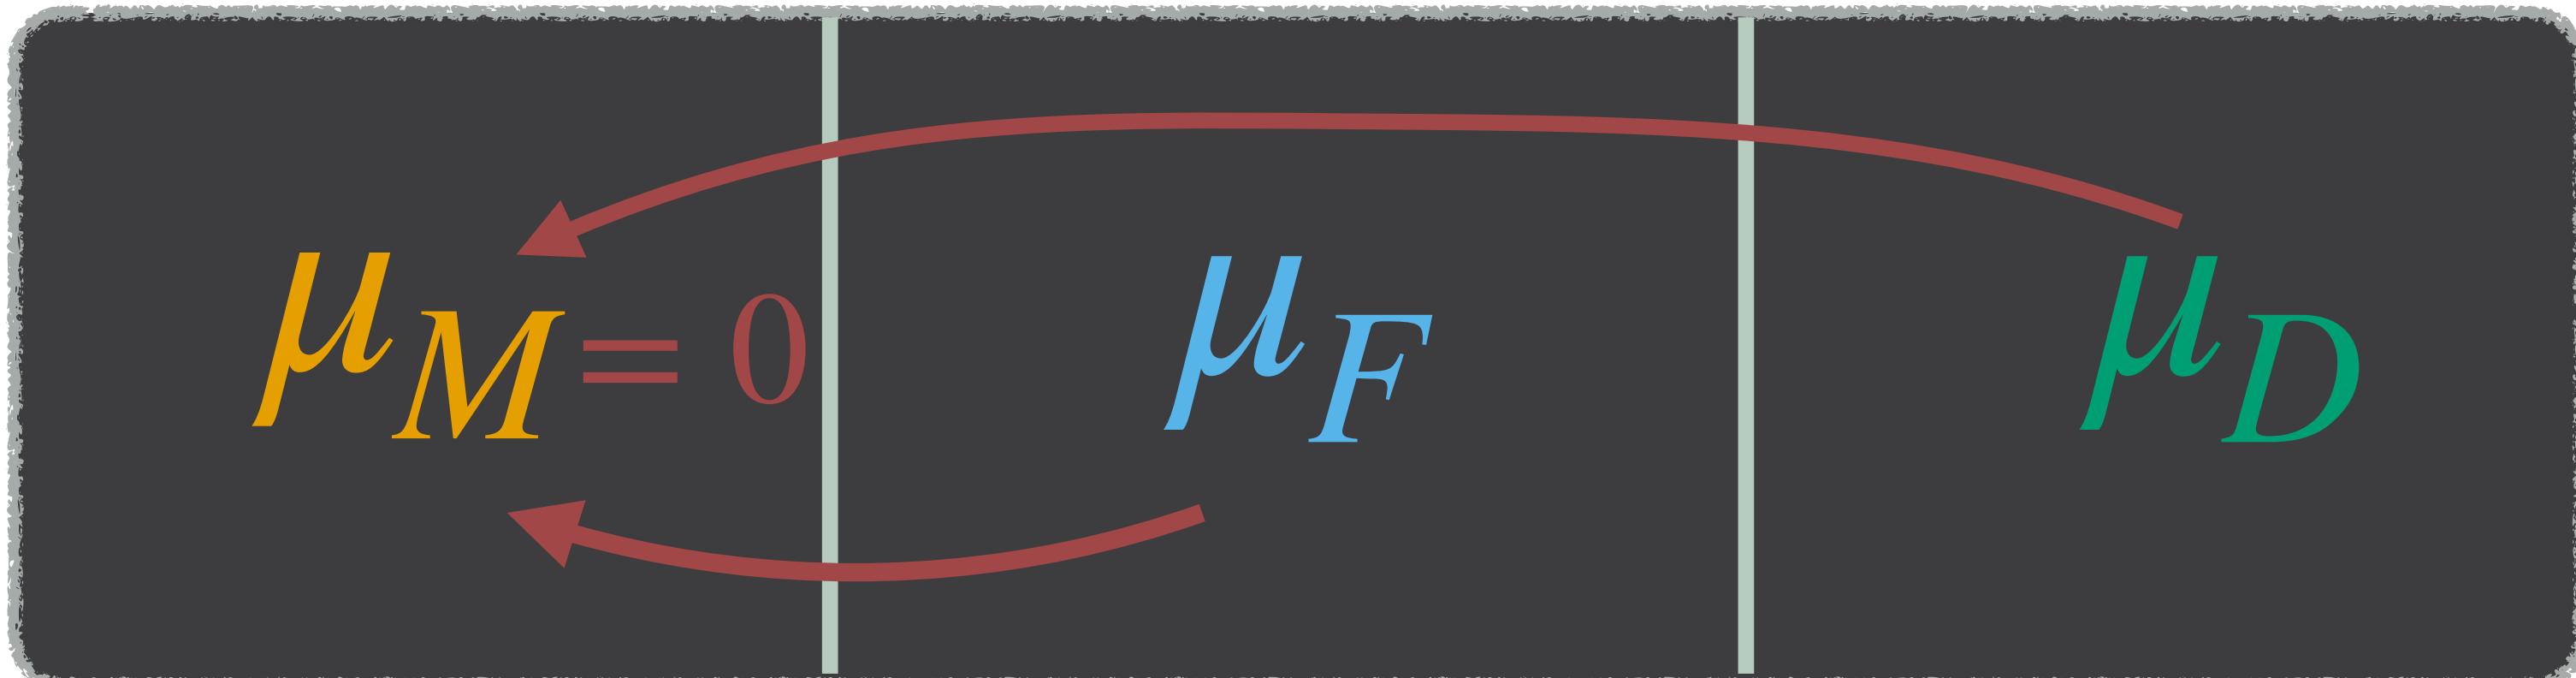
y	gender	x_0	x_1	x_2
51	M	1	0	0
59	F	1	1	0
73	D	1	0	1
:	:	:	:	:

hypotheses

$$\beta_0 = \hat{\mu}_M$$

$$\beta_1 = \hat{\mu}_F - \hat{\mu}_M$$

$$\beta_2 = \hat{\mu}_D - \hat{\mu}_M$$



Cell means coding

estimating a mean for each cell

y	gender	x_0	x_1	x_2
51	M	1	0	0
59	F	0	1	0
73	D	0	0	1
:	:	:	:	:

hypotheses

$$\beta_0 = \hat{\mu}_M$$

$$\beta_1 = \hat{\mu}_F$$

$$\beta_2 = \hat{\mu}_D$$

$$\mu_M = 0$$

$$\mu_F = 0$$

$$\mu_D = 0$$

Simple difference coding

estimating a mean for each cell

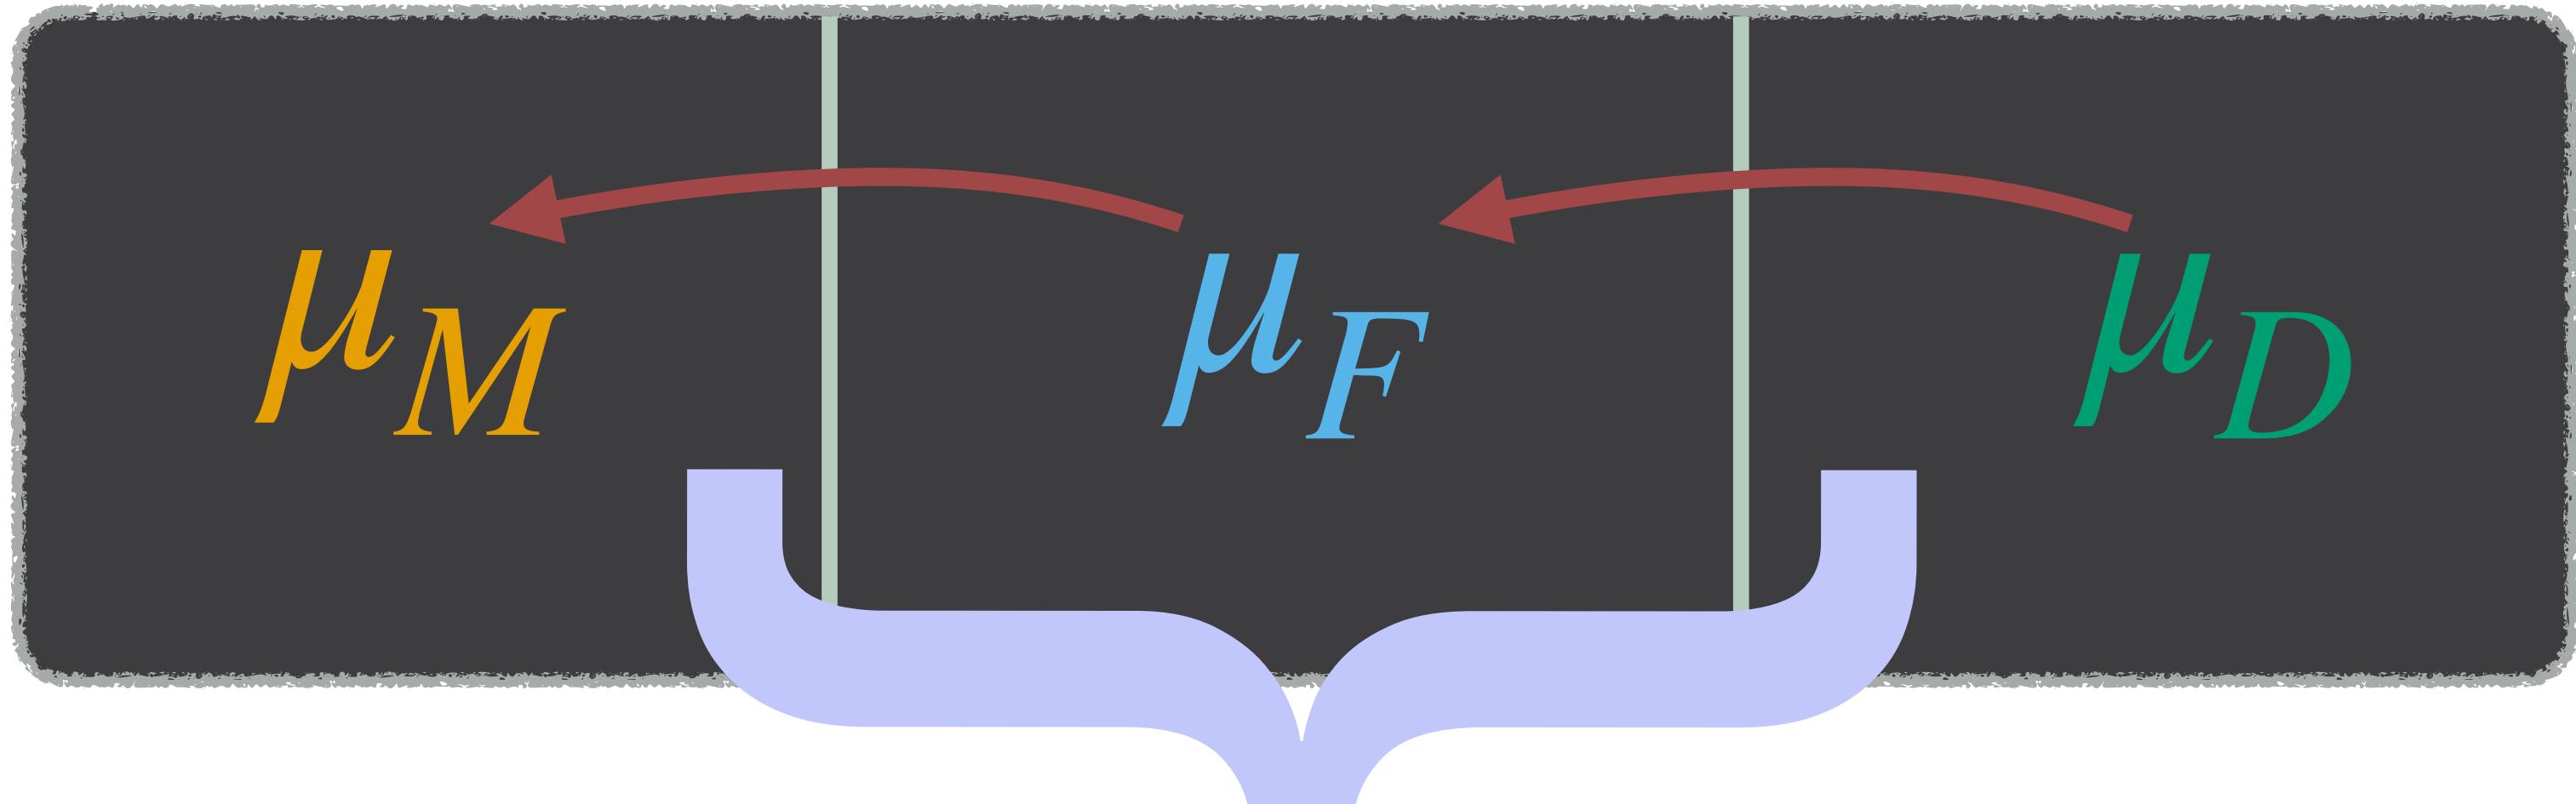
y	gender	x_0	x_1	x_2
51	M	1	- $\frac{2}{3}$	- $\frac{1}{3}$
59	F	1	$\frac{1}{3}$	- $\frac{1}{3}$
73	D	1	$\frac{1}{3}$	$\frac{2}{3}$
:	:	:	:	:

hypotheses

$$\beta_0 = \bar{\mu}$$

$$\beta_1 = \hat{\mu}_F - \hat{\mu}_M$$

$$\beta_2 = \hat{\mu}_D - \hat{\mu}_F$$



$$\bar{\mu} = 0$$

Sum coding

comparing $k - 1$ cells to the grand mean

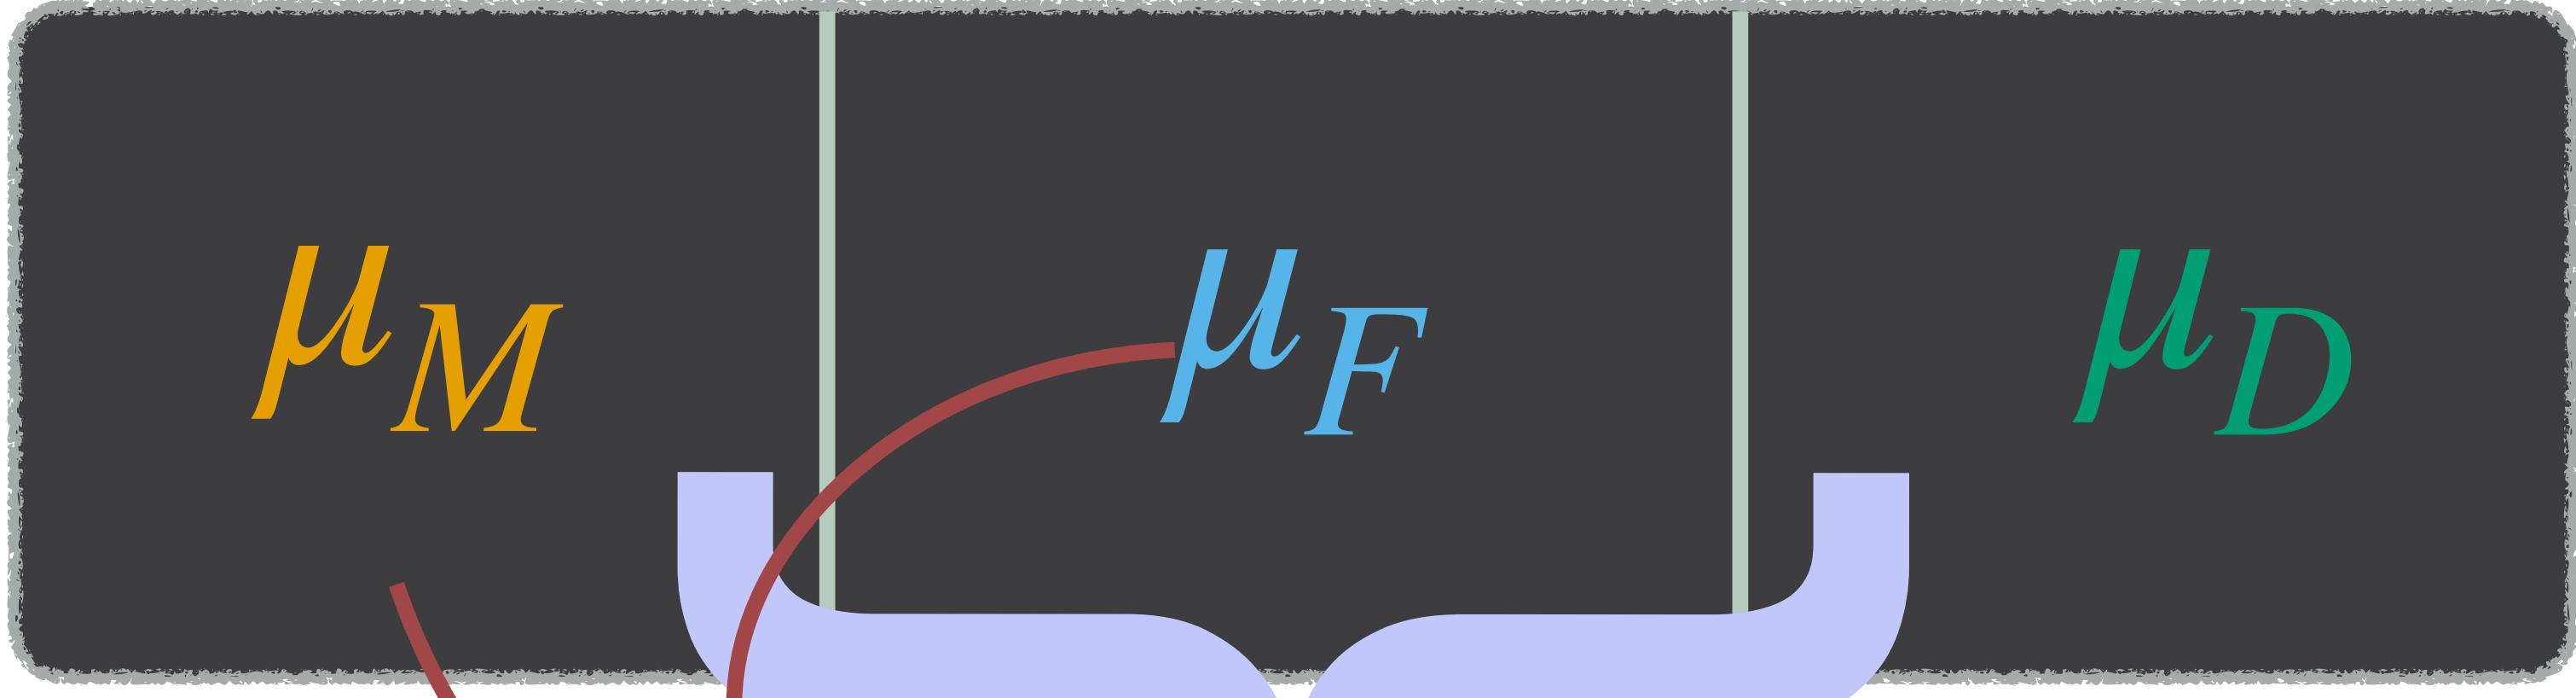
y	gender	x_0	x_1	x_2
51	M	1	1	0
59	F	1	0	1
73	D	1	-1	-1
:	:	:	:	:

hypotheses

$$\beta_0 = \bar{\mu}$$

$$\beta_1 = \hat{\mu}_M - \bar{\mu}$$

$$\beta_2 = \hat{\mu}_F - \bar{\mu}$$



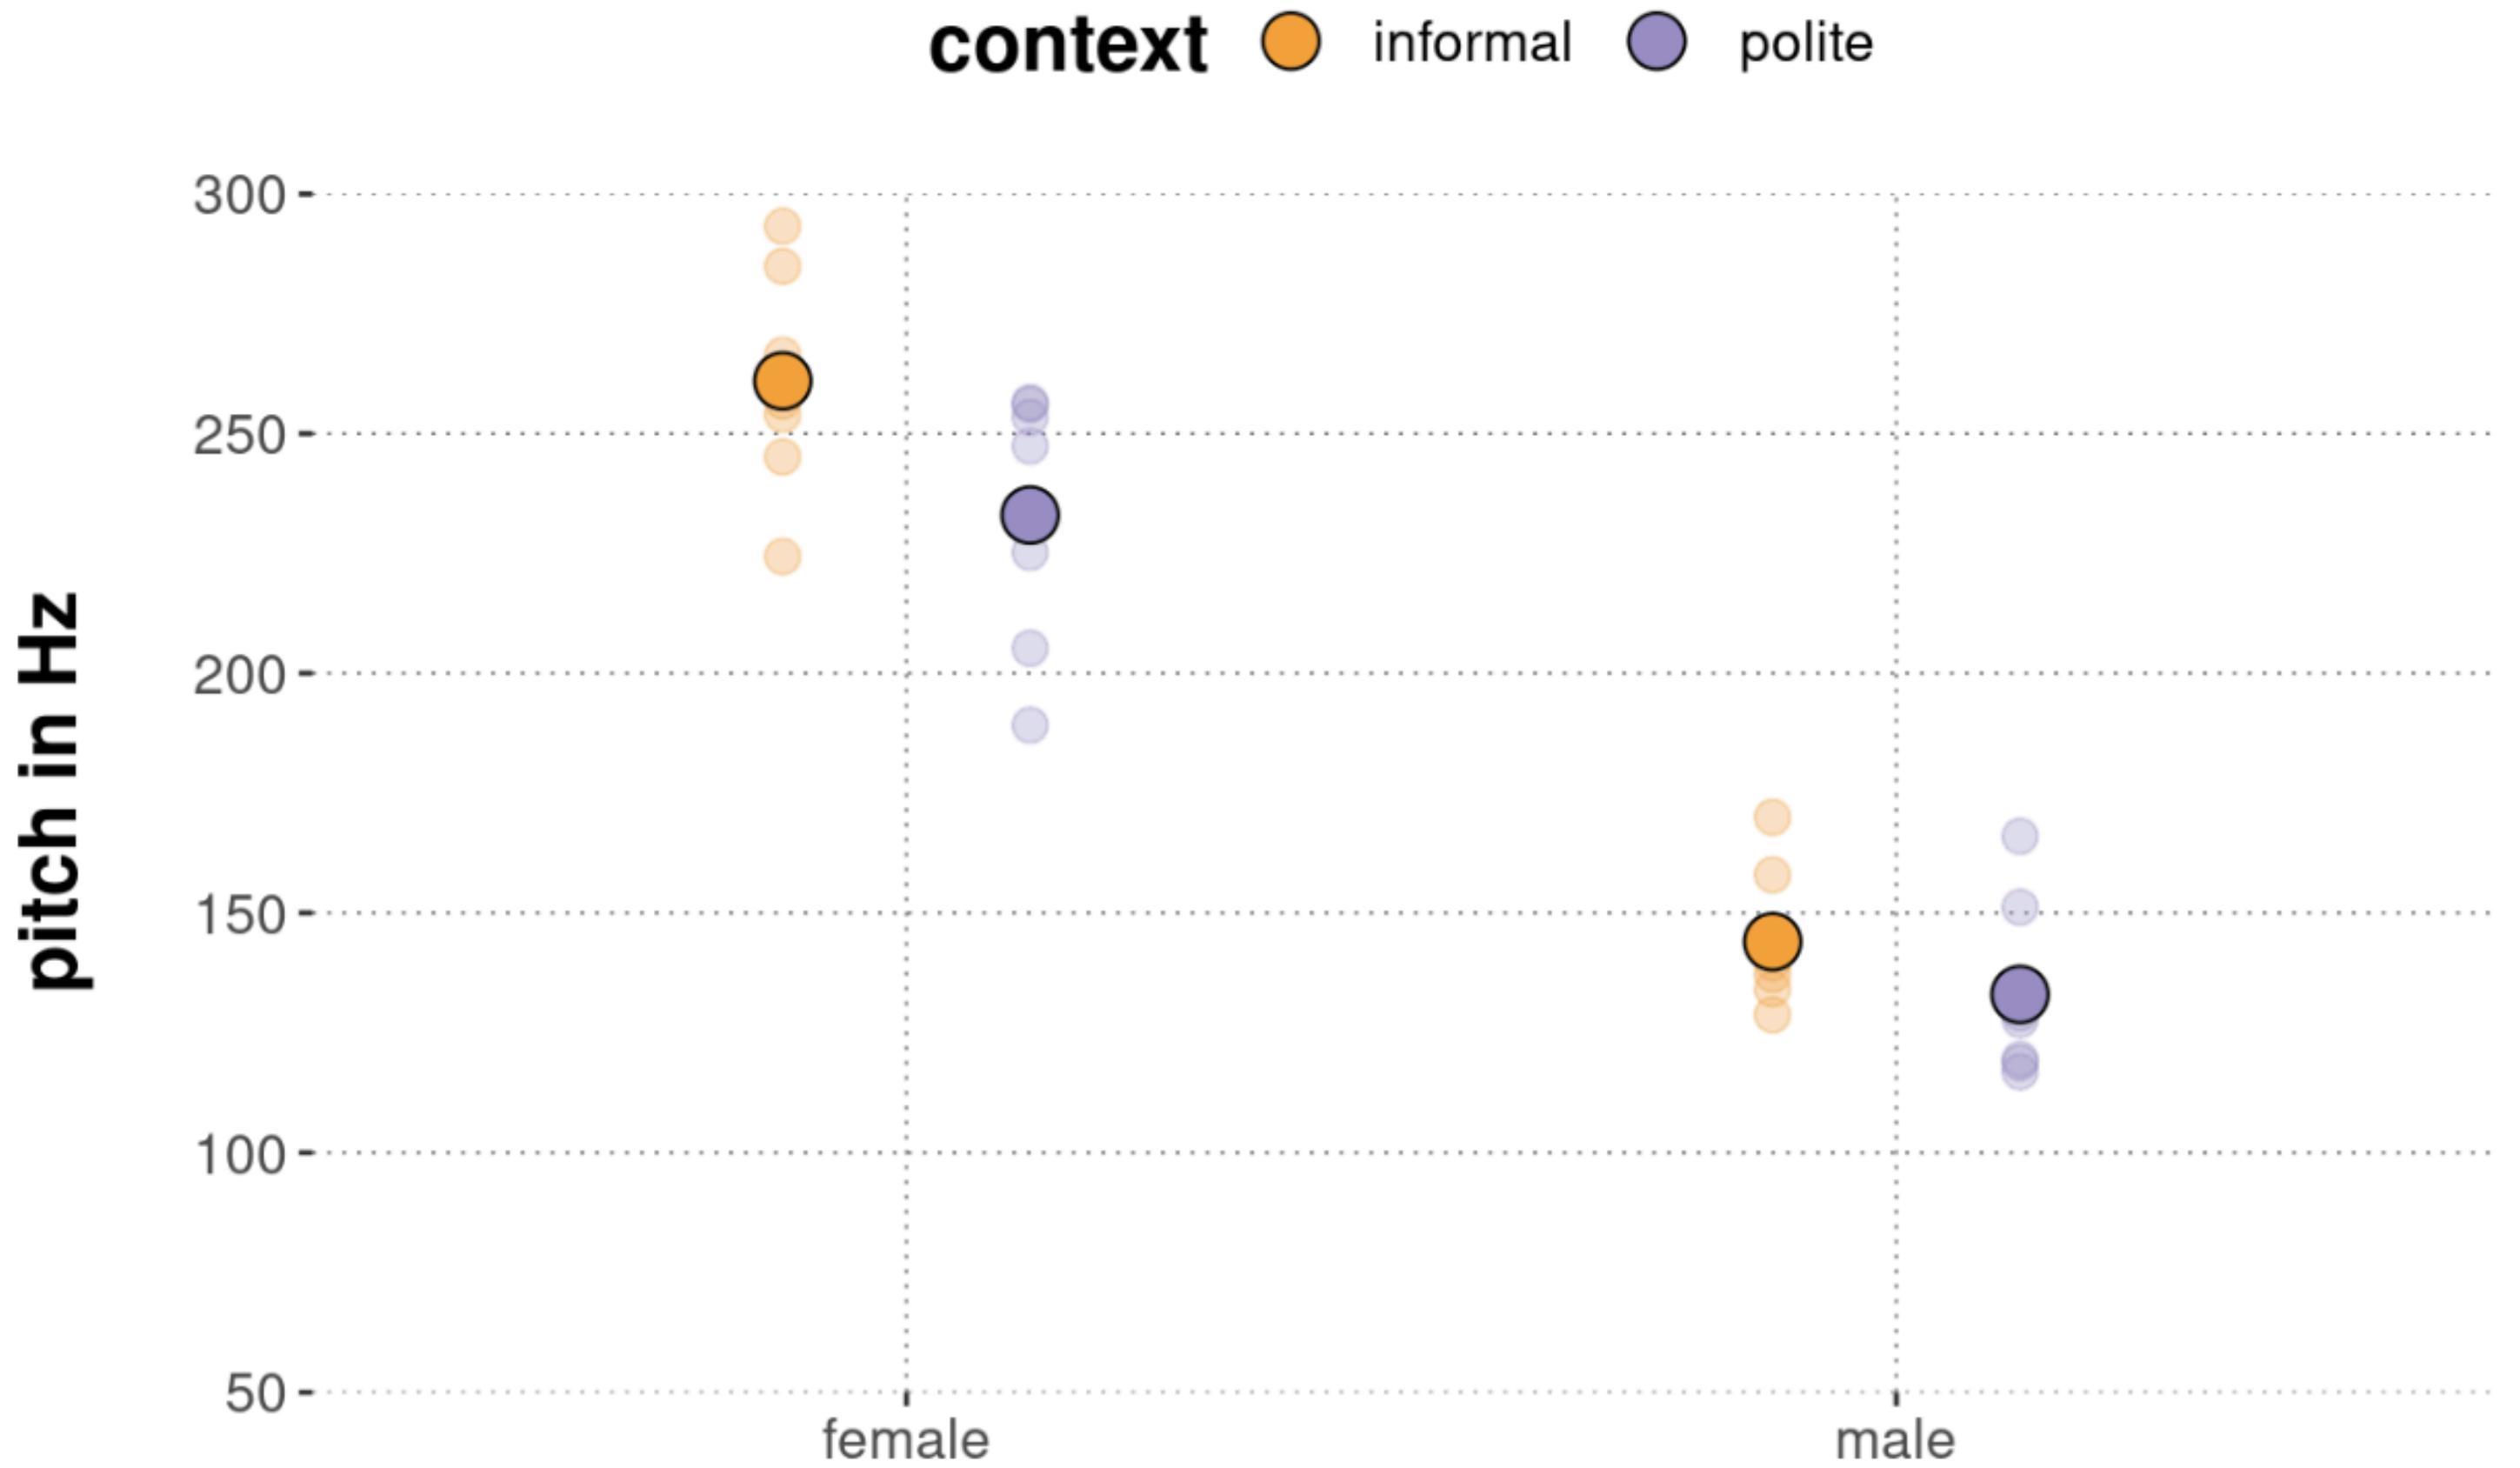
$$\bar{\mu} = 0$$

Case study: pitch in context

data from Winter & Grawunder (2012)

```
politeness_data <- aida::data_polite  
politeness_data %>% head(5)
```

```
## # A tibble: 5 × 5  
##   subject gender sentence context  pitch  
##   <chr>    <chr>   <chr>    <chr>    <dbl>  
## 1 F1       F       S1       pol      213.  
## 2 F1       F       S1       inf      204.  
## 3 F1       F       S2       pol      285.  
## 4 F1       F       S2       inf      260.  
## 5 F1       F       S3       pol      204.
```

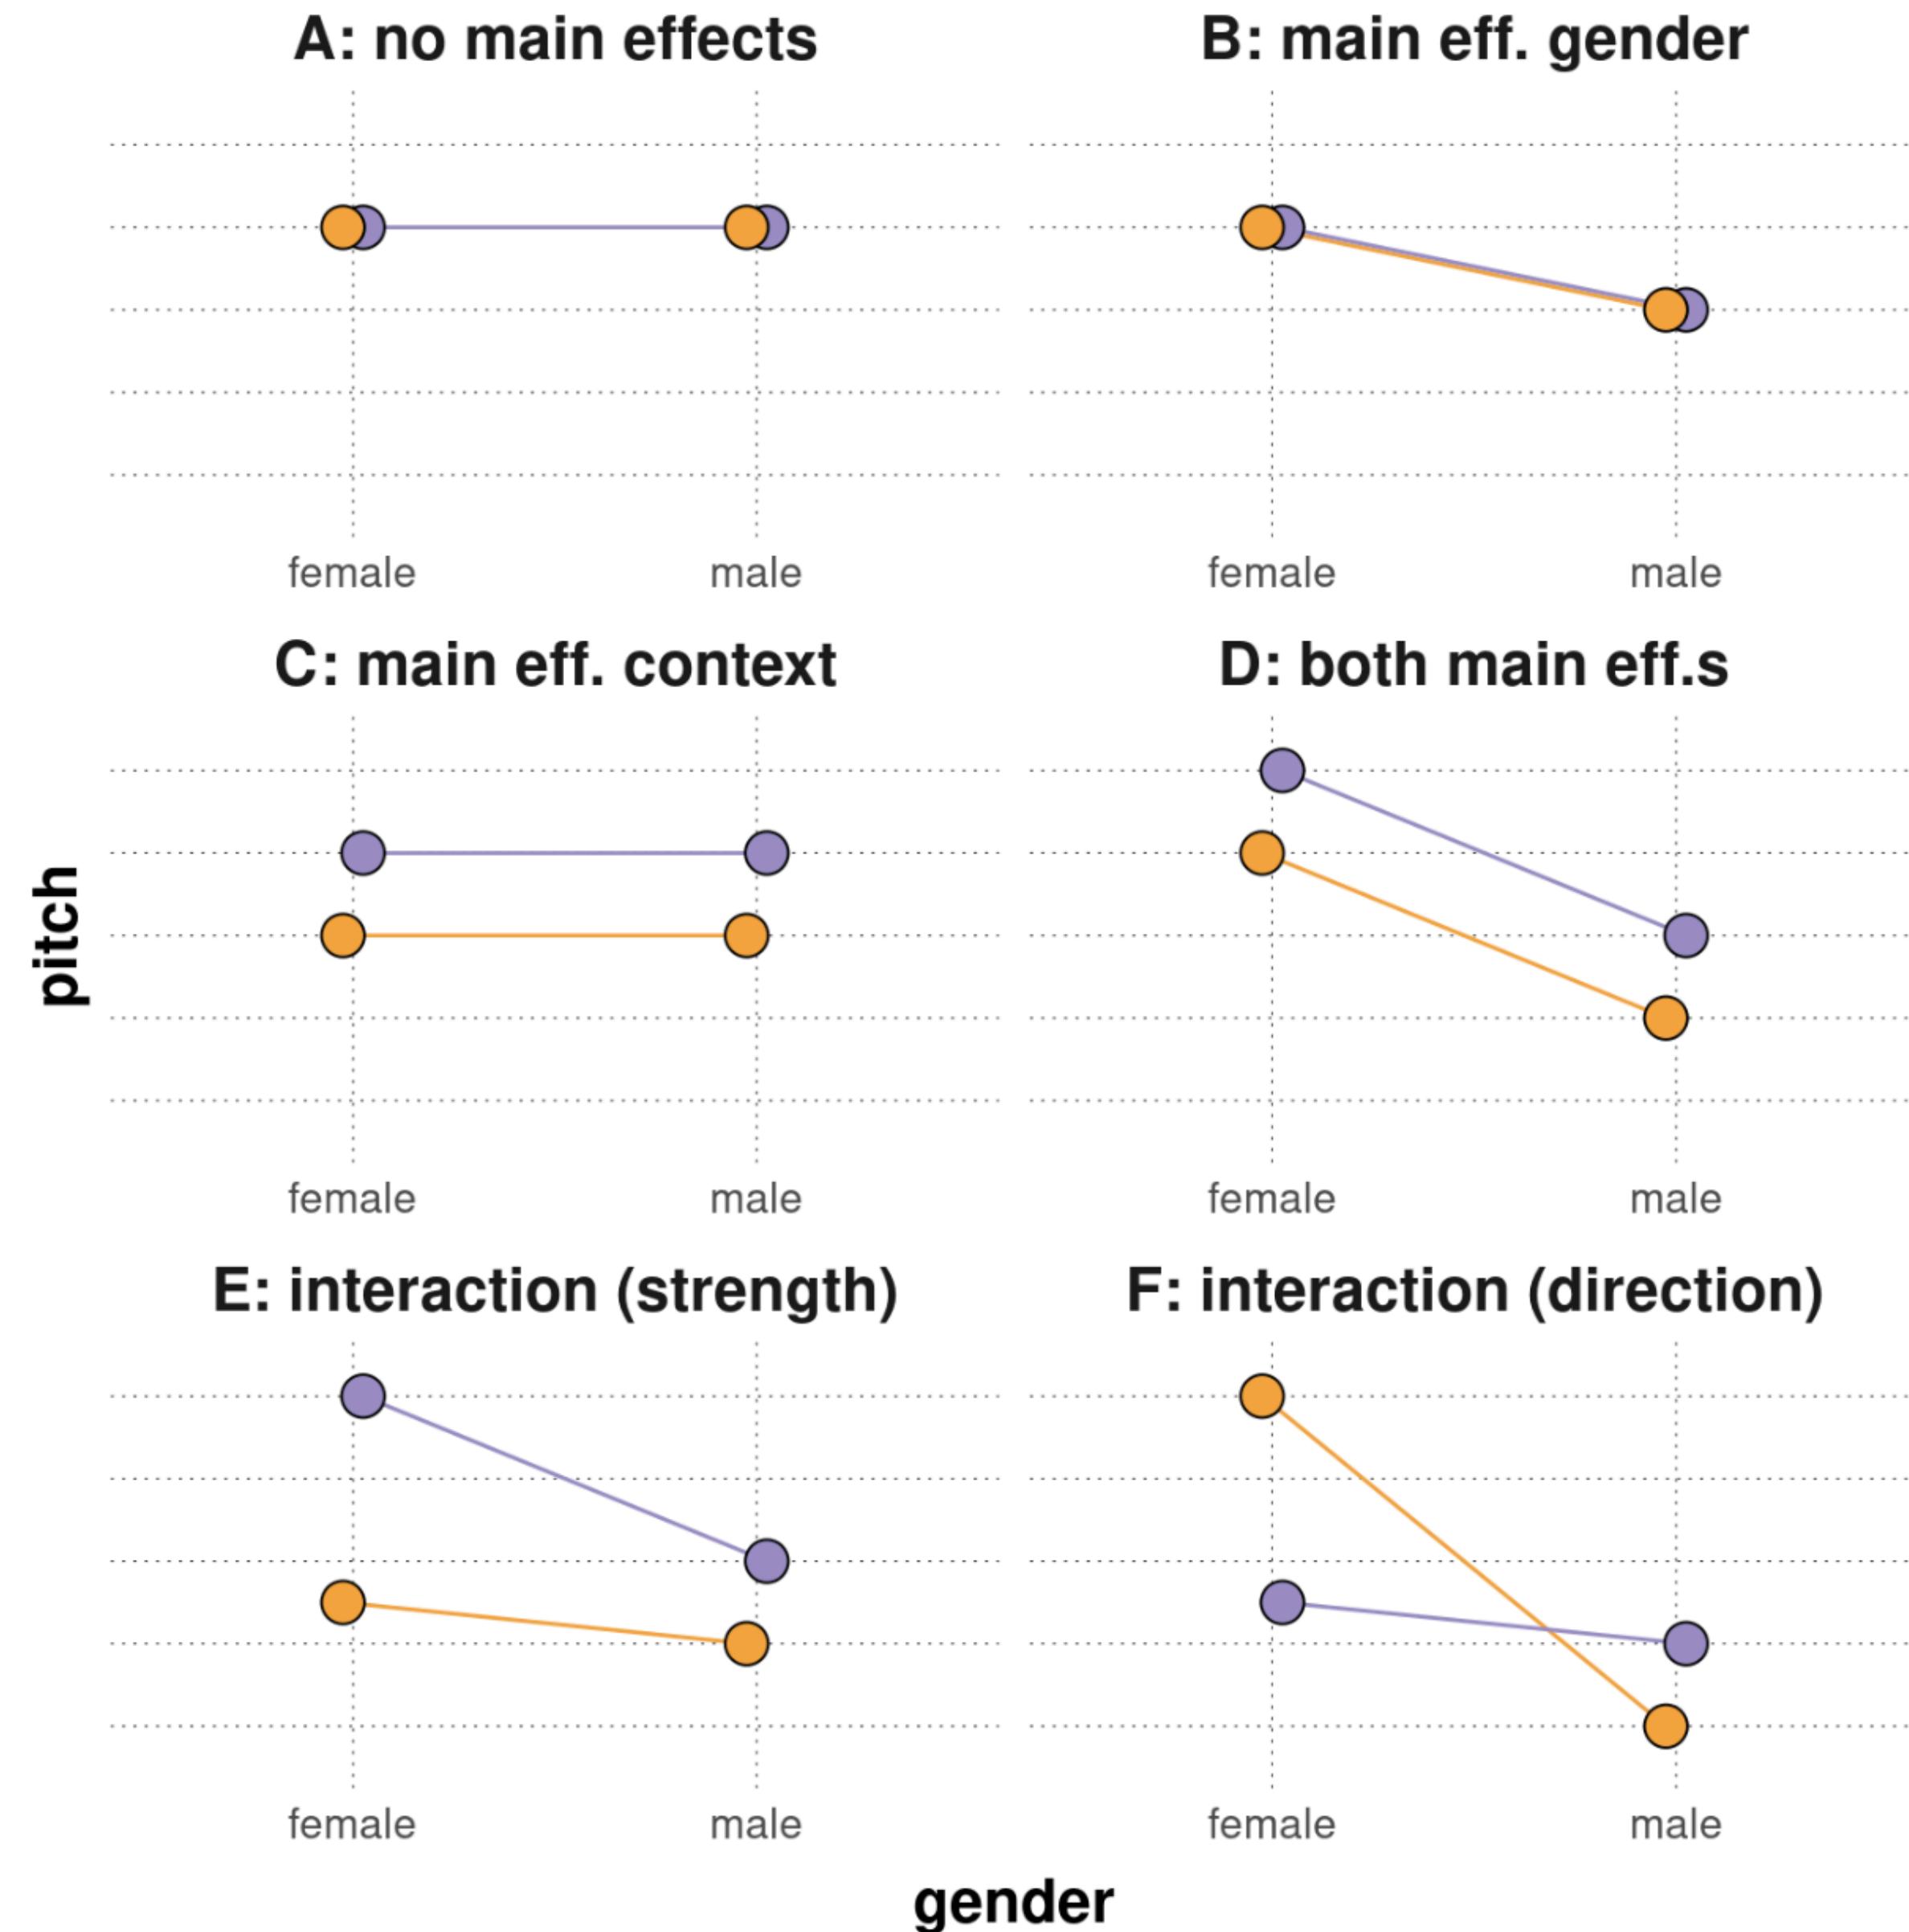


Dummy coding for 2x2 design

	informal	polite
♀	β_0 reference level	$\beta_0 + \beta_{\text{pol}}$
♂	$\beta_0 + \beta_{\text{male}}$	$\beta_0 + \beta_{\text{male}} + \beta_{\text{pol}} + \beta_{\text{pol\&male}}$

Main effects & interactions in 2x2 designs

context ● informal ● polite



Bayesian regression

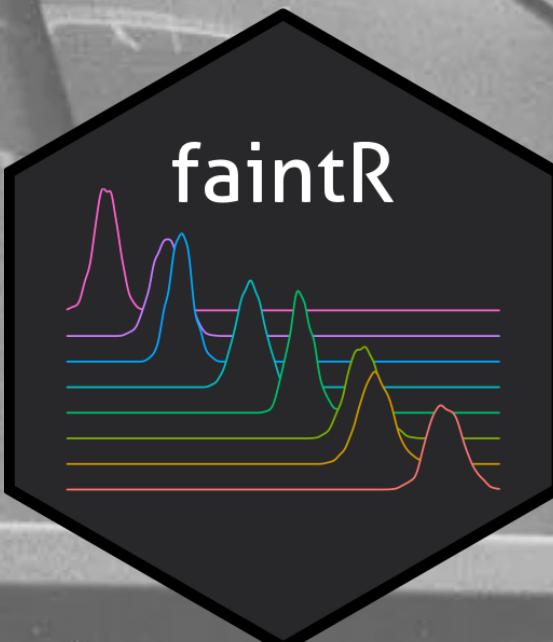
```
# here, we only use fixed effects  
fit_dummy_FE <- brm(  
  pitch ~ gender * context,  
  data = politeness_df,  
  cores = 4,  
  iter = 1000  
)
```

	informal	polite
♀	β_0 reference level	$\beta_0 + \beta_{\text{pol}}$
♂	$\beta_0 + \beta_{\text{male}}$	$\beta_0 + \beta_{\text{male}} + \beta_{\text{pol}} + \beta_{\text{pol\&male}}$

Population-Level Effects:

	Estimate	Est.Error	l-95% CI	u-95% CI
Intercept	260.56	7.87	244.15	275.21
genderM	-116.16	11.01	-137.31	-94.05
contextpol	-27.23	11.10	-48.38	-5.23
genderM:contextpol	15.77	16.05	-16.54	46.24

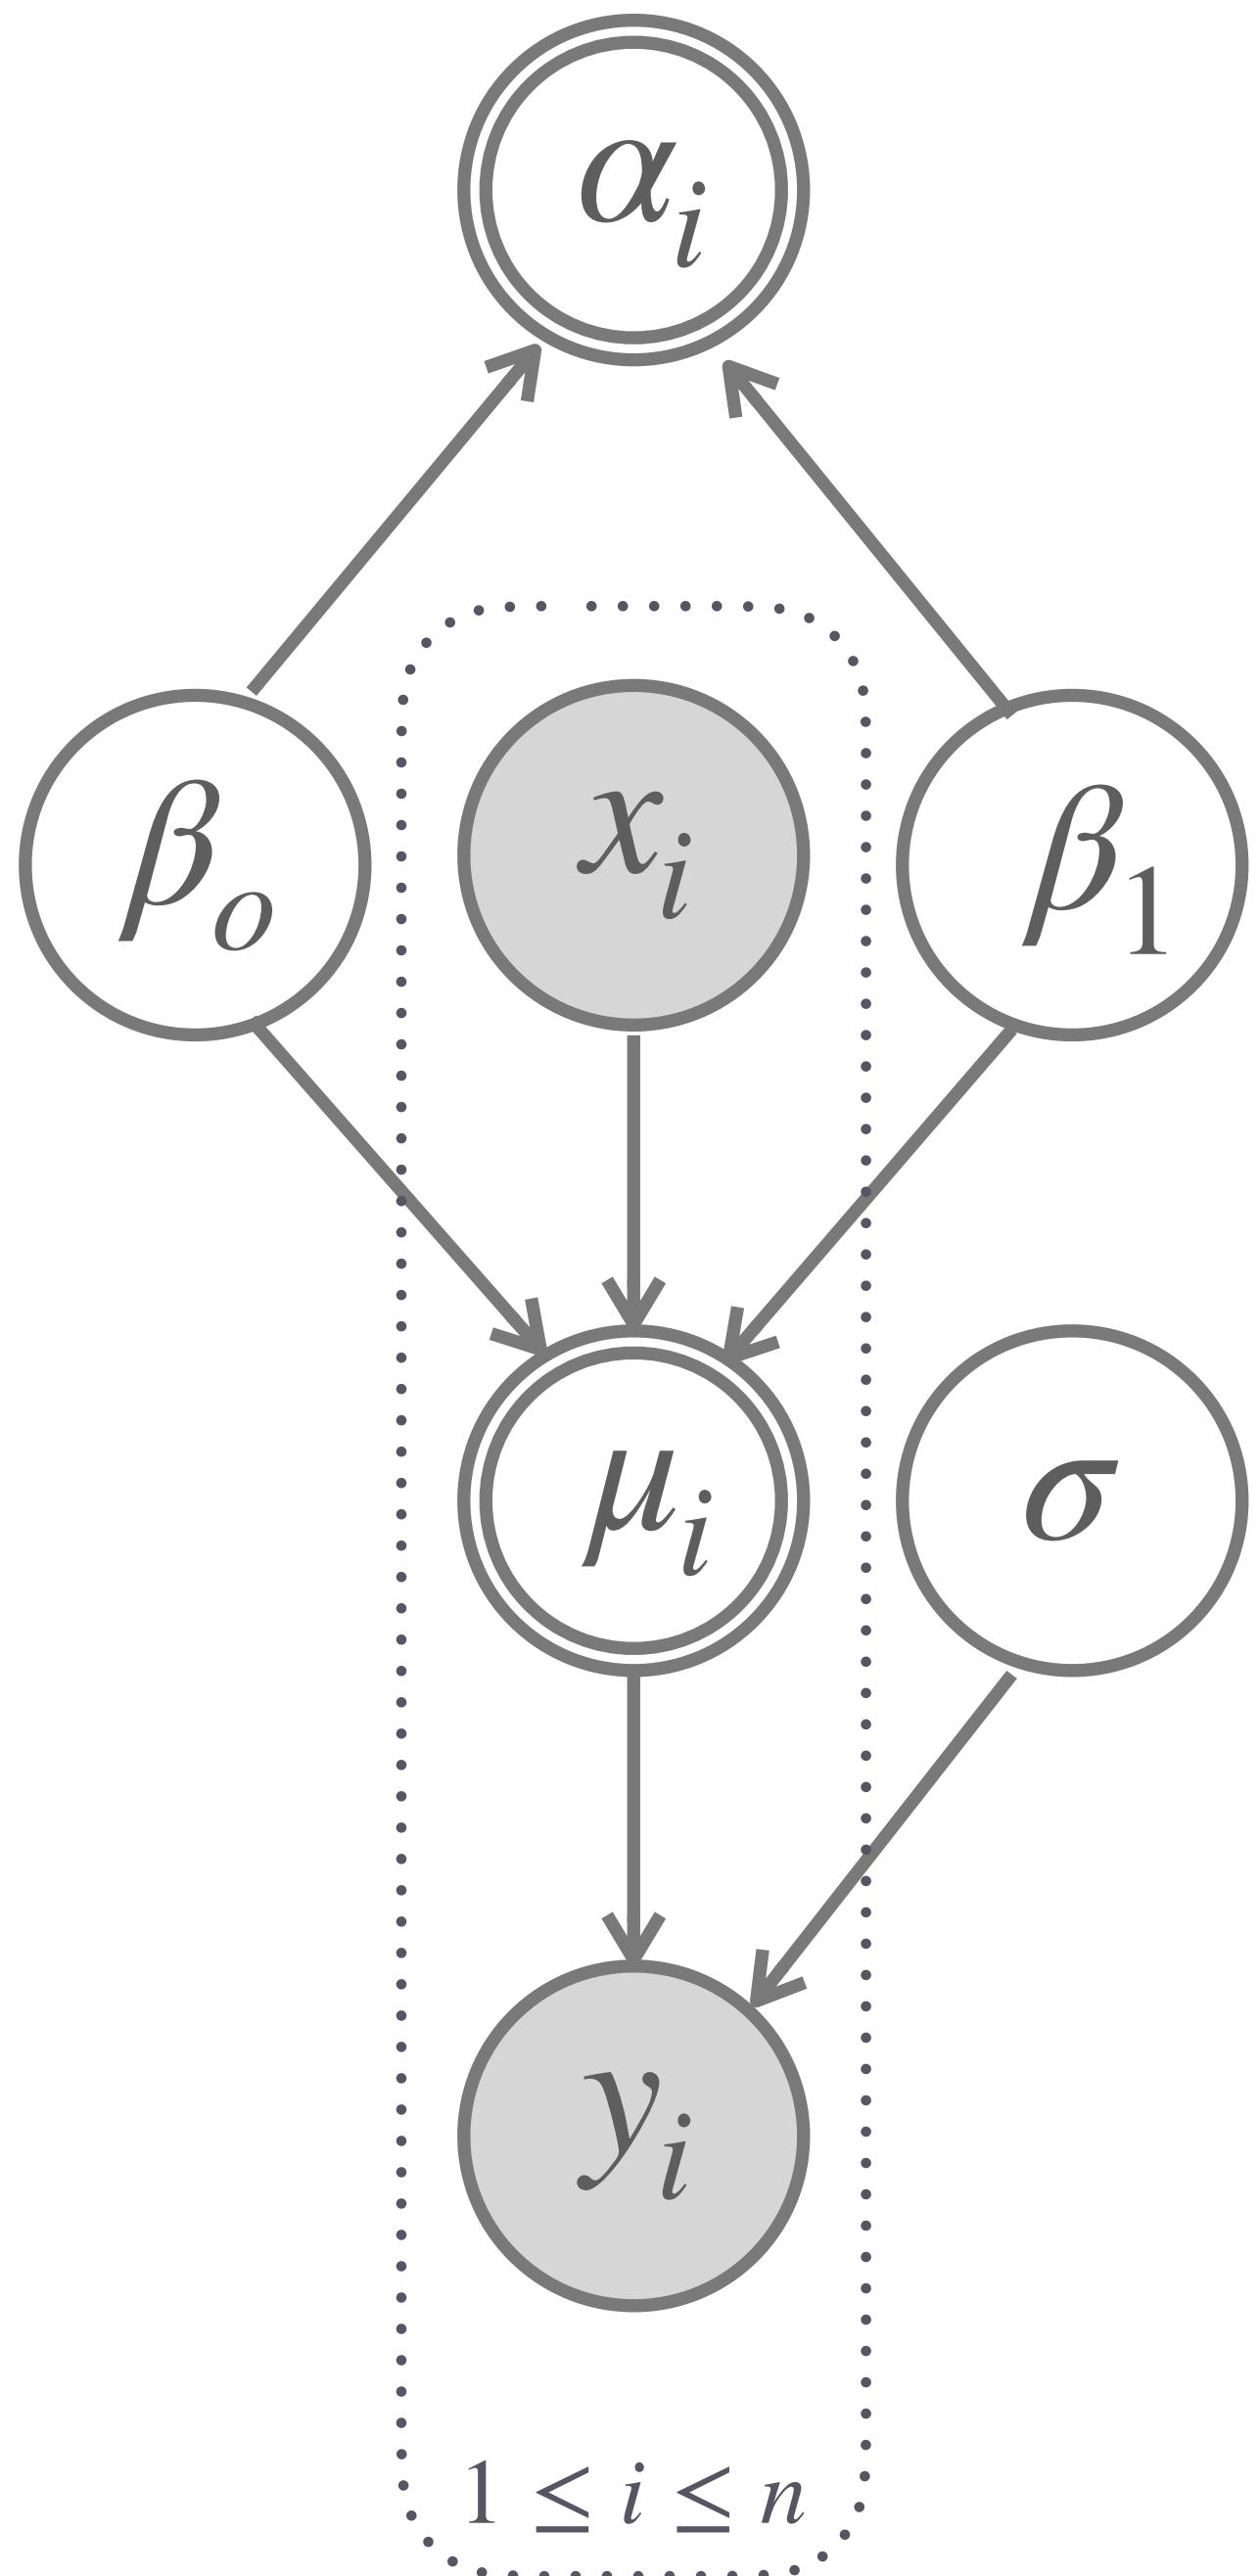
Which questions about cell mean differences can we address with this information directly?

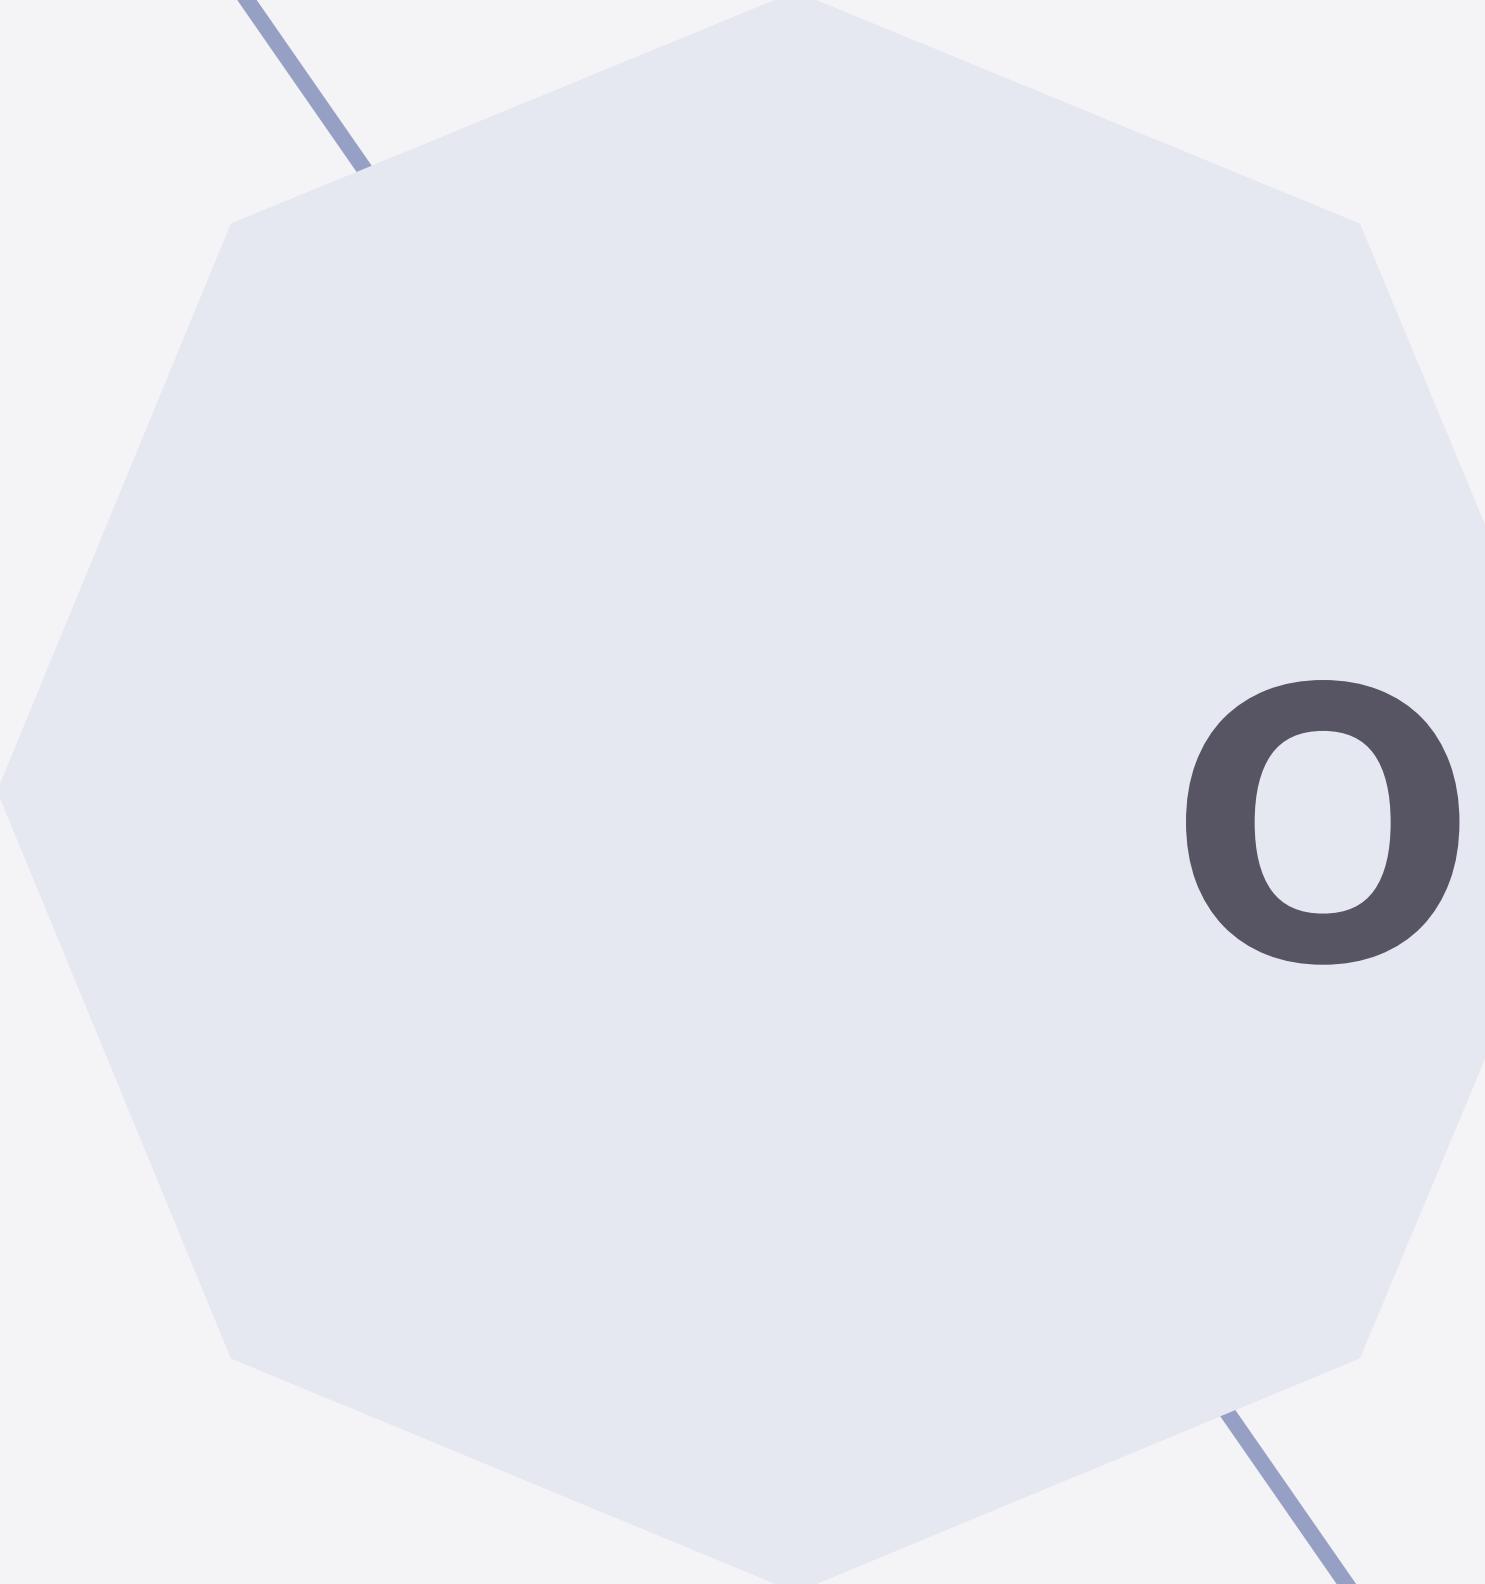


demo

Derived variables

- ▶ obtain samples from model parameters
- ▶ apply (deterministic) function to each sample
 - to derive (deterministically) a new model variable
- ▶ violà: samples from the posterior of a new “derived variable”

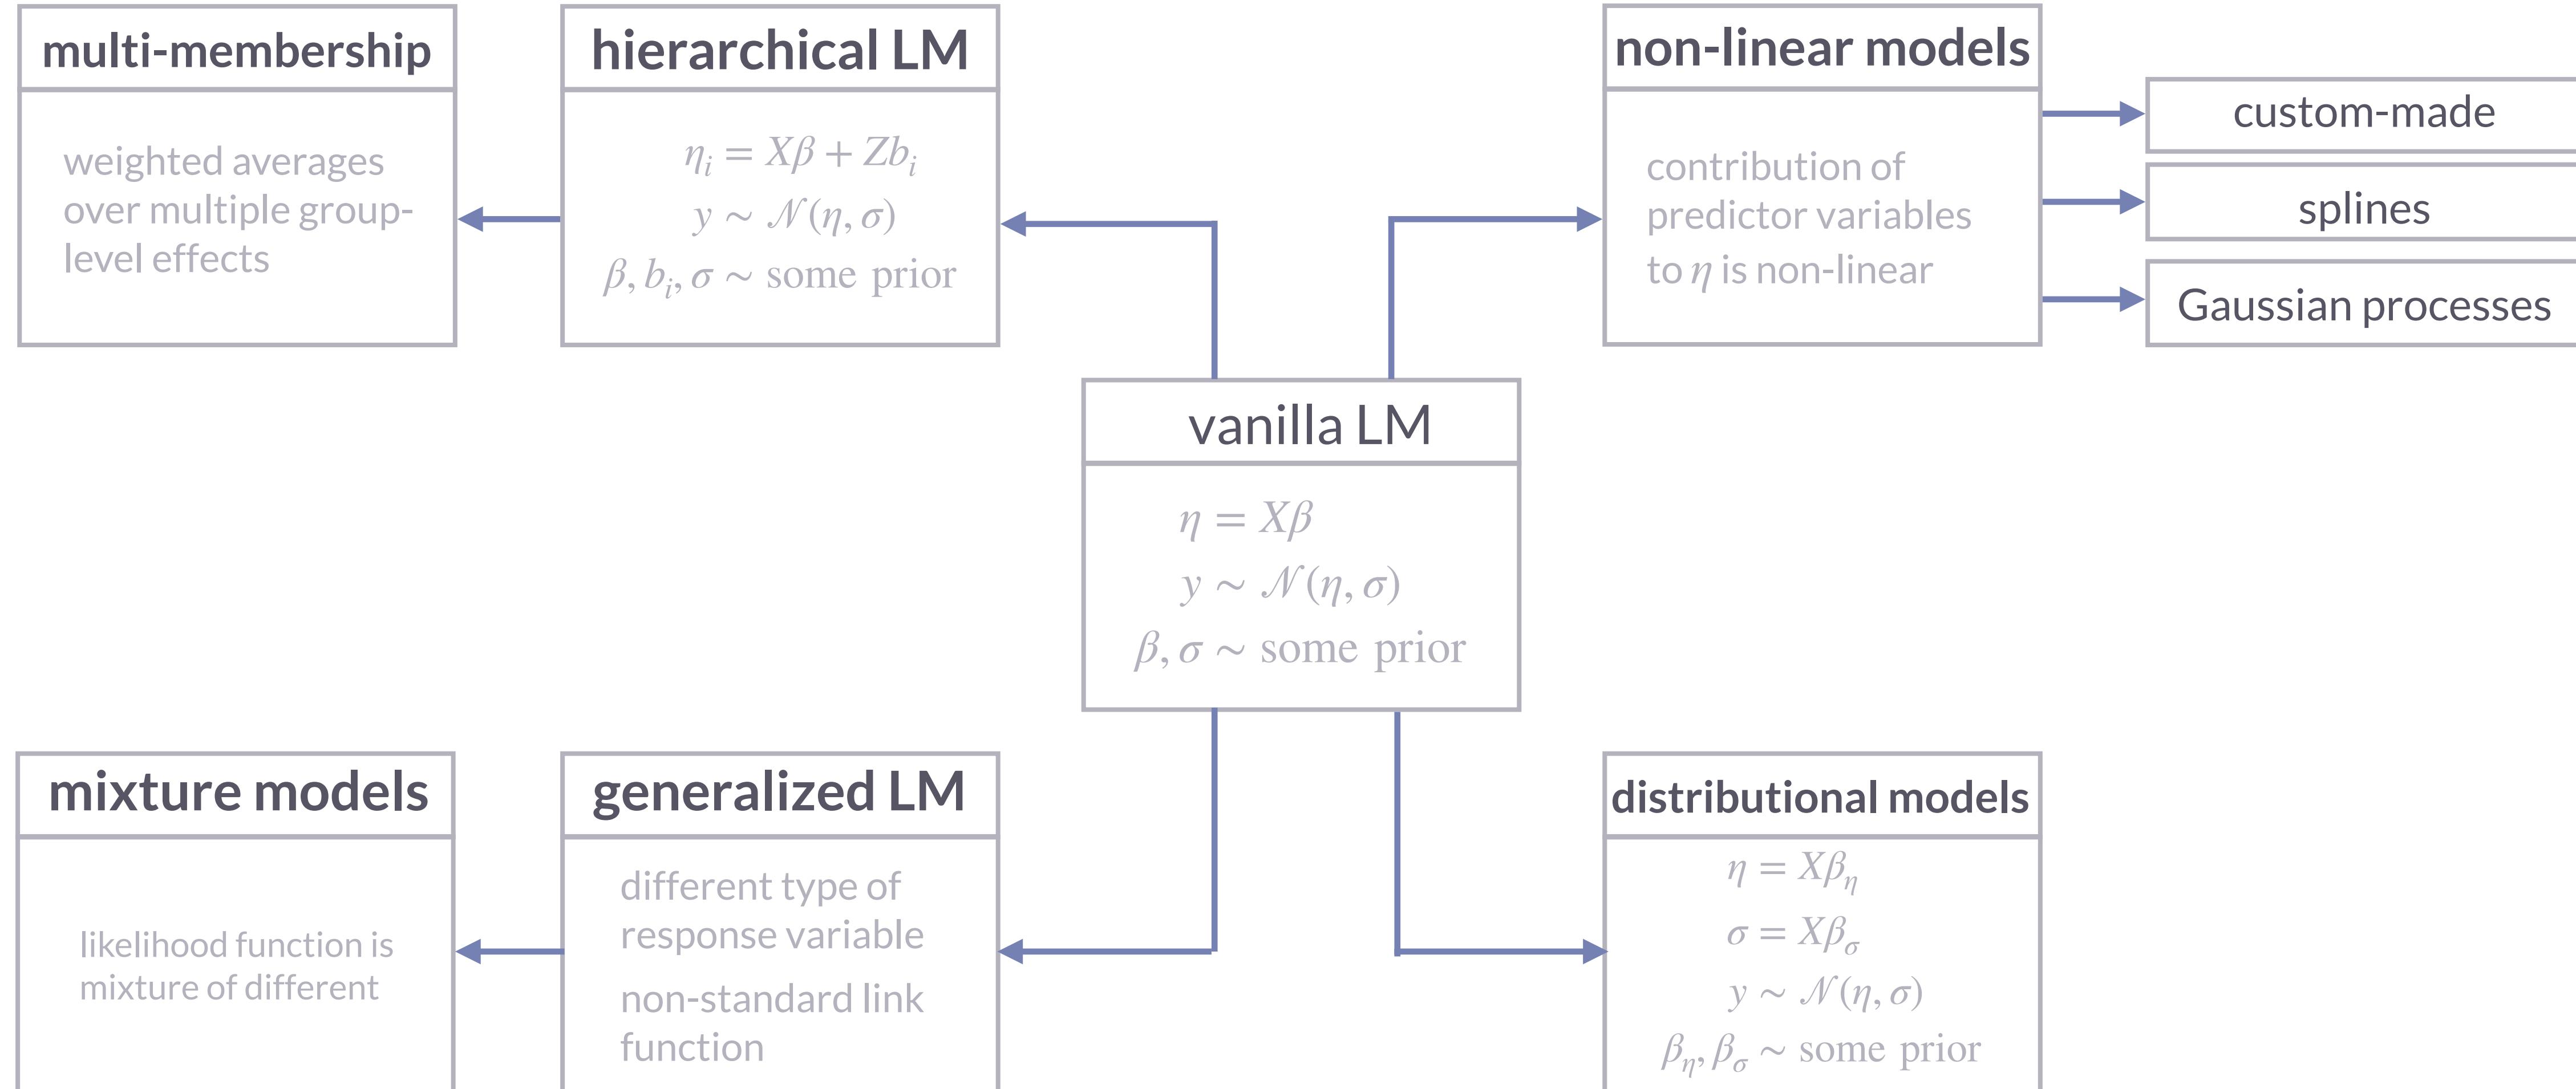




outlook

Roadmap “beyond vanilla”

common extensions of linear regression modeling



Three pillars of BDA

1. parameter estimation / inference [which parameter values are credible given data and model?]

$$\underbrace{P(\theta | D)}_{\text{posterior}} \propto \underbrace{P(\theta)}_{\text{prior}} \times \underbrace{P(D | \theta)}_{\text{likelihood}}$$

2. predictions [which future data observations are likely given my model?]

a. prior

$$P(D_{\text{pred}}) = \int P(\theta) P(D_{\text{pred}} | \theta) d\theta$$

b. posterior

$$P(D_{\text{pred}} | D_{\text{obs}}) = \int P(\theta | D_{\text{obs}}) P(D_{\text{pred}} | \theta) d\theta$$

3. model comparison [which model of two models is more likely to have generated the data?]

$$\frac{\underbrace{P(M_1 | D)}_{\text{posterior odds}}}{\underbrace{P(M_2 | D)}_{\text{posterior odds}}} = \underbrace{\frac{P(D | M_1)}{P(D | M_2)}}_{\text{Bayes factor}} \frac{\underbrace{P(M_1)}_{\text{prior odds}}}{\underbrace{P(M_2)}_{\text{prior odds}}}$$